

# Stochastic Tagging



**Bayes Rule**

**Markov Chains**

**HMM Tagging**

**Viterbi Algorithm**

B.Senthil Kumar, Asst. Prof  
Natural Language Processing  
B.E. (CSE) VII Sem

# Topics

---

- Probability
- Conditional Probability
- Bayes Rule
- HMM tagging
- Markov Chains
- Hidden Markov Models

# Introduction to Probability

---

- Experiment (trial)
  - Repeatable procedure with well-defined possible outcomes
- Sample Space (S)
  - the set of all possible outcomes
  - *finite or infinite*
  - Example
    - coin toss experiment
    - possible outcomes:  $S = \{\text{heads, tails}\}$
  - Example
    - die toss experiment
    - possible outcomes:  $S = \{1, 2, 3, 4, 5, 6\}$

# Introduction to Probability

---

- Definition of sample space depends on what we are asking
  - Sample Space (S): the set of all possible outcomes
  - Example
    - die toss experiment for whether the number is even or odd
    - possible outcomes: {even,odd}
    - *not* {1,2,3,4,5,6}

# More definitions

---

## □ Events

- an **event** is any subset of outcomes from the **sample space**

## □ Example

- die toss experiment
- let A represent the event such that the outcome of the die toss experiment is divisible by 3
- $A = \{3, 6\}$
- A is a subset of the sample space  $S = \{1, 2, 3, 4, 5, 6\}$

# Definition of Probability

---

- ▣ The probability law assigns to an event a nonnegative number
- ▣ Called  $P(A)$
- ▣ Also called the probability A
- ▣ That encodes our knowledge or belief about the collective likelihood of all the elements of A
- ▣ Probability law must satisfy certain properties

# Probability Axioms

---

## □ Nonnegativity

- $P(A) \geq 0$ , for every event  $A$

## □ Additivity

- If  $A$  and  $B$  are two disjoint events, then the probability of their union satisfies:
  - $P(A \cup B) = P(A) + P(B)$

## □ Normalization

- The probability of the entire sample space  $S$  is equal to 1, i.e.  $P(S) = 1$ .

# An example

---

- ▣ An experiment involving a single coin toss

There are two possible outcomes, H and T

Sample space S is {H,T}

If coin is fair, should assign equal probabilities to {H,T} outcomes

Since they have to sum to 1

$$P(\{H\}) = 0.5 \text{ and } P(\{T\}) = 0.5$$

$$P(\{H,T\}) = P(\{H\}) + P(\{T\}) = 1.0$$



# Another example

---

- Experiment involving 3 coin tosses
- Outcome is a 3-long string of H or T
- $S = \{HHH, HHT, HTH, HTT, THH, THT, TTH, TTT\}$
- Assume each outcome is equiprobable
  - “Uniform distribution”
- What is probability of the event that **exactly 2 heads** occur?

$$A = \{HHT, HTH, THH\}$$

3 events

$$P(A) = P(\{HHT\}) + P(\{HTH\}) + P(\{THH\})$$

union of the prob of individual events

$$= 1/8 + 1/8 + 1/8$$

$$= 3/8$$

# Probability definitions

---

□ In summary:

$$P(E) = \frac{\text{number of outcomes corresponding to event } E}{\text{total number of outcomes}}$$

Probability of drawing a spade from 52 well-shuffled playing cards:

$$\frac{13}{52} = \frac{1}{4} = 0.25$$

# Probability and part of speech tags

---

- What's the probability of a random word (from a random dictionary page) being a verb?

$$P(\text{drawing a verb}) = \frac{\text{of ways to get a verb}}{\text{all words}}$$

- How to compute each of these?

All words = just count all the words in the dictionary

# of ways to get a verb: # of words which are verbs!

If a dictionary has 50,000 entries, and 10,000 are verbs....

$P(V)$  is  $10000/50000 = 1/5 = .20$

# Conditional Probability

---

- A way to reason about the outcome of an experiment based on partial information
  - In a word guessing game the first letter for the word is a “t”. What is the likelihood that the second letter is an “h”?
  - How likely is it that a person has a disease given that a medical test was negative?
  - A spot shows up on a radar screen. How likely is it that it corresponds to an aircraft?

# An intuition

---

- Let's say A is "it's raining".
- Let's say  $P(A)$  in dry Florida is .01
- Let's say B is "it was sunny ten minutes ago"
- $P(A|B)$  means "what is the probability of it raining now if it was sunny 10 minutes ago"
- $P(A|B)$  is probably way less than  $P(A)$
- Perhaps  $P(A|B)$  is .0001
- Intuition: The knowledge about B should change our estimate of the probability of A.

# More precisely

---

- Given an experiment, a corresponding sample space  $S$ , and a probability law
- Suppose we know that the outcome is some event  $B$
- We want to quantify the likelihood that the outcome also belongs to some other event  $A$
- We need a new probability law that gives us the conditional probability of  $A$  given  $B$
- $P(A|B)$

# Conditional Probability

---

- let A and B be events in the sample space
- $P(A|B)$  = the conditional *probability* of event A *occurring given* some fixed event B *occurring*
- *definition:*  $P(A|B) = P(A \cap B) / P(B)$


# Conditional probability

---

▣  $P(A|B) = P(A \cap B)/P(B)$

▣ Or

$$P(A|B) = \frac{P(A, B)}{P(B)}$$



*Note:  $P(A, B) = P(A|B) \cdot P(B)$*

*Also:  $P(A, B) = P(B, A)$*



# Independence

---

- ▣ What is  $P(A,B)$  if  $A$  and  $B$  are independent?
- ▣  $P(A,B)=P(A) \cdot P(B)$  iff  $A,B$  independent.

$$P(\text{heads,tails}) = P(\text{heads}) \cdot P(\text{tails}) = 0.5 \cdot 0.5 = 0.25$$

*Note:  $P(A|B)=P(A)$  iff  $A,B$  independent*

*Also:  $P(B|A)=P(B)$  iff  $A,B$  independent*

# Bayes Theorem

---

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

- **Idea:** The probability of an A conditional on another event B is generally different from the probability of B conditional on A. There is a definite relationship between the two.

# Deriving Bayes Rule

---

The probability of event A given event B is

$$P(A|B) = \frac{P(A,B)}{P(B)}$$

# Deriving Bayes Rule

---

The probability of event B given event A is

$$P(B|A) = \frac{P(A,B)}{P(A)}$$

# Deriving Bayes Rule

---

$$P(A|B) = \frac{P(A,B)}{P(B)}$$



$$P(A|B)P(B) = P(A,B)$$

$$P(B|A) = \frac{P(A,B)}{P(A)}$$



$$P(B|A)P(A) = P(A,B)$$

# Deriving Bayes Rule

---

$$P(A|B) = \frac{P(A,B)}{P(B)}$$



$$P(A|B)P(B) = P(A,B)$$



$$P(A|B)P(B) = P(B|A)P(A)$$



$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

$$P(B|A) = \frac{P(A,B)}{P(A)}$$



$$P(B|A)P(A) = P(A,B)$$



$$P(A|B)P(B) = P(B|A)P(A)$$

# Bayes Rule

---

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

the theorem may be paraphrased as:

Conditional/Posterior probability =

(LIKELIHOOD multiplied by PRIOR) divided by NORMALIZING  
CONSTANT

# Hidden Markov Model (HMM) Tagging

---

- ▣ Using an HMM to do POS tagging
- ▣ HMM is a special case of Bayesian inference
- ▣ It is also related to the “noisy channel” model in ASR (Automatic Speech Recognition)



# POS tagging as a sequence classification task

---

- Given a sentence (an “observation” or “sequence of observations”)
  - *Secretariat is expected to race tomorrow*
  - sequence of  $n$  words  $w_1 \dots w_n$ .
- What is the best sequence of tags which corresponds to this sequence of observations?
- Probabilistic/Bayesian view:
  - Consider all possible sequences of tags
  - Out of this universe of sequences, choose the tag sequence which is most probable given the observation sequence of  $n$  words  $w_1 \dots w_n$ .

# Getting to HMM

---

- Let  $T = t_1, t_2, \dots, t_n$
- Let  $W = w_1, w_2, \dots, w_n$
- Goal: Out of all sequences of tags  $t_1 \dots t_n$ , get the the most probable sequence of POS tags  $T$  underlying the observed sequence of

words  $w_1, w_2, \dots, w_n$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n)$$

- Hat ^ means “our estimate of the best = the most probable tag sequence”
- $\operatorname{Argmax}_x f(x)$  means “the  $x$  such that  $f(x)$  is maximized”  
it maximizes our estimate of the best tag sequence

# Getting to HMM

---

- This equation is guaranteed to give us the best tag sequence

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n)$$

- But how do we make it operational? How do we compute this value?
- Intuition of Bayesian classification:
  - Use Bayes rule to transform it into a set of other probabilities that are easier to compute
  - Thomas Bayes: British mathematician (1702-1761)

# Bayes Rule

---

$$P(x|y) = \frac{P(y|x)P(x)}{P(y)}$$

Breaks down any conditional probability  $P(x|y)$  into three other probabilities

$P(x|y)$ : The conditional probability of an event  $x$  assuming that  $y$  has occurred

# Bayes Rule

---

$$P(x|y) = \frac{P(y|x)P(x)}{P(y)}$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \frac{P(w_1^n | t_1^n) P(t_1^n)}{P(w_1^n)}$$

We can drop the denominator: it does not change for each tag sequence; we are looking for the best tag sequence for the same observation, for the same fixed set of words

# Bayes Rule

---

$$P(x|y) = \frac{P(y|x)P(x)}{P(y)}$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \frac{P(w_1^n | t_1^n) P(t_1^n)}{P(w_1^n)}$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(w_1^n | t_1^n) P(t_1^n)$$

# Likelihood and prior

---

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \overbrace{P(w_1^n | t_1^n)}^{\text{likelihood}} \overbrace{P(t_1^n)}^{\text{prior}}$$

# Likelihood and prior

## Further Simplifications

---

1. the probability of a word appearing depends only on its own POS tag, i.e, independent of other words around it

$$P(w_1^n | t_1^n) \approx \prod_{i=1}^n P(w_i | t_i)$$

2. **BIGRAM** assumption: the probability of a tag appearing depends only on the previous tag

$$P(t_1^n) \approx \prod_{i=1}^n P(t_i | t_{i-1})$$

3. The most probable tag sequence estimated by the bigram tagger

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n) \approx \operatorname{argmax}_{t_1^n} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

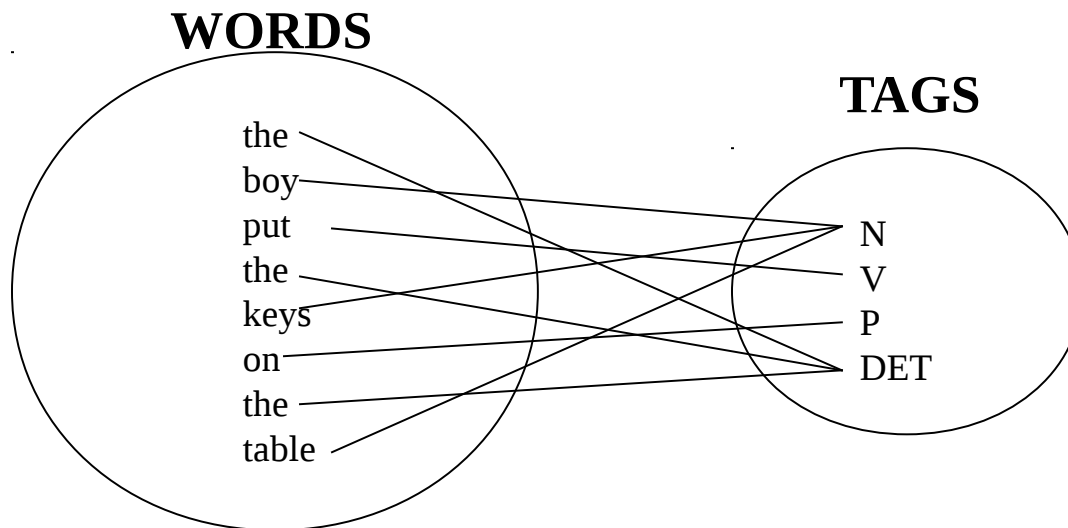


# Likelihood ratio

## Further Simplifications

1. the probability of a word appearing depends only on its own POS tag, i.e, independent of other words around it

$$P(w_1^n | t_1^n) \approx \prod_{i=1}^n P(w_i | t_i)$$



# Prior probability

## Further Simplifications

---

2. **BIGRAM** assumption: the probability of a tag appearing depends only on the previous tag

$$P(t_1^n) \approx \prod_{i=1}^n P(t_i | t_{i-1})$$

Bigrams are groups of two written letters, two syllables, or two words; they are a special case of N-gram.

Bigrams are used as the basis for simple statistical analysis of text

The bigram assumption is related to the first-order Markov assumption

# Likelihood and prior

## Further Simplifications

3. The most probable tag sequence estimated by the bigram tagger

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n) \approx \operatorname{argmax}_{t_1^n} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \overbrace{P(w_1^n | t_1^n)}^{\text{likelihood}} \overbrace{P(t_1^n)}^{\text{prior}}$$
$$P(w_1^n | t_1^n) \approx \prod_{i=1}^n P(w_i | t_i) \qquad P(t_1^n) \approx \prod_{i=1}^n P(t_i | t_{i-1})$$

bigram assumption

# Two kinds of probabilities (1)

---

- Tag transition probabilities  $p(t_i|t_{i-1})$ 
  - Determiners likely to precede adjs and nouns
    - That/DT flight/NN
    - The/DT yellow/JJ hat/NN
    - So we expect  $P(\text{NN}|\text{DT})$  and  $P(\text{JJ}|\text{DT})$  to be high
    - But  $P(\text{DT}|\text{JJ})$  to be:?

# Two kinds of probabilities (1)

---

- Tag transition probabilities  $p(t_i|t_{i-1})$ 
  - Compute  $P(NN|DT)$  by counting in a labeled corpus:

$$P(t_i|t_{i-1}) = \frac{C(t_{i-1}, t_i)}{C(t_{i-1})}$$

$$P(NN|DT) = \frac{C(DT, NN)}{C(DT)} = \frac{56,509}{116,454} = .49$$

# of times DT is followed by NN

# Two kinds of probabilities (2)

---

□ Word **likelihood probabilities**  $p(w_i|t_i)$

- $P(\text{is}|\text{VBZ})$  = probability of VBZ (3sg pres verb) being “is”

If we were expecting a third person singular verb, how likely is it that this verb would be *is*?

- Compute  $P(\text{is}|\text{VBZ})$  by counting in a labeled corpus:

$$P(w_i|t_i) = \frac{C(t_i, w_i)}{C(t_i)}$$

$$P(\text{is}|\text{VBZ}) = \frac{C(\text{VBZ}, \text{is})}{C(\text{VBZ})} = \frac{10,073}{21,627} = .47$$



# An Example: the verb “race”

---

- Secretariat/**NNP** is/**VBZ** expected/**VBN** to/**TO** **race**/**VB**  
tomorrow/**NR**
- People/**NNS** continue/**VB** to/**TO** inquire/**VB** the/**DT** reason/**NN**  
for/**IN** the/**DT** **race**/**NN** for/**IN** outer/**JJ** space/**NN**
- How do we pick the right tag?

# Disambiguating “race”

---





# Disambiguating “race”

---

- ▣  $P(\text{NN}|\text{TO}) = .00047$

$$P(\text{VB}|\text{TO}) = .83$$

The tag transition probabilities  $P(\text{NN}|\text{TO})$  and  $P(\text{VB}|\text{TO})$ :

‘How likely are we to expect verb/noun given the previous tag TO?’

- ▣  $P(\text{race}|\text{NN}) = .00057$

$$P(\text{race}|\text{VB}) = .00012$$

Lexical likelihoods from the Brown corpus for ‘race’ given a POS tag  
NN or VB.

# Disambiguating “race”

---

▣  $P(\text{NR}|\text{VB}) = .0027$

$$P(\text{NR}|\text{NN}) = .0012$$

tag sequence probability of an adverb occurring given the previous  
tag verb / noun

$$P(\text{VB}|\text{TO})P(\text{NR}|\text{VB})P(\text{race}|\text{VB}) = .83 \times .0027 \times .00012 = .00000027$$

$$P(\text{NN}|\text{TO})P(\text{NR}|\text{NN})P(\text{race}|\text{NN}) = .$$

$$00047 \times .0012 \times .00057 = .00000000032$$

Multiply the lexical likelihoods with the tag sequence probabilities:

Hence the *race* is tagged as **verb** !

# Hidden Markov Models

---

- ▣ What we've described with these two kinds of probabilities is a Hidden Markov Model (HMM)
- ▣ Let's just spend a bit of time tying this into the model
- ▣ In order to define HMM, we will first introduce the Markov Chain, or observable Markov Model.







# Definitions


---

- A **weighted finite-state automaton** adds probabilities to the arcs
  - The sum of the probabilities leaving any arc must sum to one
- A **Markov chain** is a special case of a WFST in which the input sequence uniquely determines which states the automaton will go through
- Markov chains can't represent inherently ambiguous problems
  - Useful for assigning probabilities to unambiguous sequences

# Markov chain = “First-order observable Markov Model”

**Table 1:** Probabilities  $p(q_{n+1}|q_n)$  of tomorrow's weather based on today's weather

	Tomorrow's weather		
Today's weather			
	0.8	0.05	0.15
	0.2	0.6	0.2
	0.2	0.3	0.5



# Hidden Markov Model

---

- For Markov chains, the output symbols are the same as the states.
  - See **rainy** weather: we're in state **rainy**
- But in part-of-speech tagging (and other things)
  - The output symbols are **words**
  - But the hidden states are **part-of-speech tags**
- So we need an extension!
- A **Hidden Markov Model** is an extension of a Markov chain that allows both observed events (like words as input) and hidden events (like pos tags)

# Hidden Markov Model

---

- ▣ States  $Q = q_1, q_2 \dots q_N$ ;
- ▣ Observations  $O = o_1, o_2 \dots o_N$ ;
  - Each observation is a symbol from a vocabulary  $V = \{v_1, v_2, \dots v_V\}$
- ▣ Transition probabilities (prior)
  - Transition probability matrix  $A = \{a_{ij}\}$ . Probability of moving from state  $i$  to state  $j$
- ▣ Observation likelihoods (likelihood)
  - probability matrix  $B = \{b_i(o_t)\}$   
a sequence of observation likelihoods, each expressing the probability of an observation  $O_t$  being generated from a state  $i$ .
- ▣ A special start and end state

# HMM Taggers

---


- An HMM has two kinds of probabilities
  - A transition probabilities (PRIOR) (slide 35)
  - B observation likelihoods (LIKELIHOOD) (slide 35)
- HMM Taggers choose the tag sequence which maximizes the product of word likelihood and tag sequence probability

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \overbrace{P(w_1^n | t_1^n)}^{\text{likelihood}} \overbrace{P(t_1^n)}^{\text{prior}}$$

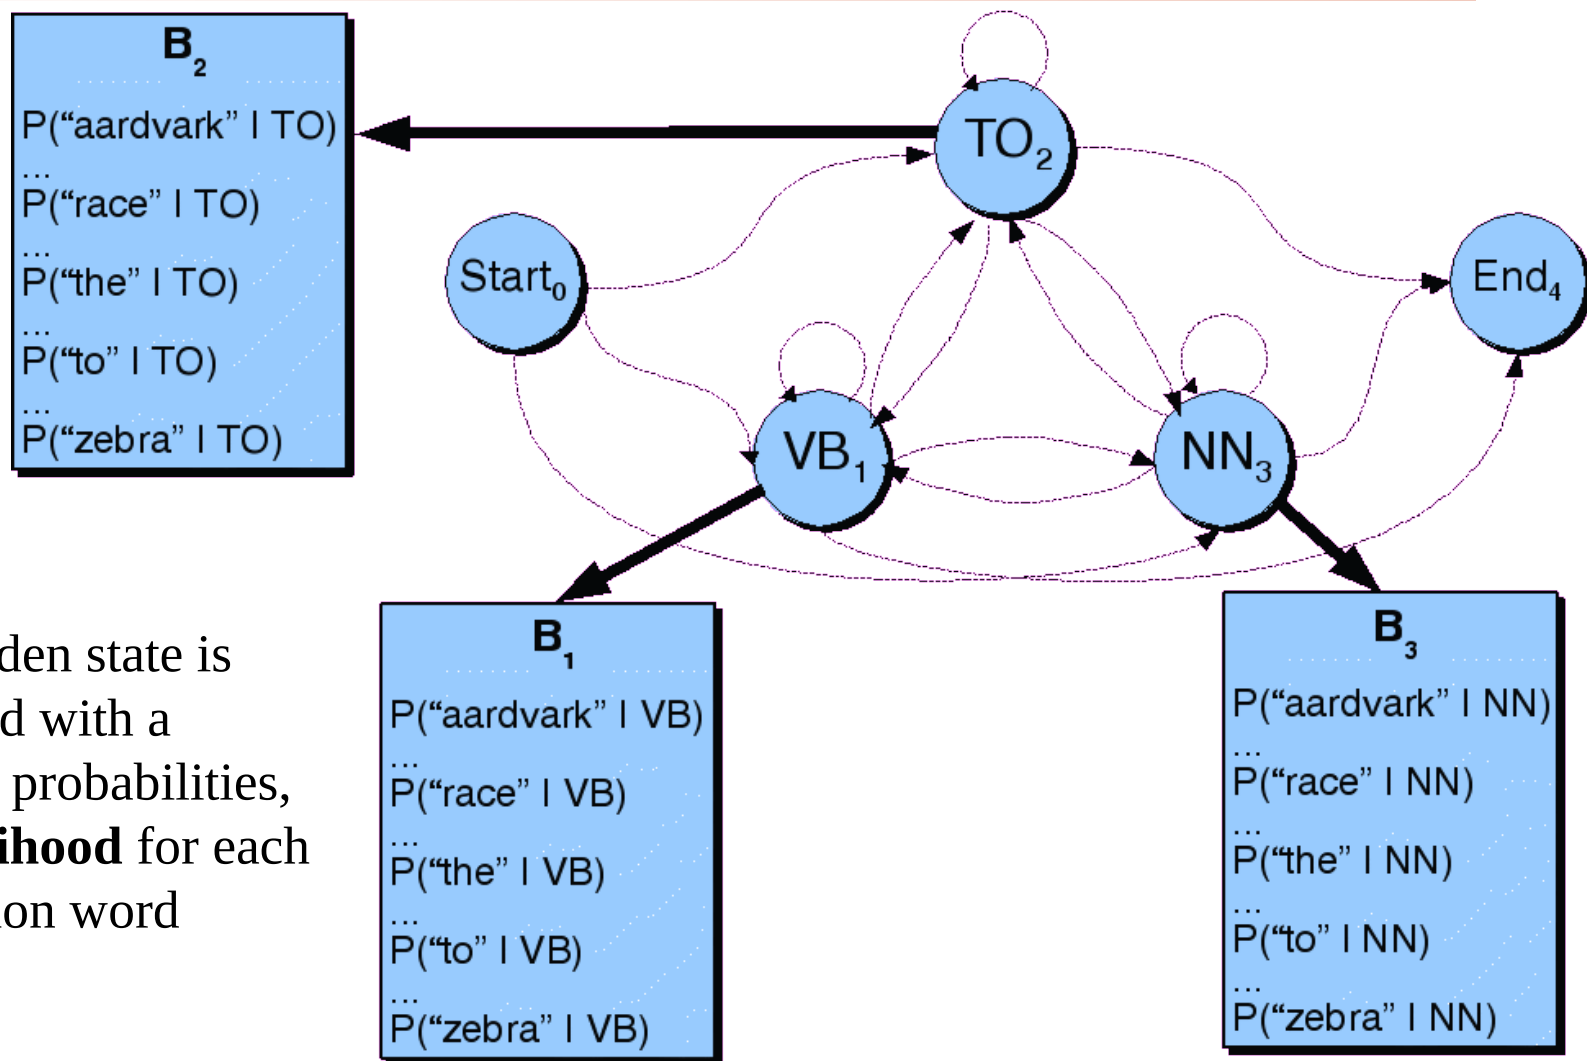


# Markov chain corresponding to hidden states of HMM, showing A probs

Transition probabilities are used to compute **prior probability**



# B observation likelihoods for HMM



Each hidden state is associated with a vector of probabilities, one **likelihood** for each observation word

# Viterbi Algorithm

---

- Any model that contains hidden variables, the task of determining which sequence of variables is the underlying source of some sequence of observations is called the **decoding task**.
- The Viterbi algorithm is the most common decoding algorithm used for HMMs.
- Let HMM be defined by two tables: Fig 5.15, Fig 5.16

# Viterbi Algorithm

- Fig 5.15 – the *transition probabilities* between hidden states (pos tags) –  $a_{ij}$  probability

	<b>VB</b>	<b>TO</b>	<b>NN</b>	<b>PPSS</b>
<b>&lt;s&gt;</b>	.019	.0043	.041	.067
<b>VB</b>	.0038	.035	.047	.0070
<b>TO</b>	.83	0	.00047	0
<b>NN</b>	.0040	.016	.087	.0045
<b>PPSS</b>	.23	.00079	.0012	.00014

**Figure 5.15** Tag transition probabilities (the  $a$  array,  $p(t_i|t_{i-1})$ ) computed from the 87-tag Brown corpus without smoothing. The rows are labeled with the conditioning event; thus  $P(PPSS|VB)$  is .0070. The symbol  $<s>$  is the start-of-sentence symbol.

# Viterbi Algorithm

- Fig 5.16 – the observation likelihoods of words given tags (i.e.,  $b_i(O_t)$  probabilities)

	<b>I</b>	<b>want</b>	<b>to</b>	<b>race</b>
<b>VB</b>	0	.0093	0	.00012
<b>TO</b>	0	0	.99	0
<b>NN</b>	0	.000054	0	.00057
<b>PPSS</b>	.37	0	0	0

**Figure 5.16** Observation likelihoods (the  $b$  array) computed from the 87-tag Brown corpus without smoothing.

# Viterbi Algorithm

---

- Viterbi sets up a probability matrix – one **column** for each **observation t** and one **row** for each **state** in the state graph.
- The algorithm first creates N state columns.
- Begin in first column and move on column by column.
- For each cell,  $viterbi[s,t]$  is computed by taking the maximum over the extensions of all the paths that lead to the current cell.

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t)$$

# Viterbi Algorithm

---

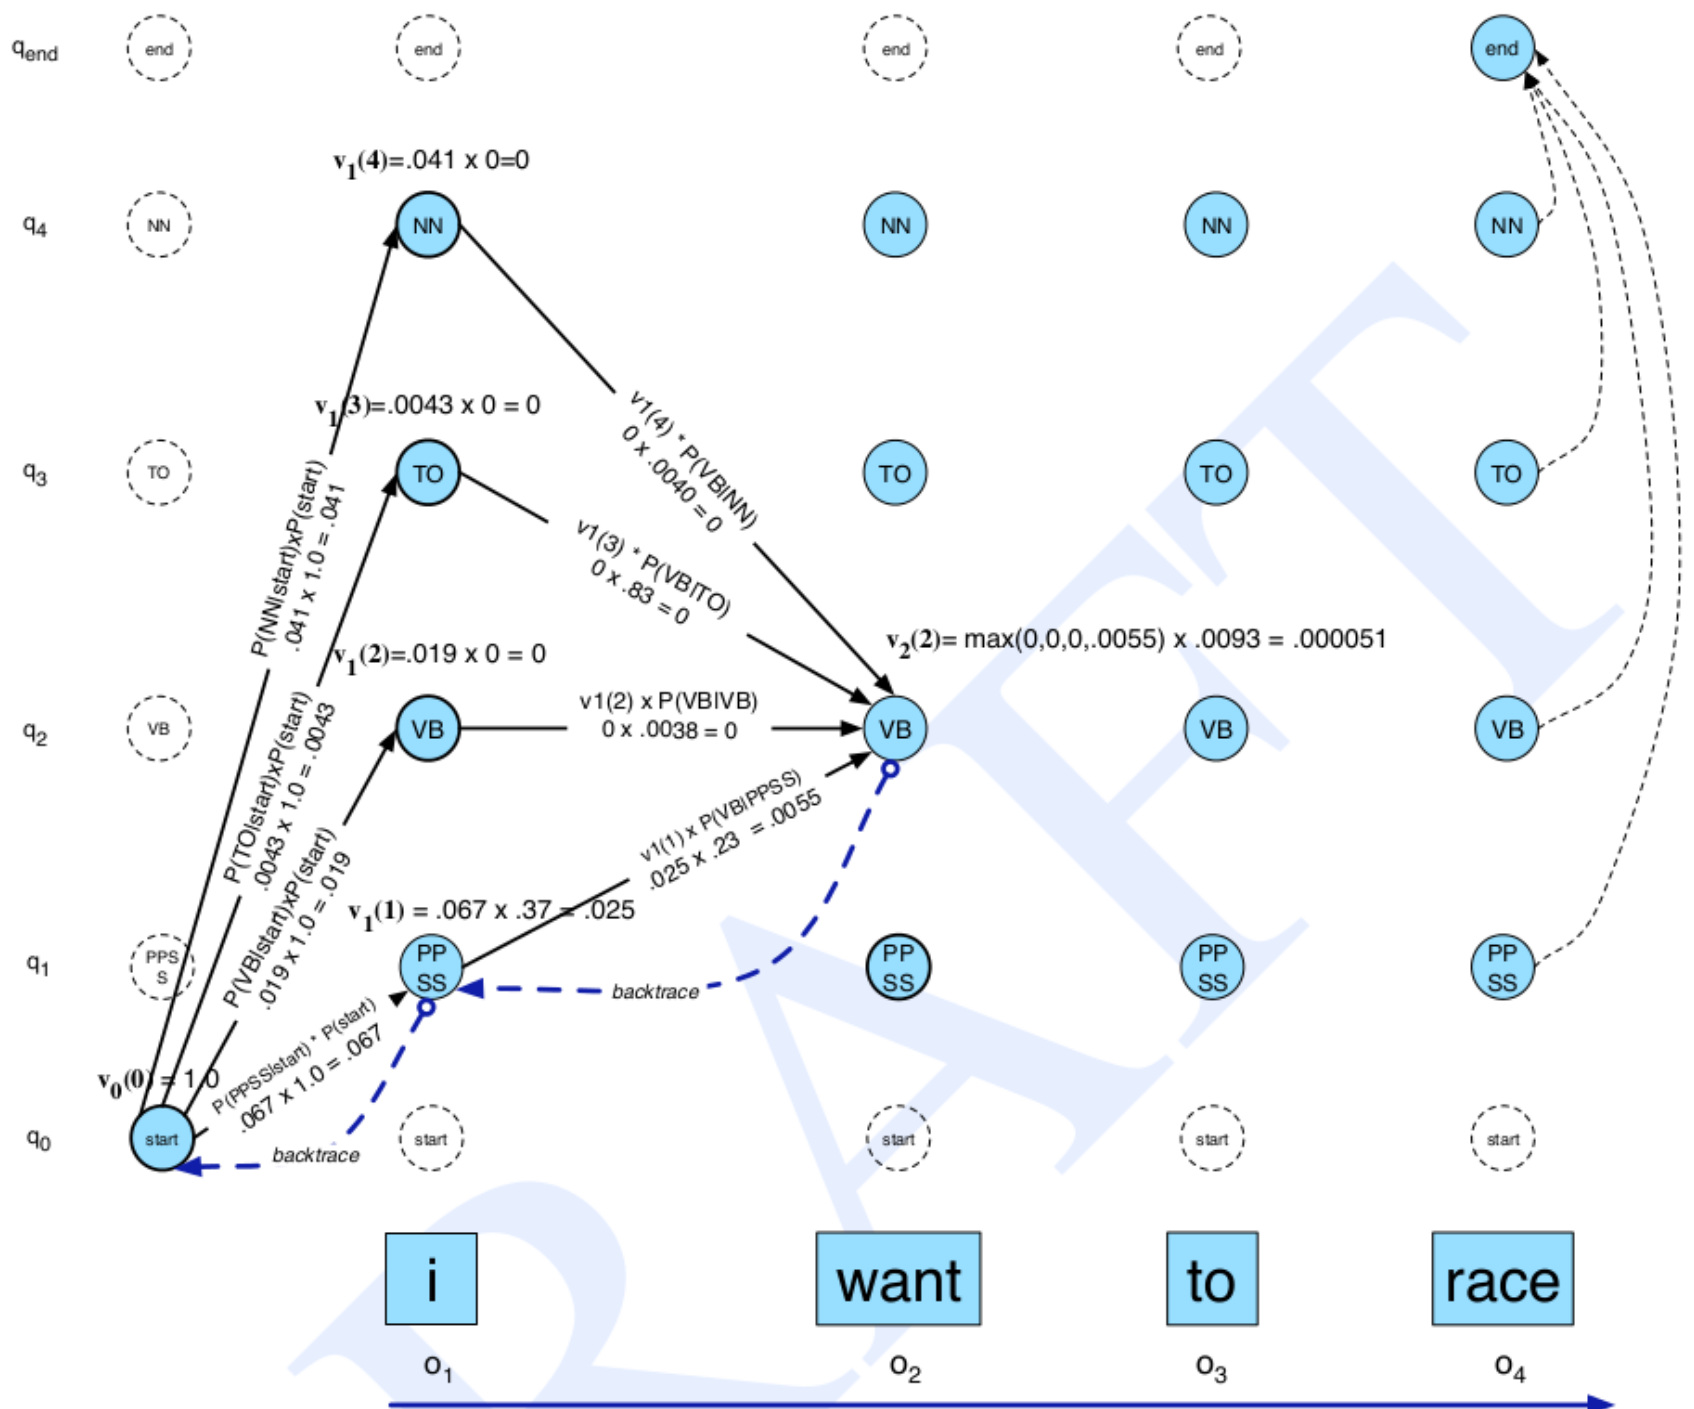
- ▣ For each cell,  $viterbi[s,t]$  is computed by taking the maximum over the extensions of all the paths that lead to the current cell.

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t)$$

$v_{t-1}(i)$  the **previous Viterbi path probability** from the previous time step

$a_{ij}$  the **transition probability** from previous state  $q_i$  to current state  $q_j$

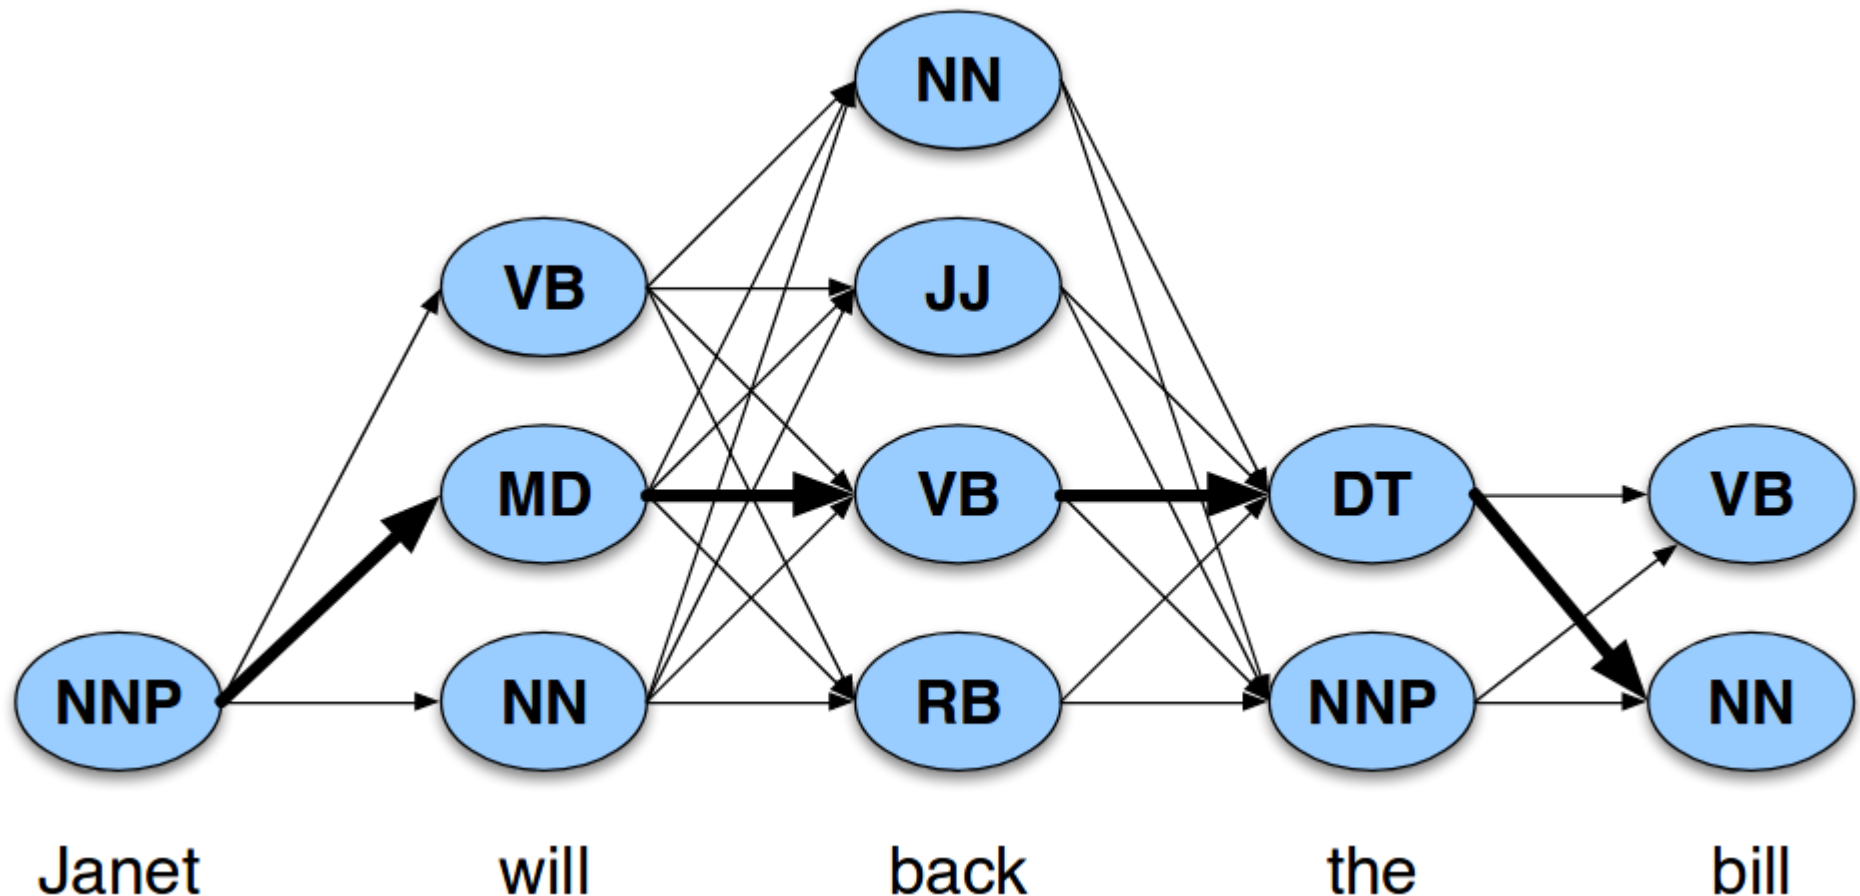
$b_j(o_t)$  the **state observation likelihood** of the observation symbol  $o_t$  given the current state  $j$





# Viterbi Algorithm

- Example 2: *Janet will back the bill*

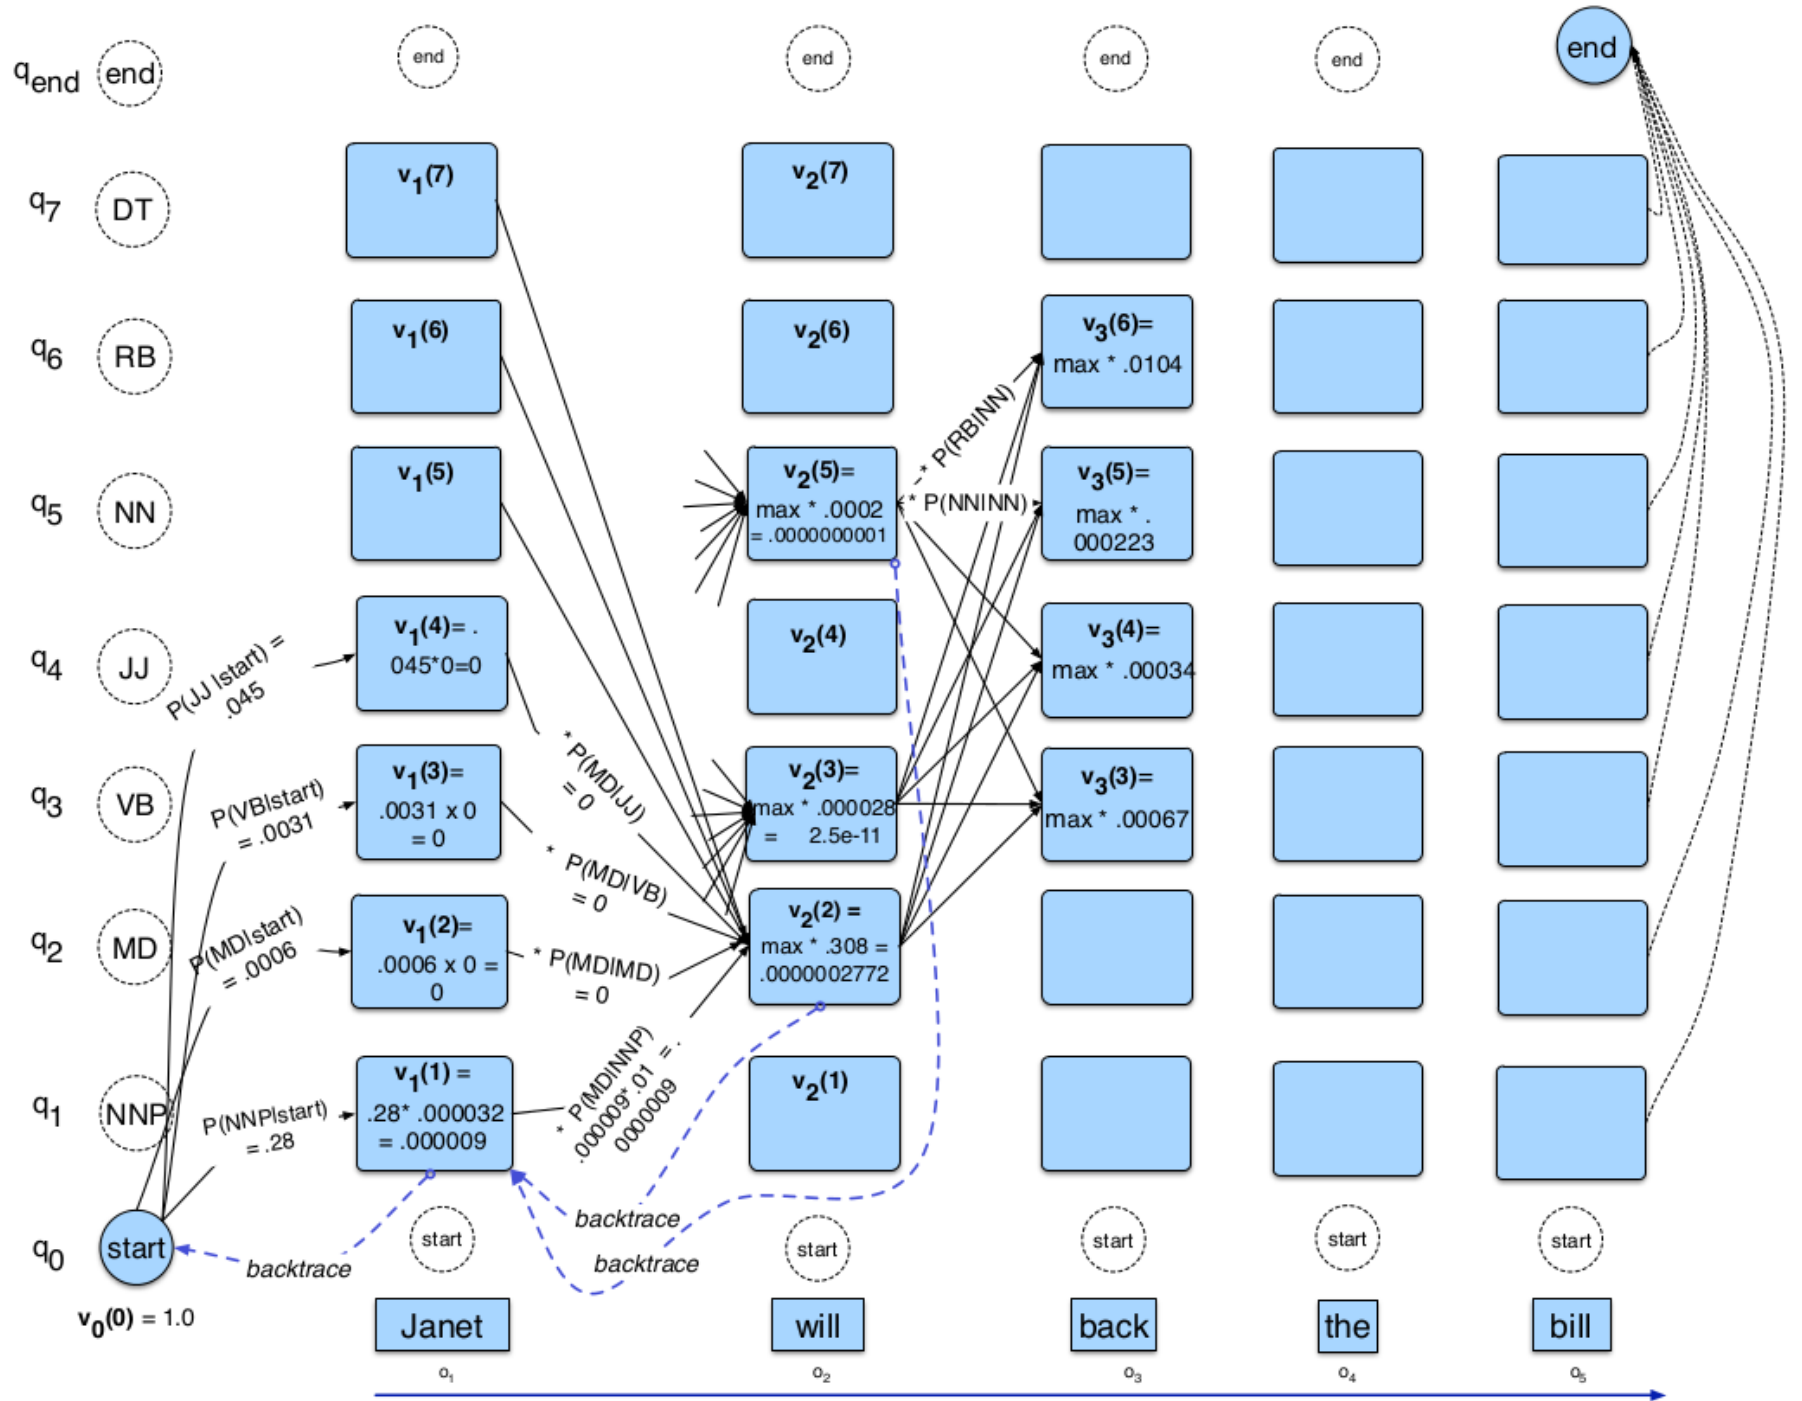


	<b>NNP</b>	<b>MD</b>	<b>VB</b>	<b>JJ</b>	<b>NN</b>	<b>RB</b>	<b>DT</b>
< <i>s</i> >	0.2767	0.0006	0.0031	0.0453	0.0449	0.0510	0.2026
<b>NNP</b>	0.3777	0.0110	0.0009	0.0084	0.0584	0.0090	0.0025
<b>MD</b>	0.0008	0.0002	0.7968	0.0005	0.0008	0.1698	0.0041
<b>VB</b>	0.0322	0.0005	0.0050	0.0837	0.0615	0.0514	0.2231
<b>JJ</b>	0.0366	0.0004	0.0001	0.0733	0.4509	0.0036	0.0036
<b>NN</b>	0.0096	0.0176	0.0014	0.0086	0.1216	0.0177	0.0068
<b>RB</b>	0.0068	0.0102	0.1011	0.1012	0.0120	0.0728	0.0479
<b>DT</b>	0.1147	0.0021	0.0002	0.2157	0.4744	0.0102	0.0017

**Figure 10.5** The  $A$  transition probabilities  $P(t_i|t_{i-1})$  computed from the WSJ corpus without smoothing. Rows are labeled with the conditioning event; thus  $P(VB|MD)$  is 0.7968.

	<b>Janet</b>	<b>will</b>	<b>back</b>	<b>the</b>	<b>bill</b>
<b>NNP</b>	0.000032	0	0	0.000048	0
<b>MD</b>	0	0.308431	0	0	0
<b>VB</b>	0	0.000028	0.000672	0	0.000028
<b>JJ</b>	0	0	0.000340	0.000097	0
<b>NN</b>	0	0.000200	0.000223	0.000006	0.002337
<b>RB</b>	0	0	0.010446	0	0
<b>DT</b>	0	0	0	0.506099	0

**Figure 10.6** Observation likelihoods  $B$  computed from the WSJ corpus without smoothing.



# References

---

- ▣ Chapter 5. Part-of-Speech Tagging  
Speech and Language Processing by Jurafsky and Martin