

Introduction to OpenMP

- Introduction
- OpenMP basics
- OpenMP directives, clauses, and library routines

Programming with Threads

Several thread libraries, more being created

- PThreads is the POSIX Standard
 - Relatively low level
 - Programmer expresses thread management and coordination
 - Programmer decomposes parallelism and manages schedule
 - Portable but possibly slow
 - Most widely used for systems-oriented code, and also used for some kinds of application code
- OpenMP is newer standard
 - Higher-level support for scientific programming on shared memory architectures

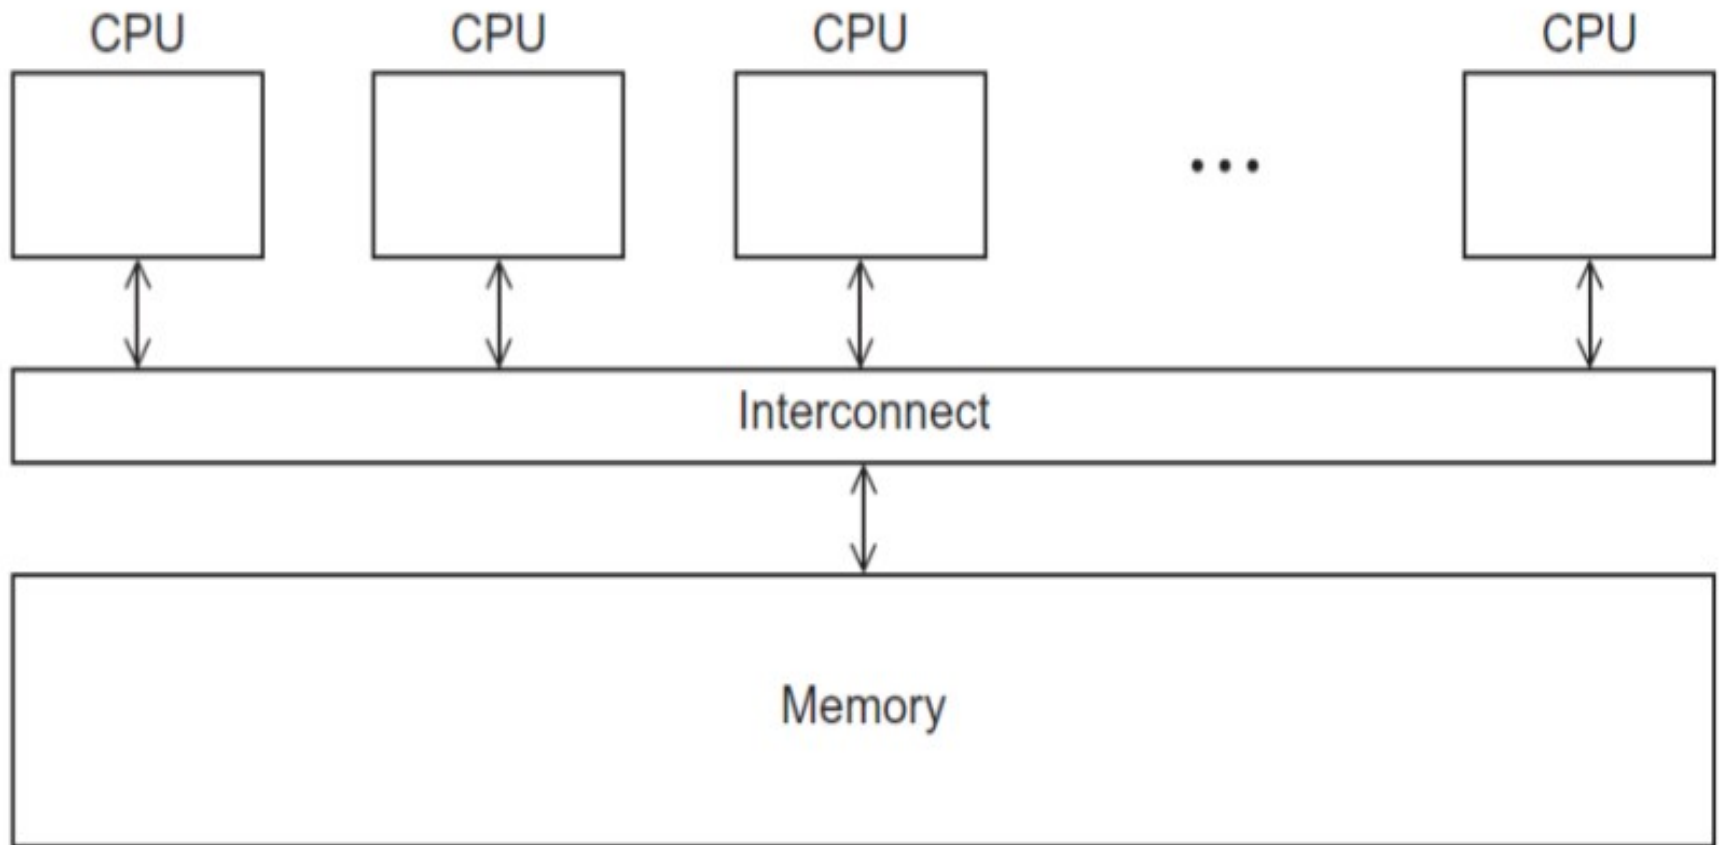
Introduction to OpenMP

- Introduction
- OpenMP basics
- OpenMP directives, clauses, and library routines

OpenMP

- An API for shared-memory parallel programming.
- MP = multiprocessing
- Designed for systems in which each thread or process can potentially have access to all available memory.
- System is viewed as a collection of cores or CPU's, all of which have access to main memory.

A Shared Memory system



Definition of OpenMP

- Application Program Interface (API) for Shared Memory Parallel Programming
- Directive based approach with library support
- Targets existing applications and widely used languages:
 - Fortran API released October `97
 - C, C++ API released October `98
- Multi-vendor/platform support

Pragmas

- Special preprocessor instructions.
- Typically added to a system to allow behaviors that aren't part of the basic C specification.
- Compilers that don't support the pragmas ignore them

`#pragma`

```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
} /* Hello */

```



```
gcc -g -Wall -fopenmp -o omp_hello omp_hello . c
```

```
./ omp_hello 4
```

running with 4 threads

compiling

Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

possible
outcomes


Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4

Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4

OpenMP Pragmas

- `# pragma omp parallel`
 - Most basic parallel directive.
 - The number of threads that run the following structured block of code is determined by the run-time system.

A process forking and joining two threads



Clause

- Text that modifies a directive.
- The num_threads clause can be added to a parallel directive.
- It allows the programmer to specify the number of threads that should execute the following block.


```
# pragma omp parallel num_threads ( thread_count )
```

Of note...

- There may be system-defined limitations on the number of threads that a program can start.
- The OpenMP standard doesn't guarantee that this will actually start `thread_count` threads.
- Most current systems can start hundreds or even thousands of threads.
- Unless we're trying to start a lot of threads, we will almost always get the desired number of threads.

Some terminology

- In OpenMP parlance the collection of threads executing the parallel block — the original thread and the new threads — is called a team, the original thread is called the master, and the additional threads are called slaves.



In case the compiler doesn't support OpenMP

```
# include <omp.h>
```



```
#ifdef _OPENMP
```


```
# include <omp.h>
```

```
#endif
```

In case the compiler doesn't support OpenMP

```
# ifdef _OPENMP
    int my_rank = omp_get_thread_num ( );
    int thread_count = omp_get_num_threads ( );
# e l s e
    int my_rank = 0;
    int thread_count = 1;
# endif
```


The trapezoidal rule



The approximate area under the curve is found by adding the area of the trapezoids.

Serial algorithm

```
/* Input:  a, b, n */  
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++) {  
    x_i = a + i*h;  
    approx += f(x_i);  
}  
approx = h*approx;
```

A First OpenMP Version


- 1) We identified two types of tasks:
 - a) computation of the areas of individual trapezoids **[different operations in a single iteration]**, and
 - b) adding the areas of trapezoids.
- 2) There is no communication among the tasks in the first collection, but each task in the first collection communicates with task 1b.

A First OpenMP Version

3) We assumed that there would be many more trapezoids than cores.

- So we aggregated tasks by assigning a contiguous block of trapezoids to each thread (and a single thread to each core).

Assignment of trapezoids to threads



Time	Thread 0	Thread 1
0	<code>global_result = 0 to register</code>	<code>finish my_result</code>
1	<code>my_result = 1 to register</code>	<code>global_result = 0 to register</code>
2	<code>add my_result to global_result</code>	<code>my_result = 2 to register</code>
3	<code>store global_result = 1</code>	<code>add my_result to global_result</code>
4		<code>store global_result = 2</code>

Unpredictable results when two (or more) threads attempt to simultaneously execute:

`global_result += my_result ;`



Mutual exclusion

```
# pragma omp critical  
global_result += my_result ;
```




only one thread can execute
the following structured block at a time

How to compile and run OpenMP programs?

- Gcc 4.2 and above supports OpenMP 3.0
 - gcc -fopenmp a.c
 - Try example1.c
- To run: 'a.out'
 - To change the number of threads:
 - setenv OMP_NUM_THREADS 4 (tcsh) or export OMP_NUM_THREADS=4(bash)

OpenMP execution model



- OpenMP uses the **fork-join model** of parallel execution.
 - All OpenMP programs begin with a single **master thread**.
 - The master thread executes sequentially until a **parallel region** is encountered, when it creates a **team of parallel threads** (FORK).
 - When the team threads complete the parallel region, they synchronize and terminate, leaving only the master thread that executes sequentially (JOIN).

OpenMP general code structure

```
#include <omp.h>
```

```
main () {
```

```
    int var1, var2, var3;
```

```
    Serial code
```

```
    . . .
```

```
    /* Beginning of parallel section. Fork a team of threads. Specify  
       variable scoping*/
```

```
    #pragma omp parallel private(var1, var2) shared(var3)
```

```
    {
```

```
        /* Parallel section executed by all threads */
```

```
        . . .
```

```
        /* All threads join master thread and disband*/
```

```
    }
```

```
    Resume serial code
```

```
    . . .
```

```
}
```

Scope of Variables



Scope


- In serial programming, the scope of a variable consists of those parts of a program in which the variable can be used.
- In OpenMP, the scope of a variable refers to the set of threads that can access the variable in a parallel block.

Scope in OpenMP

- A variable that can be accessed by all the threads in the team has **shared** scope.
- A variable that can only be accessed by a single thread has **private** scope.
- The default scope for variables declared before a parallel block is **shared**.

Data model

- Private and shared variables



P = private data space
G = global data space


- Variables in the global data space are accessed by all parallel threads (**shared** variables).
- Variables in a thread's private space can only be accessed by the thread (**private** variables)
- several variations, depending on the initial values and whether the results are copied outside the region.

```

#pragma omp parallel for private( privIndx,
                                privDbl )
for ( i = 0; i < arraySize; i++ ) {
    for ( privIndx = 0; privIndx < 16; privIndx
        ++ ) { privDbl = ( (double) privIndx ) / 16;
        y[i] = sin( exp( cos( - exp( sin(x[i]) ) ) ) ) )
        + cos( privDbl );
    }
}

```

Parallel for loop index is
Private by default.



execution context for "arrayUpdate_II"

Reduction Clause



Reduction Clause

We need this more complex version to add each thread's local calculation to get *global_result*.

```
void Trap(double a, double b, int n, double* global_result_p);
```

Although we'd prefer this.

```
double Trap(double a, double b, int n);
```



```
global_result = Trap(a, b, n);
```

Reduction Clause

If we use this, there's no critical section!

```
double Local_trap(double a, double b, int n);
```

If we fix it like this...

```
global_result = 0.0;
# pragma omp parallel num_threads(thread_count)
{
#   pragma omp critical
    global_result += Local_trap(double a, double b, int n);
}
```

... we force the threads to execute sequentially.

Reduction Clause

We can avoid this problem by declaring a private variable inside the parallel block and moving the critical section after the function call.

```
global_result = 0.0;
# pragma omp parallel num_threads(thread_count)
{
    double my_result = 0.0;  /* private */

    my_result += Local_trap(double a, double b, int n);
# pragma omp critical
    global_result += my_result;
}
```

Reduction Operators

- ⌘ A **reduction operator** is a binary operation (such as addition or multiplication).
- ⌘ A **reduction** is a computation that repeatedly applies the same reduction operator to a sequence of operands in order to get a single result.
- ⌘ All of the intermediate results of the operation should be stored in the same variable: the reduction variable.

Reduction Operators

A reduction clause can be added to a parallel directive.

```
reduction(<operator>: <variable list>)
```



`+, *, -, &, |, ^, &&, ||`

```
global_result = 0.0;  
# pragma omp parallel num_threads(thread_count) \  
    reduction(+: global_result)  
global_result += Local_trap(double a, double b, int n);
```

OpenMP directives

- Format:

`#pragma omp directive-name [clause,...] newline`
(use ‘\’ for multiple lines)

- Example:

`#pragma omp parallel default(shared) private(beta,pi)`

- Scope of a directive is one block of statements
`{ ... }`




THE “PARALLEL FOR” DIRECTIVE

Parallel for

- Forks a team of threads to execute the following **structured block**.
- However, the structured block following the parallel for directive must be a **for loop**.
- Furthermore, with the **parallel for directive** the system parallelizes the for loop by **dividing the iterations of the loop** among the threads.

Parallel for

```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```



```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+: approx)  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```

Legal forms for parallelizable for statements

for $\left(\begin{array}{l} \text{index} = \text{start} ; \text{index} < \text{end} ; \text{index} += \text{incr} \\ \text{index} = \text{start} ; \text{index} <= \text{end} ; \text{index} += \text{incr} \\ \text{index} = \text{start} ; \text{index} >= \text{end} ; \text{index} -= \text{incr} \\ \text{index} = \text{start} ; \text{index} > \text{end} ; \text{index} -= \text{incr} \end{array} \right)$

Caveats

- The variable **index** must have integer or pointer type (e.g., it can't be a float).
- The expressions **start**, **end**, and **incr** must have a compatible type. For example, if **index** is a pointer, then **incr** must have integer type.

Caveats

- The expressions **start**, **end**, and **incr** must not change during execution of the loop.
- During execution of the loop, the variable **index** can only be modified by the —increment expression|| in the **for** statement.

Data dependencies

```
fibonacci[0] = fibonacci[1] = 1;  
for (i = 2; i < n; i++)  
    fibonacci[i] = fibonacci[i-1] + fibonacci[i-2];
```

note 2 threads

```
fibonacci[0] = fibonacci[1] = 1;  
# pragma omp parallel for num_threads(2)  
for (i = 2; i < n; i++)  
    fibonacci[i] = fibonacci[i-1] + fibonacci[i-2];
```

1 1 2 3 5 8 13 21 34 55

this is correct

but sometimes
we get this

1 1 2 3 5 8 0 0 0 0

What happened?

1. OpenMP compilers **don't check for dependences** among iterations in a loop that's being parallelized with a parallel for directive.
2. A loop in which the results of one or more iterations depend on other iterations **cannot, in general, be correctly parallelized by OpenMP.**

Estimating π (For Eg:)

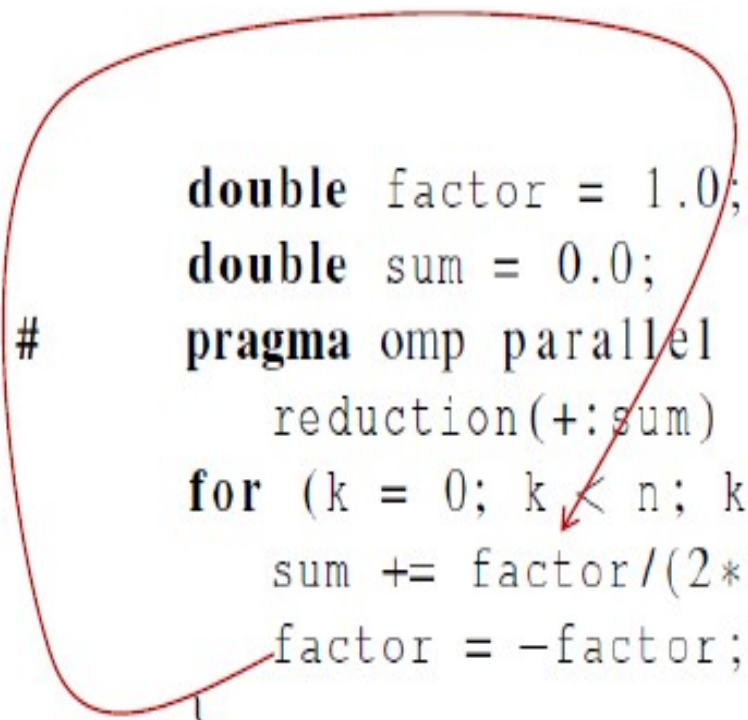
$$\pi = 4 \left[1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

```
double factor = 1.0;  
double sum = 0.0;  
for (k = 0; k < n; k++) {  
    sum += factor / (2 * k + 1);  
    factor = -factor;  
}  
pi_approx = 4.0 * sum;
```

OpenMP solution #1


loop dependency

```
# double factor = 1.0;
double sum = 0.0;
#pragma omp parallel for num_threads(thread_count) \
    reduction(+:sum)
    for (k = 0; k < n; k++) {
        sum += factor/(2*k+1);
        factor = -factor;
    }
pi_approx = 4.0*sum;
```



OpenMP solution #2

```
double sum = 0.0;
# pragma omp parallel for num_threads(thread_count) \
    reduction(+:sum) private(factor)
for (k = 0; k < n; k++) {
    if (k % 2 == 0)
        factor = 1.0;
    else
        factor = -1.0;
    sum += factor/(2*k+1);
}
```



Insures factor has private scope.

The default clause

Lets the programmer specify the scope of each variable in a block.

default (none)

- With this clause the compiler will require that we specify the scope of each variable we use in the block and that has been declared outside the block.

The default clause

```
double sum = 0.0;
# pragma omp parallel for num_threads(thread_count) \
    default(none) reduction(+:sum) private(k, factor) \
    shared(n)
for (k = 0; k < n; k++) {
    if (k % 2 == 0)
        factor = 1.0;
    else
        factor = -1.0;
    sum += factor/(2*k+1);
}
```

Forking Pthreads

Signature:

```
int pthread_create(pthread_t *,
                  const pthread_attr_t *,
                  void * (*)(void *),
                  void *);
```

Example call:

```
errcode = pthread_create(&thread_id,
                        &thread_attribute,
                        &thread_fun, &fun_arg);
```

- `thread_id` is the thread id or handle (used to halt, etc.)
- `thread_attribute` various attributes
 - standard default values obtained by passing a NULL pointer
- `thread_fun` the function to be run (takes and returns void*)
- `fun_arg` an argument can be passed to `thread_fun` when it starts
- `errorcode` will be set to nonzero if the create operation fails

Forking Pthreads, cont.

The effect of `pthread_create`

- Master thread actually causes the operating system to create a new thread
- Each thread executes a specific function, `thread_fun`
 - The same thread function is executed by all threads that are created, representing the thread's *computation decomposition*
- For the program to perform different work in different threads, the arguments passed at thread creation distinguish the thread's “id” and any other unique features of the thread.

Simple Threading Example

```
int main() {  
    pthread_t threads[16];  
    int tn;  
    for(tn=0; tn<16; tn++) {  
        pthread_create(&threads[tn], NULL, ParFun, NULL);  
    }  
    for(tn=0; tn<16 ; tn++) {  
        pthread_join(threads[tn], NULL);  
    }  
    return 0;  
}
```

Compile using gcc ... -lpthread

This code creates 16 threads that execute the function “ParFun”.

Note that thread creation is costly, so it is important that ParFun do a lot of work in parallel to amortize this cost.

Slide source: Jim Demmel and Kathy Yelick



MORE ABOUT LOOPS IN OPENMP: SORTING

Bubble Sort

```
for (list_length = n; list_length >= 2; list_length--)  
  for (i = 0; i < list_length-1; i++)  
    if (a[i] > a[i+1]) {  
      tmp = a[i];  
      a[i] = a[i+1];  
      a[i+1] = tmp;  
    }
```



Serial Odd-Even Transposition Sort

```
for (phase = 0; phase < n; phase++)  
    if (phase % 2 == 0)  
        for (i = 1; i < n; i += 2)  
            if (a[i-1] > a[i]) Swap(&a[i-1], &a[i]);  
    else  
        for (i = 1; i < n-1; i += 2)  
            if (a[i] > a[i+1]) Swap(&a[i], &a[i+1]);
```

Serial Odd-Even Transposition Sort

Phase	Subscript in Array			
	0	1	2	3
0	9	\leftrightarrow 7	8	\leftrightarrow 6
	7	9	6	8
1	7	9	\leftrightarrow 6	8
	7	6	9	8
2	7	\leftrightarrow 6	9	\leftrightarrow 8
	6	7	8	9
3	6	7	\leftrightarrow 8	9
	6	7	8	9

First OpenMP Odd-Even Sort

```

for (phase = 0; phase < n; phase++) {
    if (phase % 2 == 0)
#       pragma omp parallel for num_threads(thread_count) \
        default(none) shared(a, n) private(i, tmp)
        for (i = 1; i < n; i += 2) {
            if (a[i-1] > a[i]) {
                tmp = a[i-1];
                a[i-1] = a[i];
                a[i] = tmp;
            }
        }
    else
#       pragma omp parallel for num_threads(thread_count) \
        default(none) shared(a, n) private(i, tmp)
        for (i = 1; i < n-1; i += 2) {
            if (a[i] > a[i+1]) {
                tmp = a[i+1];
                a[i+1] = a[i];
                a[i] = tmp;
            }
        }
}
}

```

Second OpenMP Odd-Even Sort

```

# pragma omp parallel num_threads(thread_count) \
    default(none) shared(a, n) private(i, tmp, phase)
for (phase = 0; phase < n; phase++) {
    if (phase % 2 == 0)
#        pragma omp for
        for (i = 1; i < n; i += 2) {
            if (a[i-1] > a[i]) {
                tmp = a[i-1];
                a[i-1] = a[i];
                a[i] = tmp;
            }
        }
    else
#        pragma omp for
        for (i = 1; i < n-1; i += 2) {
            if (a[i] > a[i+1]) {
                tmp = a[i+1];
                a[i+1] = a[i];
                a[i] = tmp;
            }
        }
}

```

Odd-even sort with two parallel **for** directives and two **for** directives.

(Times are in seconds.)


thread_count	1	2	3	4
Two parallel for directives	0.770	0.453	0.358	0.305
Two for directives	0.732	0.376	0.294	0.239



Hello World! (1)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
```

declares the various Pthreads
functions, constants, types, etc.



```
/* Global variable: accessible to all threads */
int thread_count;
```

```
void *Hello(void* rank); /* Thread function */
```

```
int main(int argc, char* argv[]) {
    long thread; /* Use long in case of a 64-bit system */
    pthread_t* thread_handles;
```

```
/* Get number of threads from command line */
thread_count = strtol(argv[1], NULL, 10);
```

```
thread_handles = malloc (thread_count*sizeof(pthread_t));
```


Hello World! (2)

```
for (thread = 0; thread < thread_count; thread++)  
    pthread_create(&thread_handles[thread], NULL,  
        Hello, (void*) thread);  
  
printf("Hello from the main thread\n");  
  
for (thread = 0; thread < thread_count; thread++)  
    pthread_join(thread_handles[thread], NULL);  
  
free(thread_handles);  
return 0;  
} /* main */
```

Hello World! (3)

```
void *Hello(void* rank) {  
    long my_rank = (long) rank;  /* Use long in case of 64-bit system */  
  
    printf("Hello from thread %ld of %d\n", my_rank, thread_count);  
  
    return NULL;  
} /* Hello */
```

References

1. www.cs.fsu.edu/~xyuan/cda5125/lect9_openmp.ppt
2. http://eclass.uoa.gr/modules/document/file.php/D186/%CE%8E%CE%BB%CE%B7%202011-12/PachecoChapter_5.pdf
3. www.cs.utah.edu/~mhall/cs4961f11/CS4961-L5.ppt