

OpenMP Execution Model

Execution Model

- OpenMP API uses the **fork-join model** of parallel execution.
- **Multiple threads of execution** perform tasks defined implicitly or explicitly by openMP directives
- It is intended to support programs that will execute correctly both as parallel programs (**multiple threads of execution and a full openMP support library**) and as sequential programs (**directives ignored and a simple openMP stubs library**).

Execution Model(cont...)

- An OpenMP program begins as a single thread of execution, called an **initial thread**.
- An initial thread executes sequentially, as if enclosed in an implicit task region, called an **initial task region** that is defined by an implicit **inactive parallel region** that surrounds the **entire target region**.
- When a target construct is encountered, **the target region is executed by the implicit device task**.
- The task that encounters the target construct **waits** at the end of the construct until execution of the region completes.
 - If a target device doesn't exist, or the target device is not supported by the implementation, or the target device cannot execute the target construct then the **target region is executed by the host device**.

Execution Model(cont...)

- The teams construct creates a *league of thread teams* where the **masters thread** of each team executes the region.
 - Each of these master threads is an initial thread, and executes sequentially, as if enclosed in an implicit task region that is defined by an **implicit parallel region** that surrounds the entire teams region.
- When any thread encounters a parallel construct, the thread creates a team of itself and zero or more additional threads and becomes the master of the new team.
 - A **set of implicit tasks**, one per thread, is generated.
 - The code for each task is defined by the code **inside the parallel construct**.

Execution Model(cont...)

- Each task is assigned to a different thread in the team and becomes tied, ie., it is always executed by the thread to which it is initially assigned.
 - The task region of the task being executed by the encountering thread is suspended, and each member of the new team executes its implicit task.
 - There is an **implicit barrier** at the end of the parallel construct.
 - Only the master thread resumes execution beyond the end of the parallel construct, resuming the task region that was suspended upon encountering the parallel construct.
 - Any no. of parallel constructs can be specified in a single program.

Execution Model : nested parallelism

- Parallel regions may be **arbitrarily nested inside each other**.
- If **nested parallelism is disabled**, or is not supported by the OpenMP implementation, then the **new team** that is created by a thread encountering a parallel construct inside a parallel region **will consist only of the encountering thread**. (ie. Encountering thread alone will be present there and executes the implementation).
- However, if **nested parallelism is supported and enabled**, then the new team can consist of more than one thread.
 - A parallel construct may include a **proc_bind** clause to specify the places to use for the threads in the team within the parallel region.

Execution Model : Work Sharing Constructs

- When any team encounters a work sharing construct, the work inside the **construct is divided among the members of the team**, and executed cooperatively instead of being executed by every thread.
- There is a **default barrier at the end** of each work sharing construct
- **Redundant execution of code** by every thread in the team resumes after the end of the work sharing construct.
- When any thread encounters a task construct, **a new explicit task is generated**.
- Execution of explicitly generated tasks is **assigned to one of the threads in the current team**, subject to the thread's availability to execute work.

Execution Model : Work Sharing Constructs (cont...)

- Thus, execution of the new task could be **immediate**, or **deferred** until later according to task scheduling constraints and thread availability.
- Threads are allowed to **suspend the current task region** at a task scheduling point in order **to execute a different task**.
- If the suspended task region is for **tied task**, the initially assigned thread resumes execution of the suspended task region.
- If the suspended task region is for **untied task**, then any thread may resume its execution.

Execution Model : Work Sharing Constructs (cont...)

- Completion of **all explicit tasks** bound to a given parallel region is guaranteed **before the master thread leaves the implicit barrier** at the end of the region.
- Completion of a **subset of all explicit tasks** bound to a given parallel region may be **specified through the use of task synchronization constructs**.
- Completion of **all explicit tasks** bound to a implicit parallel region is guaranteed by the time the program exits.

Execution Model : construct-type clause

- When **any thread encounters a SIMD construct**, the iterations of the loop associated with the construct may be **executed concurrently using the SIMD lanes** that are available to the thread.
- The **cancel construct** can alter the previously described flow of execution in an OpenMP region.
- Effect of the cancel construct depends on its **construct-type clause**.
- If a task encounters a cancel construct with a task group construct-type clause, then the **task activates cancellation and continues execution at the end of its task region**, which implies completion of that task.

Parallel region construct

- A block of code that will be executed by multiple threads.

```
#pragma omp parallel [clause ...]  
{  
    .....  
} (implied barrier)
```

Clauses: *if (expression)*, *private (list)*, *shared (list)*, *default (shared | none)*, *reduction (operator: list)*, *firstprivate(list)*, *lastprivate(list)*

- *if (expression)*: only in parallel if expression evaluates to true
- *private(list)*: everything private and local (no relation with variables outside the block).
- *shared(list)*: data accessed by all threads
- *default (none|shared)*

- The reduction clause:

```
Sum = 0.0;
```

```
#pragma parallel default(none) shared (n, x) private (I) reduction(+ : sum)  
{  
    For(I=0; I<n; I++) sum = sum + x(I);  
}
```

- Updating sum must avoid racing condition
 - With the reduction clause, OpenMP generates code such that the race condition is avoided.
- Firstprivate(list): variables are initialized with the value before entering the block
- Lastprivate(list): variables are updated going out of the block.

Work-sharing constructs


- `#pragma omp for [clause ...]`
- `#pragma omp section [clause ...]`
- `#pragma omp single [clause ...]`
- The work is distributed over the threads
- Must be enclosed in parallel region
- No implied barrier on entry, implied barrier on exit (unless specified otherwise)

The omp for directive: example

```
#pragma omp parallel default(none) \  
    shared(n,a,b,c,d) private(i)  
{  
    #pragma omp for nowait  
    for (i=0; i<n-1; i++)  
        b[i] = (a[i] + a[i+1])/2;  
  
    #pragma omp for nowait  
    for (i=0; i<n; i++)  
        d[i] = 1.0/c[i];  
  
} /*-- End of parallel region --*/  
    (implied barrier)
```

- Schedule clause (decide how the iterations are executed in parallel):

schedule (static | dynamic | guided [, chunk])



The omp session clause - example

```
#pragma omp parallel default(none) \  
    shared(n,a,b,c,d) private(i)  
{  
    #pragma omp sections nowait  
    {  
        #pragma omp section  
        for (i=0; i<n-1; i++)  
            b[i] = (a[i] + a[i+1])/2;  
  
        #pragma omp section  
        for (i=0; i<n; i++)  
            d[i] = 1.0/c[i];  
  
    } /*-- End of sections --*/  
  
} /*-- End of parallel region --*/
```


```
#pragma omp parallel  
#pragma omp for  
    for (...)
```



```
#pragma omp parallel for  
for (...)
```

Single PARALLEL loop

```
#pragma omp parallel  
#pragma omp sections  
{ ... }
```



```
#pragma omp parallel sections  
{ ... }
```

Single PARALLEL sections

Synchronization: barrier

```
For(I=0; I<N; I++)  
    a[I] = b[I] + c[I];
```

```
For(I=0; I<N; I++)  
    d[I] = a[I] + b[I]
```

Both loops are in parallel region
With no synchronization in between.
What is the problem?

Fix:

```
For(I=0; I<N; I++)  
    a[I] = b[I] + c[I];  
  
#pragma omp barrier  
  
For(I=0; I<N; I++)  
    d[I] = a[I] + b[I]
```

Critical session

```
For(I=0; I<N; I++) {  
    .....  
    sum += A[I];  
    .....  
}
```

Cannot be parallelized if sum is shared.

Fix:

```
For(I=0; I<N; I++) {  
    .....  
    #pragma omp critical  
    {  
        sum += A[I];  
    }  
    .....  
}
```

OpenMP environment variables

- OMP_NUM_THREADS
- OMP_SCHEDULE

OpenMP runtime environment

- `omp_get_num_threads`
- `omp_get_thread_num`
- `omp_in_parallel`
- Routines related to locks
-

OpenMP example

- See pi.c

Sequential Matrix Multiply

```
For (I=0; I<n; I++)
```

```
    for (j=0; j<n; j++)
```

```
        c[I][j] = 0;
```

```
        for (k=0; k<n; k++)
```

```
            c[I][j] = c[I][j] + a[I][k] * b[k][j];
```


OpenMP Matrix Multiply

```
#pragma omp parallel for private(j, k)
For (I=0; I<n; I++)
    for (j=0; j<n; j++)
        c[I][j] = 0;
        for (k=0; k<n; k++)
            c[I][j] = c[I][j] + a[I][k] * b[k][j];
```

Travelling Salesman Problem(TSP)

- The map is represented as a graph with nodes representing cities and edges representing the distances between cities.
- A special node (cities) is the starting point of the tour.
- Travelling salesman problem is to find the circle (starting point) that covers all nodes with the smallest distance.
- This is a well known NP-complete problem.

Sequential TSP

```
Init_q(); init_best();  
While ((p = dequeue()) != NULL) {  
    for each expansion by one city {  
        q = addcity (p);  
        if (complete(q)) {update_best(q);}   
        else enqueue(q);  
    }  
}
```

OpenMP TSP

```
Do_work() {  
  While ((p = dequeue()) != NULL) {  
    for each expansion by one city {  
      q = addcity (p);  
      if (complete(q)) {update_best(q);}  
      else enqueue(q);  
    }  
  }  
}
```

```
main() {  
  init_q(); init_best();  
  #pragma omp parallel for  
  for (i=0; i < NPROCS; i++)  
    do_work();  
}
```

Sequential SOR

```
for some number of timesteps/iterations {  
    for (i=0; i<n; i++)  
        for( j=1, j<n, j++ )  
            temp[i][j] = 0.25 *  
                ( grid[i-1][j] + grid[i+1][j]  
                  grid[i][j-1] + grid[i][j+1] );  
    for( i=0; i<n; i++)  
        for( j=1; j<n; j++ )  
            grid[i][j] = temp[i][j];  
}
```

- OpenMP version?

- Summary:
 - OpenMP provides a compact, yet powerful programming model for shared memory programming
 - It is very easy to use OpenMP to create parallel programs.
 - OpenMP preserves the sequential version of the program
 - Developing an OpenMP program:
 - Start from a sequential program
 - Identify the code segment that takes most of the time.
 - Determine whether the important loops can be parallelized
 - The loops may have critical sections, reduction variables, etc
 - Determine the shared and private variables.
 - Add directives

OpenMP discussion

- Ease of use
 - OpenMP takes care of the thread maintenance.
 - Big improvement over pthread.
 - Synchronization
 - Much higher constructs (critical section, barrier).
 - Big improvement over pthread.
- OpenMP is easy to use!!

OpenMP discussion

- Expressiveness
 - Data parallelism:
 - MM and SOR
 - Fits nicely in the paradigm
 - Task parallelism:
 - TSP
 - Somewhat awkward. Use OpenMP constructs to create threads. OpenMP is not much different from pthread.

OpenMP discussion

- Exposing architecture features (performance):
 - Not much, similar to the pthread approach
 - Assumption: dividing job into threads = improved performance.
 - How valid is this assumption in reality?
 - Overheads, contentions, synchronizations, etc
 - This is one weak point for OpenMP: the performance of an OpenMP program is somewhat hard to understand.

OpenMP final thoughts

- Main issues with OpenMP: performance
 - Is there any obvious way to solve this?
 - Exposing more architecture features?
 - Is the performance issue more related to the fundamental way that we write parallel program?
 - OpenMP programs begin with sequential programs.
 - May need to find a new way to write efficient parallel programs in order to really solve the problem.

References

1. www.cs.fsu.edu/~xyuan/cda5125/lect9_openmp.ppt
2. http://eclass.uoa.gr/modules/document/file.php/D186/%CE%8E%CE%BB%CE%B7%202011-12/PachecoChapter_5.pdf
3. www.cs.utah.edu/~mhall/cs4961f11/CS4961-L5.ppt
4. sddsd

