

MPI Program Execution

Sample Programs - Explanations

MPI Program Execution

- How to program distributed-memory systems using message-passing.
- In message passing programs, a program running on one core-memory pair is usually called a **process**, and 2 processes can communicate by calling functions:
 - One process calls a send function and
 - the other calls a receive function
- The implementation of message-passing is called Message-Passing Interface(MPI)

MPI Program Execution: Simple **MPI** Program

/*

Setting the path as :

export PATH="\$PATH:/home/
\$USER/.openmpi/bin"

export
LD_LIBRARY_PATH="\$LD_LIBRARY_P
ATH:/home/\$USER/.openmpi/lib/"

\$ cd MPI_Prog/

*/

```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
main(int argc, char **argv)
```

```
{
```

```
    int ierr;
```

```
    ierr = MPI_Init(&argc, &argv);
```

```
    printf("Hello world\n");
```

```
    ierr = MPI_Finalize();
```

```
}
```

MPI Program Execution

- Output:

```
$ mpicc hello.c -o hello
```

```
$ mpirun -np 2 ./hello
```

```
Hello world
```

```
Hello world
```

MPI Program Execution: Simple C program

```
#include <stdio.h>

int main(void)
{
    printf("hello, world");
    return 0;
}
```

- Lets write a program that to “hello, world” that makes some use of MPI.
- Instead of having each process simply print a message, we'll designate one process to do the output, and the other processes will send it messages, which it will print.

MPI Program Execution

- In parallel programming, it's common for the processes to be identified by nonnegative integer ranks.
- So if there are p processes, the processes will have ranks $0, 1, 2, \dots, p-1$.
- For our parallel “hello, world”, let's make process 0 the designated process, and the other processes will send it messages.

MPI Program Execution

```
#include <mpi.h>
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#include <string.h>
```

```
int main(int argc, char** argv)
```

```
{
```

```
    int q;
```

```
    // Initialize the MPI environment
```

```
    MPI_Init(NULL, NULL);
```

```
    // Find out rank, size
```

```
    int world_rank; // number the processes
```

```
    MPI_Comm_rank(MPI_COMM_WORLD,  
&world_rank);
```

```
    int world_size; // Tot number the threads
```

```
    MPI_Comm_size(MPI_COMM_WORLD,  
&world_size);
```

```
    // We are assuming at least 2 processes for  
    this task
```

```
    if (world_size < 2) {
```

```
        fprintf(stderr, "World size must be greater than  
1 for %s\n", argv[0]);
```

```
        MPI_Abort(MPI_COMM_WORLD, 1);
```

```
    }
```

MPI Program Execution

```
char greeting[50];  
  
if (world_rank != 0)  
{  
    // If we are rank != 0, send the greeting to  
    // process id !=0  
  
    sprintf(greeting,"Greeting... from process %d  
of %d!",world_rank,world_size);  
  
    MPI_Send(greeting, strlen(greeting)+1,  
MPI_CHAR, 0, 0, MPI_COMM_WORLD);  
}
```

```
else  
{  
    printf("Greetings from process %d of  
%d!\n",world_rank,world_size);  
    for(q=1;q<world_size;q++)  
    {  
        MPI_Recv(greeting, 50, MPI_CHAR, q, 0,  
MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
        printf("%s\n",greeting);  
    }  
} // end of else  
  
MPI_Finalize();  
return 0;  
} // end of main()
```


MPI Program Execution

```
char greeting[50];

if (world_rank != 0)
{
    // If we are rank != 0, send the greeting to
    // process id !=0

    sprintf(greeting,"Greeting... from process %d
of %d!",world_rank,world_size);

    MPI_Send(greeting, strlen(greeting)+1,
MPI_CHAR, 0, 0, MPI_COMM_WORLD);
}
```

Dest:
rank of
the
process
that
should
recv the
msg

tag:
nonnegative
int.
0=> messages
to be printed
1=> messages
to be used in a
computation

communicator:
collection of
processes that
can send
messages to
each other.
If 2 processes are
using diff
communicator,
msg cant be sent
and recvd

```
else
{
    printf("Greetings from process %d of
%d!\n",world_rank,world_size);

    for(q=1;q<world_size;q++)
    {
        MPI_Recv(greeting, 50, MPI_CHAR, q, 0,
MPI_COMM_WORLD, MPI_STATUS_IGNORE);

        printf("%s\n",greeting);
    }
} // end of else

MPI_Finalize();

return 0;
} // end of main()
```

MPI Program Execution

```
char greeting[50];

if (world_rank != 0)
{
    // If we are rank != 0, send the greeting to
    // process id !=0

    sprintf(greeting,"Greeting... from process %d
of %d!",world_rank,world_size);

    MPI_Send(greeting, strlen(greeting)+1,
MPI_CHAR, 0, 0, MPI_COMM_WORLD);
}
```

```
else
{
    printf("Greetings from process %d of
%d!\n",world_rank,world_size);
    for(q=1;q<world_size;q++)
    {
        MPI_Recv(greeting, 50, MPI_CHAR, q, 0,
MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("%s\n",greeting);
    }
} // end of else

MPI_Finalize();
return 0;

} // end of main()
```

Source: specifies the process from which the message should be received

tag: match the tag arg of the msg being sent

MPI Program Execution

- Output:

```
$ mpicc send_recv.c -o send_recv
```

```
$ mpirun -np 2 ./send_recv
```

Greetings from process 0 of 2!

Greeting... from process 1 of 2!

or

```
$ mpirun -np 4 ./send_recv
```

Greetings from process 0 of 4!

Greeting... from process 1 of 4!

Greeting... from process 2 of 4!

Greeting... from process 3 of 4!

- `MPI_Comm_rank` and `MPI_Comm_size` are first used to determine the world size along with the rank of the process.
- In the else statement, process zero initializes greeting message with process id 0
- As you can see in the if statement, process ranks(which are !=0) is calling `MPI_Send` to send the greeting. Message received at the `MPI_Recv` also