MapR

THE CONVERGED DATA PLATFORM

What is MapR?

- **One Platform for Big Data Applications**
- The MapR Converged Data Platform enables direct processing of files, tables, and event streams.
- The MapR Platform has the ability to leverage multiple data types and formats directly, depending on your use case.

MapR Components

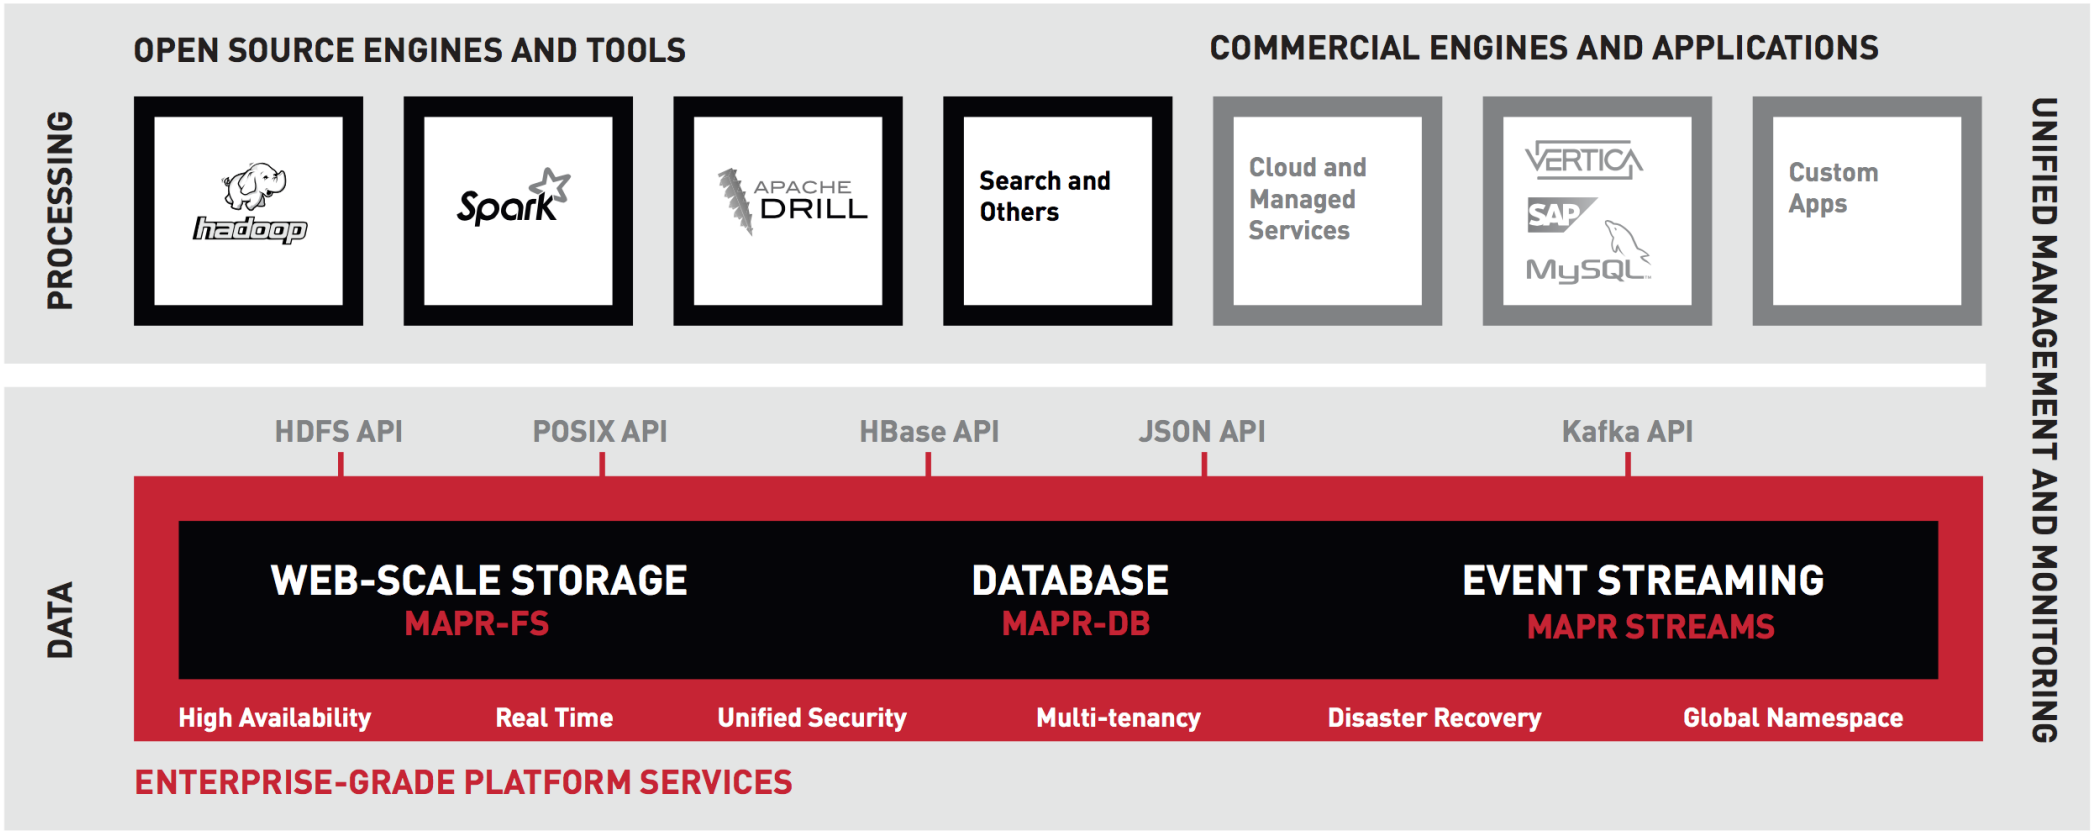

MapR-FS

- **MapR-FS provides high performance enterprise-grade storage for big data**
- MapR-FS supports the HDFS API, fast NFS access, access controls MapR ACEs, and transparent data compression.
- MapR-FS includes enterprise-grade features such as block-level mirroring for mission-critical disaster recovery as well as load balancing, and consistent snapshots for easy data recovery.

Key Features

- With MapR-FS, you get a complete POSIX file system that handles raw disk I/O for big data workloads. You can run it as an enterprise-wide NAS to run existing file system-based applications and systems. It supports:
- Converged Data Platform
- Open Source Engines
- MapR-FS

Key Features

- System-wide high availability
- Mission-critical disaster recovery
- Consistent snapshots
- Unified security

MapR - DB

- **MapR-DB adds real-time, interactive database workloads to analytical environments**
- MapR-DB is an enterprise-grade, high performance, NoSQL (“Not Only SQL”) database management system.
- Used to add real-time, operational analytics capabilities to big data applications.

Key Features

- Real-Time Operational Analytics:
- You can run MapR-DB in the same cluster as Apache™ Hadoop® and Apache Spark, letting you immediately analyze or process live, interactive data.
- Minimal Administrative Overhead:
- With MapR-DB, you can reduce your administrative overhead by eliminating maintenance tasks that are typical in other NoSQL databases

Key Features

- Fine Grained Security:
- MapR-DB lets you grant access permissions at a granular level using MapR Access Control Expressions (ACES), which are designed for flexibility and ease-of-use. You can set different permission levels for different roles at the sub-document level (for the JSON document model) or at the column level (for the wide column data model).

MapR Streams

- **MapR Streams is a global publish-subscribe event streaming system for big data**.
- It connects data producers and consumers worldwide in real time, with unlimited scale
- Publishers (data producers) write data to one or more topics in MapR Streams. Subscribers (data consumers) to the topic can read the data instantaneously, anywhere across the globe.
- MapR Streams enables MapR Converged Data Platform to integrate streams, storage and analytics in a single cluster

Key Features

- The first big data-scale streaming system to be built into a converged data platform
- The only big data streaming system to support global event replication at Internet-of-Things (IoT) scale and reliability, providing failover endpoints between up to thousands of distributed clusters.