

WS-Coordination and WS-Transactions

Transactions

- Distributed system
 - Reliability problems
 - Subject to independent failure of any of its components
 - Decentralization allows:
 - Parts of the system fail
 - Other parts remain functioning
 - Possibility of abnormal behaviors

ACID Properties

- ***Atomicity:*** The transaction completes successfully (commits) or if it fails (aborts) all of its effects are undone (rolled back)
- ***Consistency:*** Transactions produce consistent results and preserve application specific invariants
- ***Isolation:*** Intermediate states produced while a transaction is executing are not visible to other transactions
 - appear to execute serially
 - achieved by locking resources
- ***Durability:*** The effects of a committed transaction are never lost (except by a catastrophic failure)

Atomic Transactions

- can be terminated in two ways:
 - *Committed*
 - *all changes made within it are made durable*
 - *Aborted* (rolled back)
 - all changes made during the lifetime of the transaction are undone

Web Environment

- Web Service activities form a unit of work, but ACID properties are not always appropriate since Web is loosely coupled and activities are frequently long-running
 - Communication delays
 - Server loading (and hence response time) is unpredictable
 - Locks might not be retained for duration of activity
 - Servers are controlled by different organizations
 - Modules might not trust each other
 - Failures more likely

Why Atomic transactions may be too strong

- Suitable for short lived applications
- Long-lived transactions may reduce the concurrency
 - By holding onto resources for a long time
 - If it aborts
 - Much valuable work already performed will be undone
- Pure ACID transactions are not suitable for Web Services

A Suit of Specifications

- As a response to these needs in July 2002, BEA, IBM, and Microsoft released a trio of specifications designed to support business transactions over Web services:
 - BPEL4WS (Business Process Execution Language for Web Services),
 - WS-C (WS-Coordination) and
 - WS-TX (WS-Transaction)

Coordination

- Refers to the mechanism used by the components of an activity to reach common agreement:
 - How are components identified?
 - How are exceptional situations to be handled?
 - System related failures (e.g., crash, communication failure)
 - Application related exceptions (e.g., unanticipated messages)
 - What constitutes successful termination?

WS-Coordination

- WS-Coordination defines a framework for providing protocols that coordinate the actions of distributed applications
 - Provides a generic framework for coordination protocols to be plugged in
- Provides only context management - it allows contexts to be created and activities to be registered with those contexts.

WS-Coordination

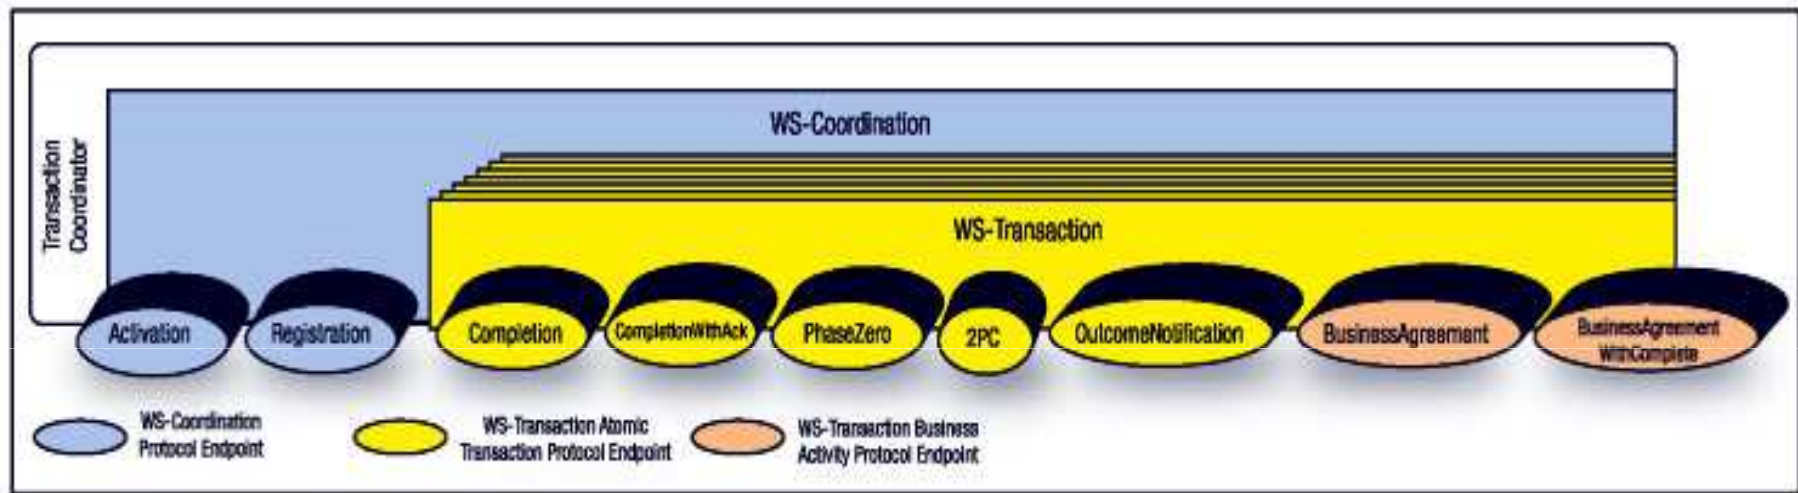


FIGURE 1 | WS-Coordination and WS-Transaction

Overview

- A coordinator supports three services:
 - **Activation Service:** part of **WS-Coordination**
 - Create a new activity with a particular CoordinationType: AtomicTransaction or BusinessActivity
 - Each type supports several termination protocols; a participant can choose the protocol appropriate to its role in the activity
 - **Registration Service:** part of **WS-Coordination**
 - Allow each participant to register for a protocol within the type
 - Different participants in the same activity might use different protocols
 - Applications register with a **coordinator** to create a **coordination context** that is **carried by all applications** within a given activity
 - **Protocol Service:** supports the execution of the protocol as specified in WS-AtomicTransaction or WS-BusinessActivity

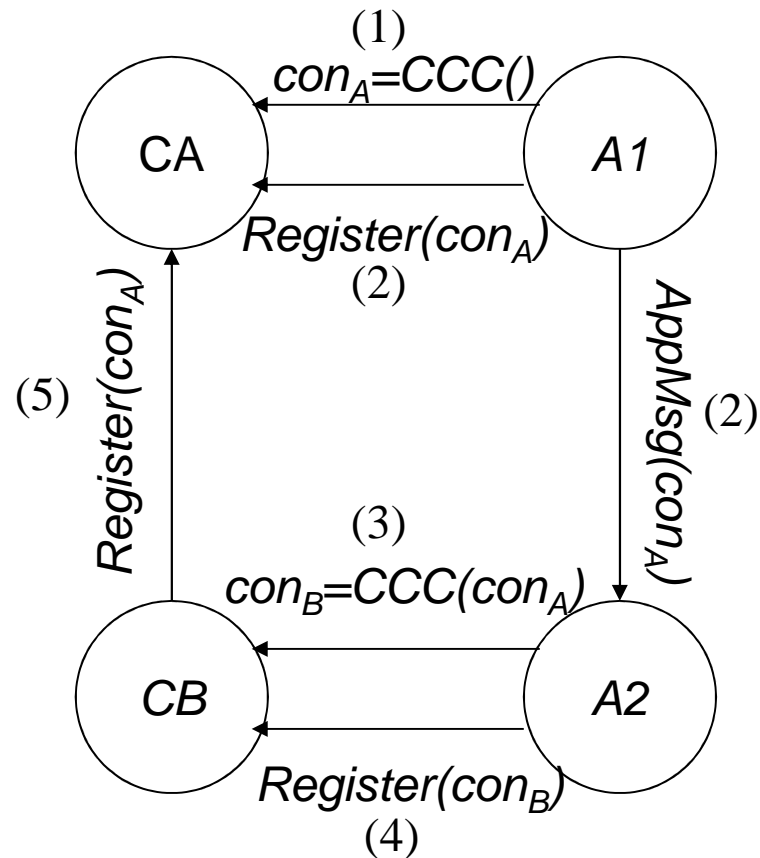
Creating an Activity 1

- *A1* invokes *CreateCoordinationContext()* to get a context, con_A , for a new activity of coordination type Q from coordinator CA
- *A1* sends an application message containing the context to *A2*
 - When SOAP is used, the context is a header block with *mustUnderstand*=“true”
- *A1* and *A2* then use CA 's registration service to register for protocols (perhaps different, but supported by Q)


Creating an Activity 2

- A2 might want to use a different coordinator
 - Reasons: performance, trust
- Protocol:
 - A2 invokes *CreateCoordinationContext()* at coordinator *CB*, passing con_A as a parameter
 - *CB* creates a new context, con_B , with the same identifier and CoordinationType, returns it to A2
 - A1 registers for a protocol with CA
 - A2 registers for a protocol with *CB* and *CB* registers for the same protocol with CA

Creating an Activity



Message exchange



Protocol tree

Registration Service

- Participant explicitly registers for a particular protocol supported by the CoordinationType.
 - Contrasts with two-phase commit for distributed transactions where registration is automatic when server is invoked
 - Participant can register several times in order to participate in activity termination in several ways.

WS-Transaction

- WS-T proposes two distinct models:
 - **Atomic transaction (AT) Model** is used to coordinate activities having a short duration and executed within limited trust domains
 - addresses "fine-grained" transactions
 - **Business Activity (BA) Model** is used to coordinate activities that are long in duration and desire to apply business logic to handle business exceptions
 - addresses "coarse-grained" transactions
 - BA defines a protocol based on compensation used to achieve distributed consensus on whether the results of a long-running message exchange should be made persistent.

Atomic Transactions

- Similar to traditional ACID transactions
- Services enroll transaction-aware resources
 - Databases
 - Message queues

The Players

- Completion Initiator
 - Signals coordinator to complete a transaction
 - Can request commit or rollback
- Coordinator
 - Responsible for coordinating a single outcome
 - Drives 2PC with participants
 - Phase 1: Ensure all participants are prepared
 - Phase 2: Notify participants of outcome
- Participants
 - can vote to abort
 - Can vote “prepared to commit”
 - Must honor coordinator’s commit decision

WS-AtomicTransaction

- A coordination type that implements transactional atomicity using two-phase commit
 - Intended for tightly-coupled, short-lived activities within an organizational structure
 - Systems should trust each other
 - They must be responsive since locks are held until protocol completes
 - A single system can (perhaps maliciously) abort the entire transaction
- Protocols supported
 - **Completion**: initiates termination, final result returned (no durable resources)
 - *commit*, *rollback*, *committed*, and *aborted* messages
 - Two-Phase Commit (with presumed abort)
 - **Volatile 2PC** (no durable resources)
 - **Durable2PC** (durable resources)
 - *prepare*, *vote*, *commit*, *rollback*, *readOnly* (for cohorts that need not participate in phase 2), *aborted*, *committed* (done), and *replay* (for cohorts dealing with failure) messages
- A participant might register for both completion and 2PC

AT: Example

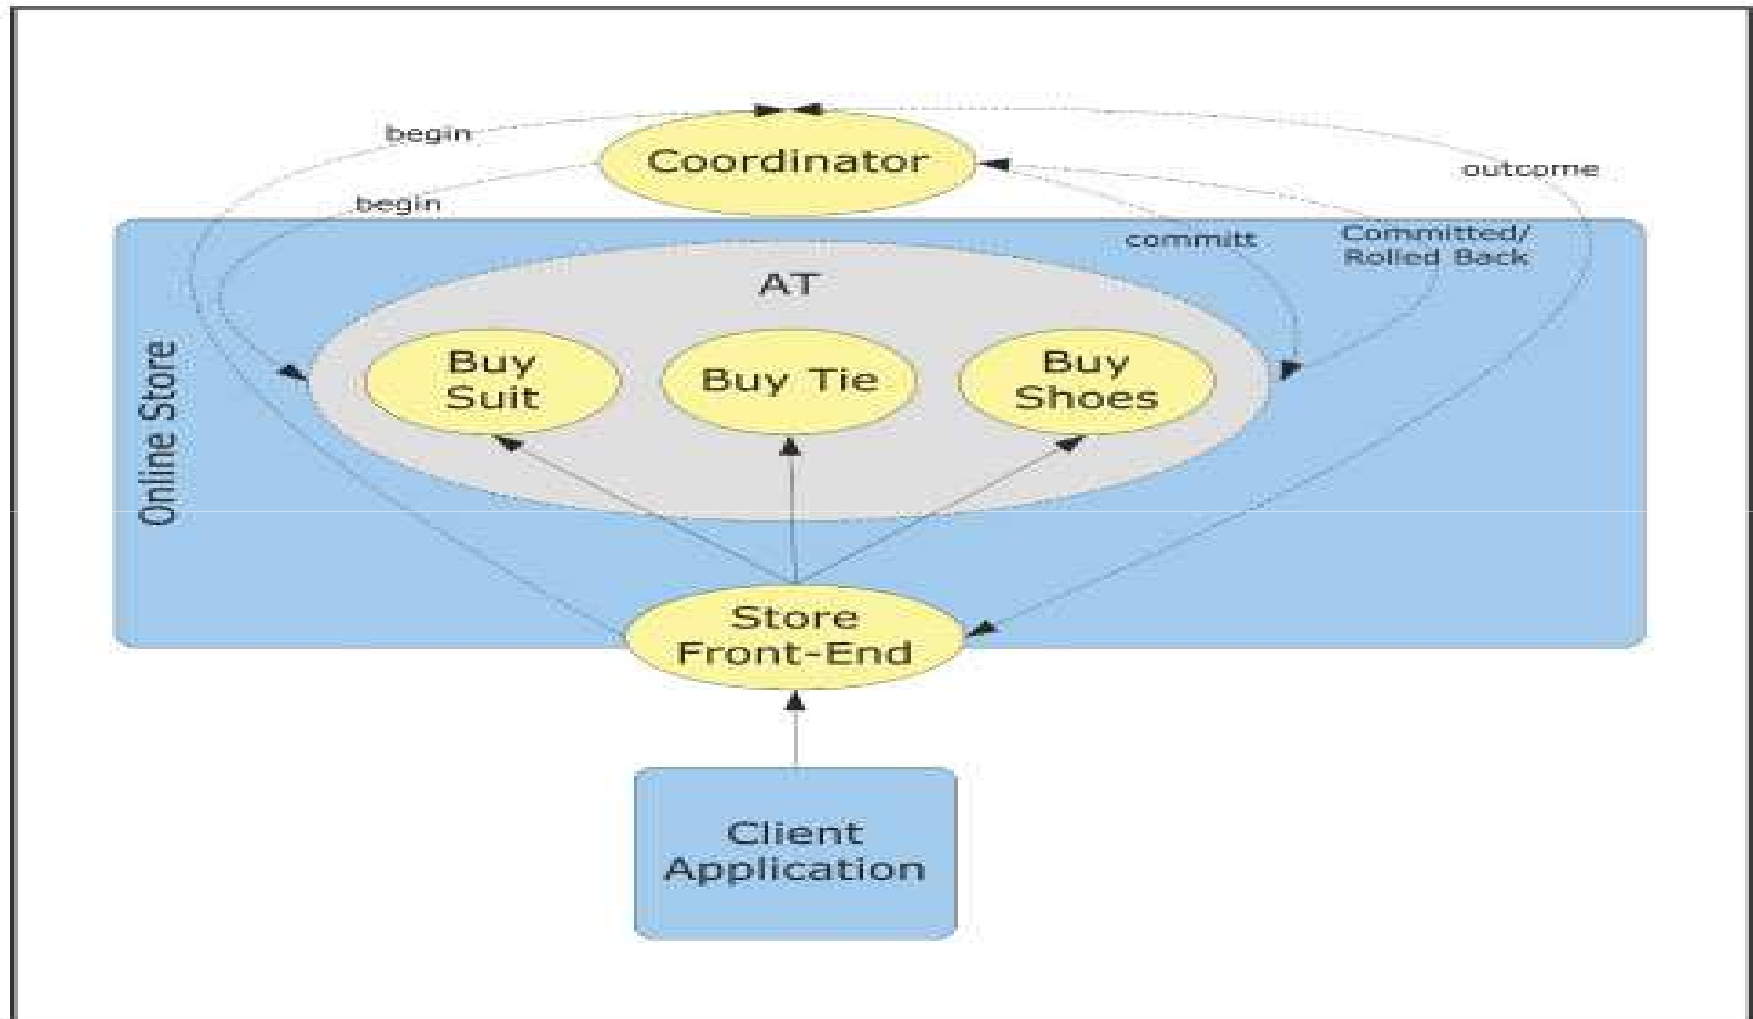


FIGURE 6 | Using an AT to ensure all-or-nothing semantics for buying formalwear

Business Activity

- Short running atomic transactions can be part of a long running business transaction
- The actions of the embedded atomic transactions are committed and made visible before the long running business transaction completes
- In the event of the long running business transaction failing, the effects of such atomic transactions need to be compensated for.

BA: Example

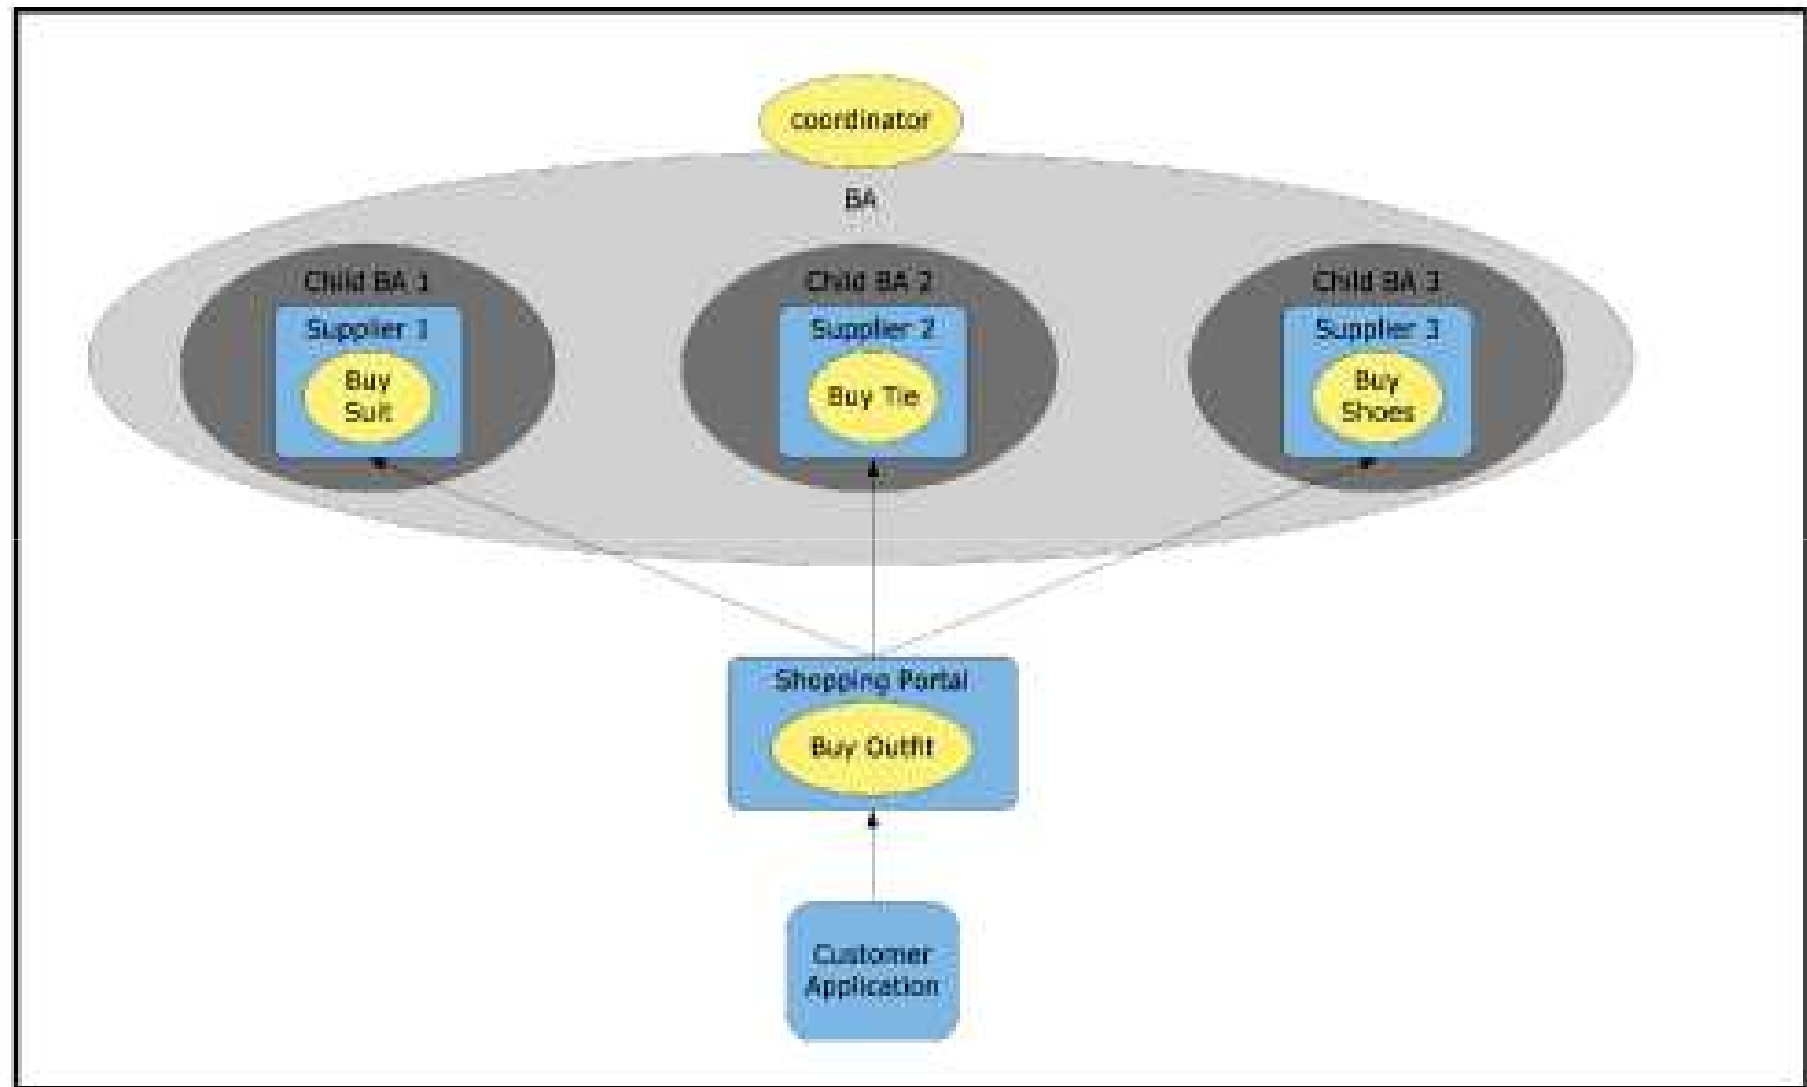


FIGURE 1 | Using a business activity to support compensating transactions across enterprise boundaries

BA: Example

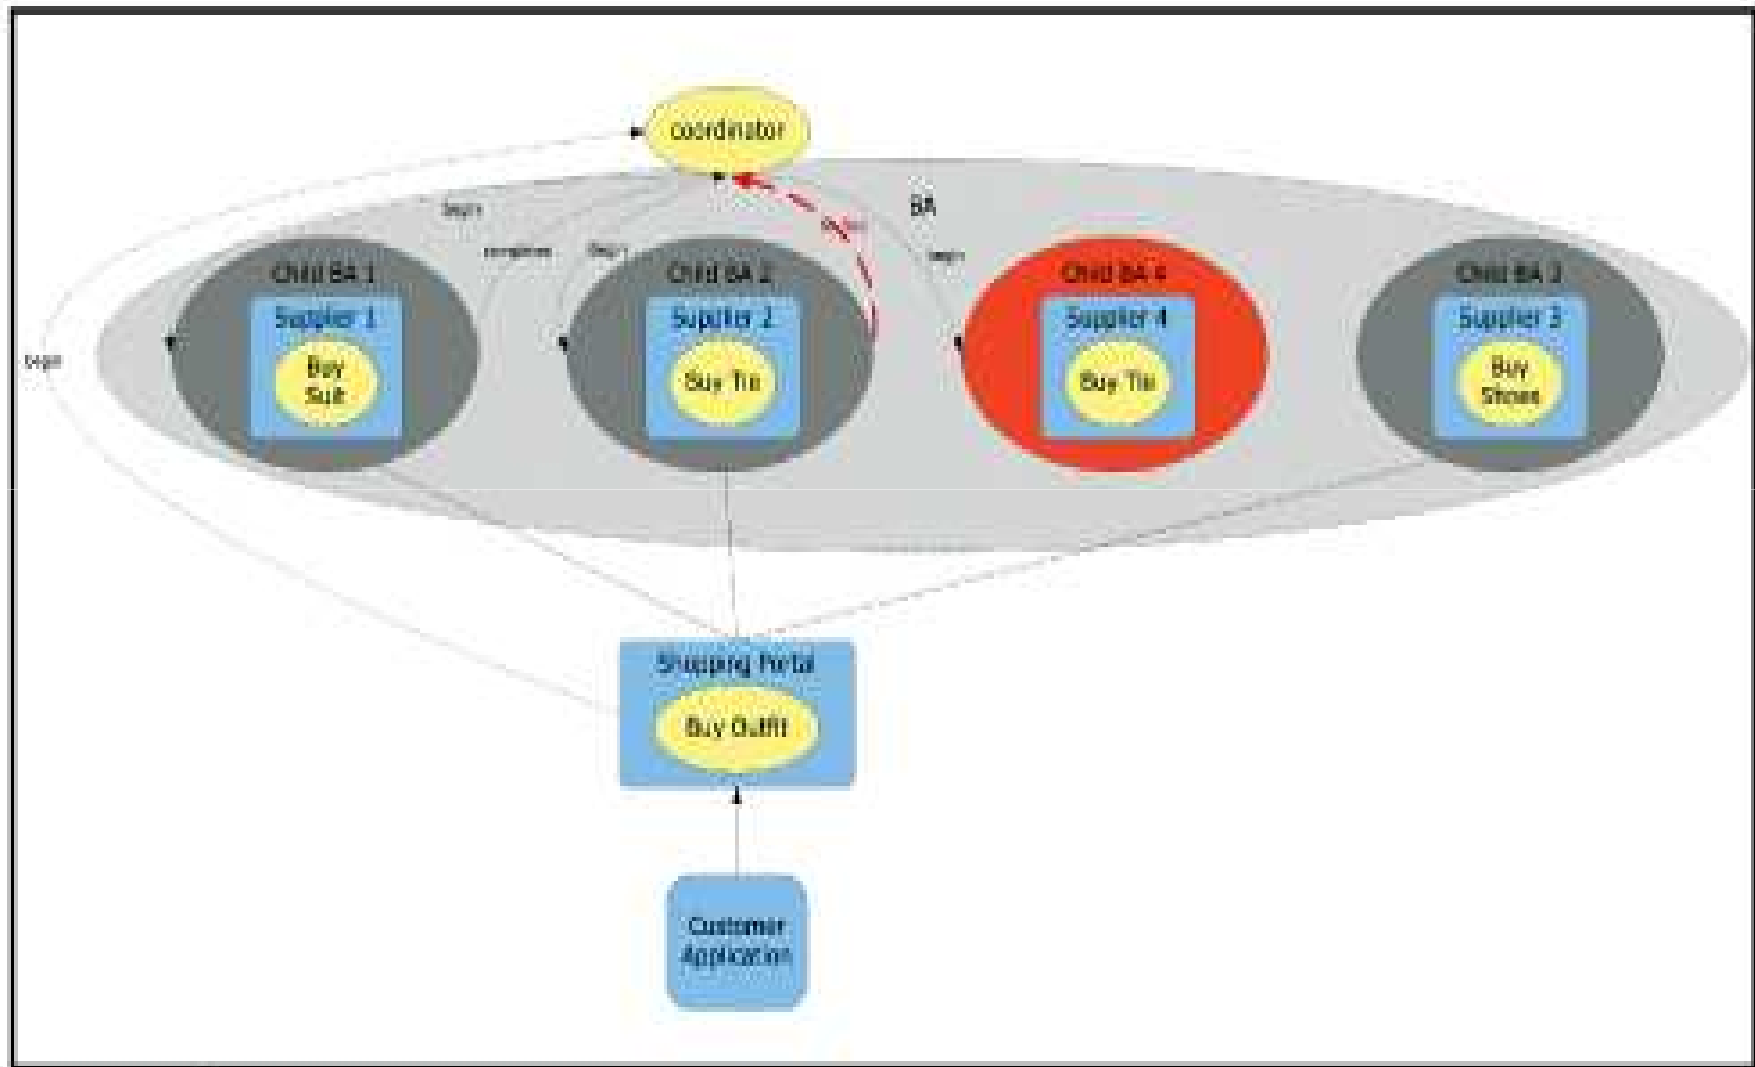


FIGURE 2 | Error handling and forward recovery

WS-BusinessActivity

- WS-BusinessActivity coordination type supports two protocols
 - BusinessAgreementWithParticipantCompletion (BAWPC)
 - Participant registering for this protocol can initiate termination
 - BusinessAgreementWithCoordinatorCompletion (BAWCC)
 - Participant registering for this protocol expects coordinator to tell it when to terminate

WS-BusinessActivity

- Components of same business activity might register for different protocols depending on who is responsible for determining completion
 - Root application might register for BAWPC, other participants for BAWCC

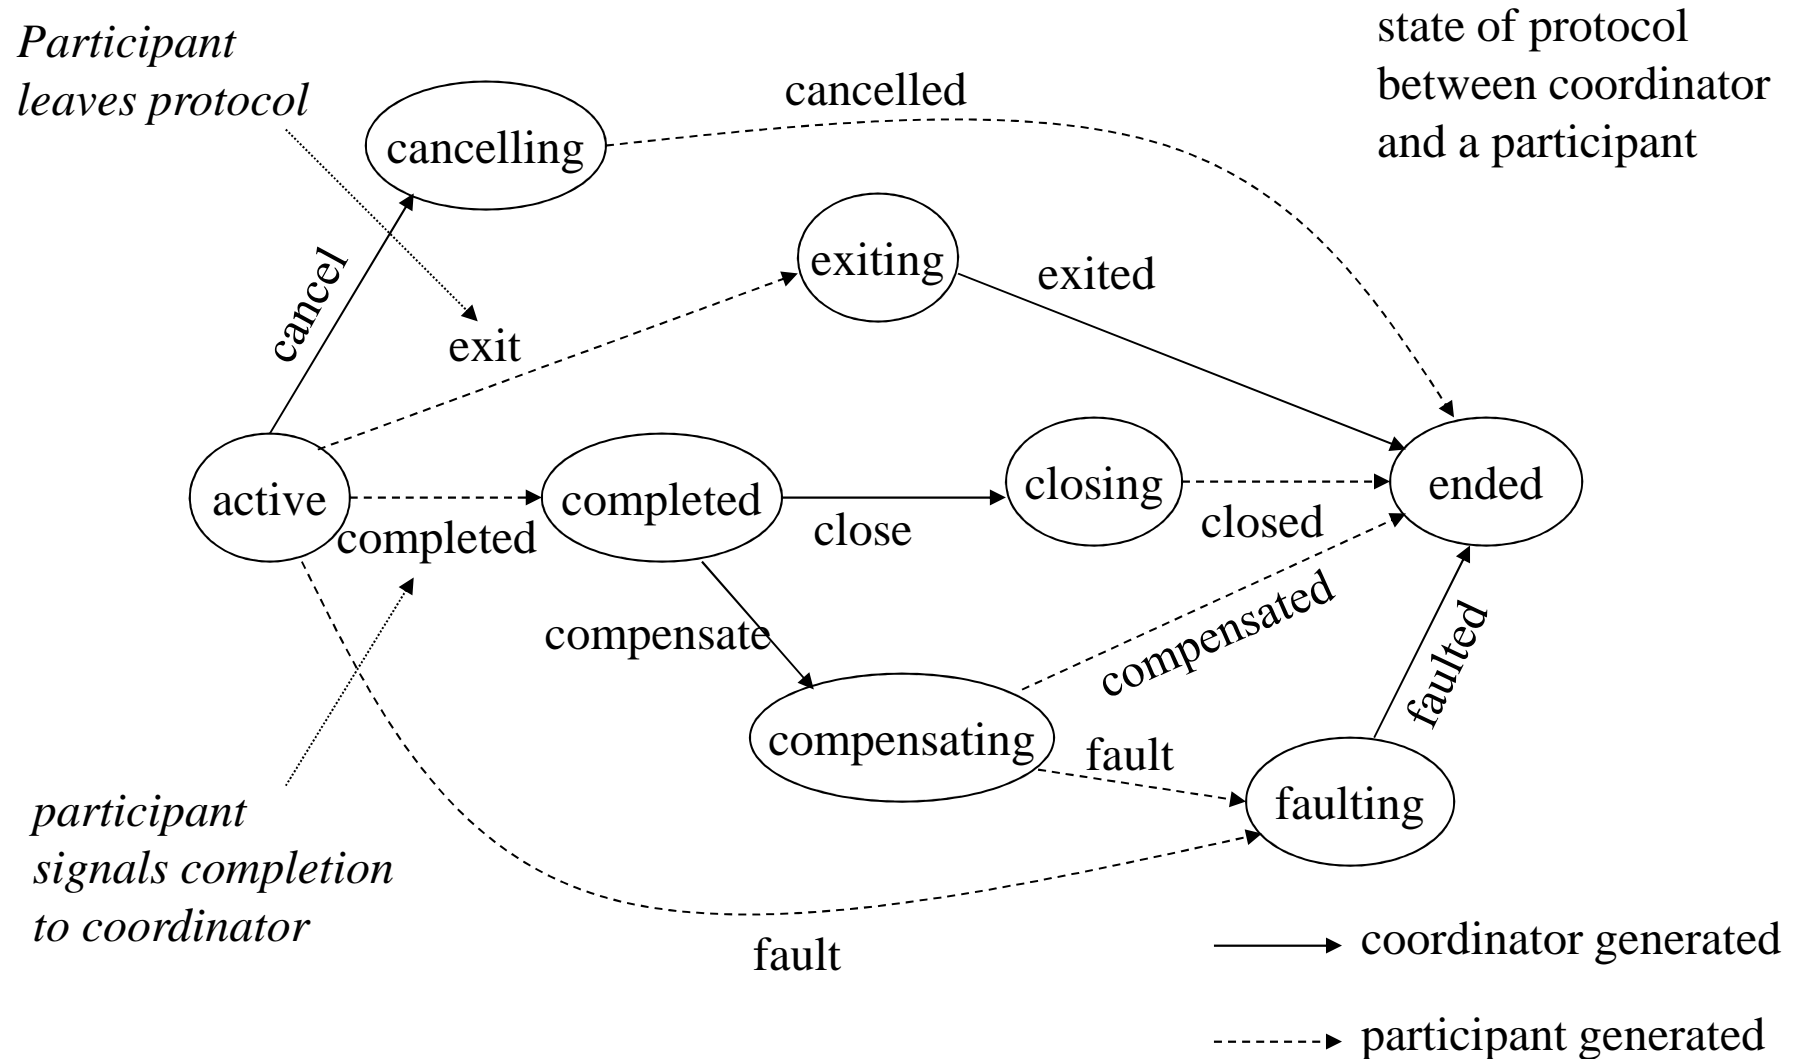
BAWPC Messages

- *Completed* – analagous to a vote message
- *Exit* – participant leaves the protocol
- *Cancel* – participant is forced out of the protocol by the coordinator
- *Close/Compensate* – in completed state coordinator decides on outcome based on coordination logic specified in context
- *Fault* – participant notifies coordinator that it has failed

Example

- Buyer sends copies of a purchase order to sellers *A*, *B*, *C*
 - *A* can't make a quote, responds with *fault* message
 - *B* and *C* make quotes using *completed* message
 - Indicates that *B* and *C* have completed successfully
- Buyer chooses *B*
 - Notifies *B* using *close* message
 - Indicates that *B* has successfully participated in a terminated business activity
 - Notifies *C* using *compensate* message
 - Indicates that *C* should rollback any action it has taken

BAWPC State Diagram



BAWCC State Diagram

