

Authentication, Authorization & Grid Security Infrastructure

Von Welch

Distributed Systems Laboratory

Univ. Of Chicago and Argonne National Laboratory





the globus project
www.globus.org

What do we want from security?

- Identity
- Authentication
- Privacy
- Integrity
- Authorization
- Single sign-on
- Delegation





Identity & Authentication

- Each entity should have an identity
 - Who are you?
 - Example: Unix login name
- Authentication:
 - Prove your identity
 - Stops masquerading imposters
- Examples:
 - Passport
 - Username and password







Message Protection

- Sending message securely
- Integrity
 - Detect whether message has been tampered
- Privacy
 - No one other than sender and receiver should be able to read message



Message Protection: Privacy




Medical Record
Patient no: 3456






the globus project
www.globus.org

Message Protection: Integrity



Run myHome/rm -f *



Run
myHome/whoami





Authorization establishes rights to do actions

- Establishing rights
- What can a particular identity do?

Examples:

- Are you allowed to read this file?
- Are you allowed to run a job on this machine?
- Unix read/write/execute permissions
- **Must authenticate first**
 - Authentication != authorization





Authorization

Examples:

- Are you allowed to be on this flight ?
 - > Passenger ?
 - > Pilot ?
- Unix read/write/execute permissions
- Must authenticate first





Single sign on

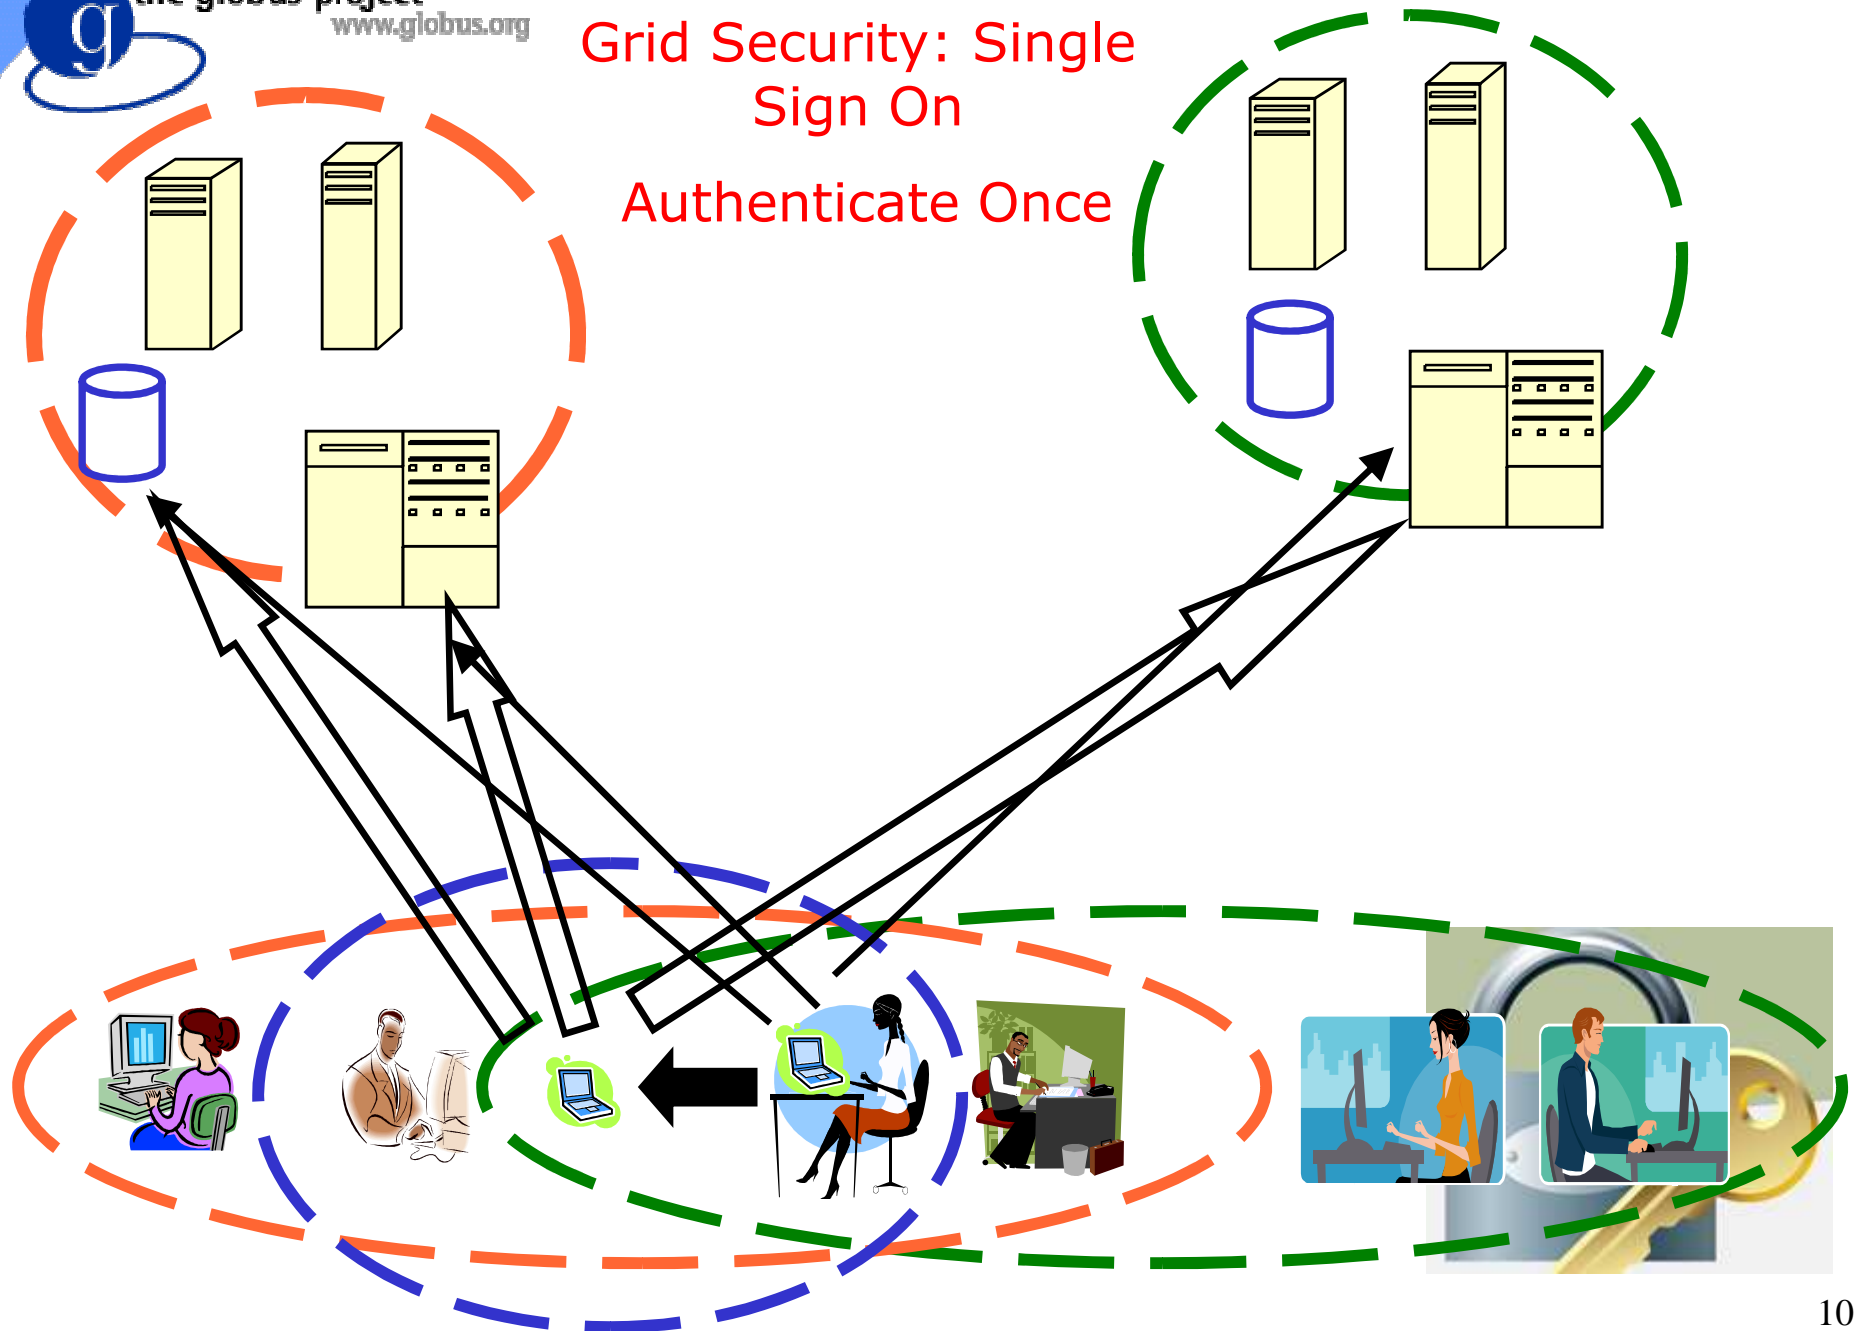
- Log on once
 - Type password once
- Use any grid resource without typing password again.
- Important for complex applications that need to use Grid resources
 - Enables easy coordination of varied resources
 - Enables automation of process
 - Allows remote processes and resources to act on user's behalf
 - Authentication and Delegation





the globus project
www.globus.org

Grid Security: Single Sign On Authenticate Once





Delegation

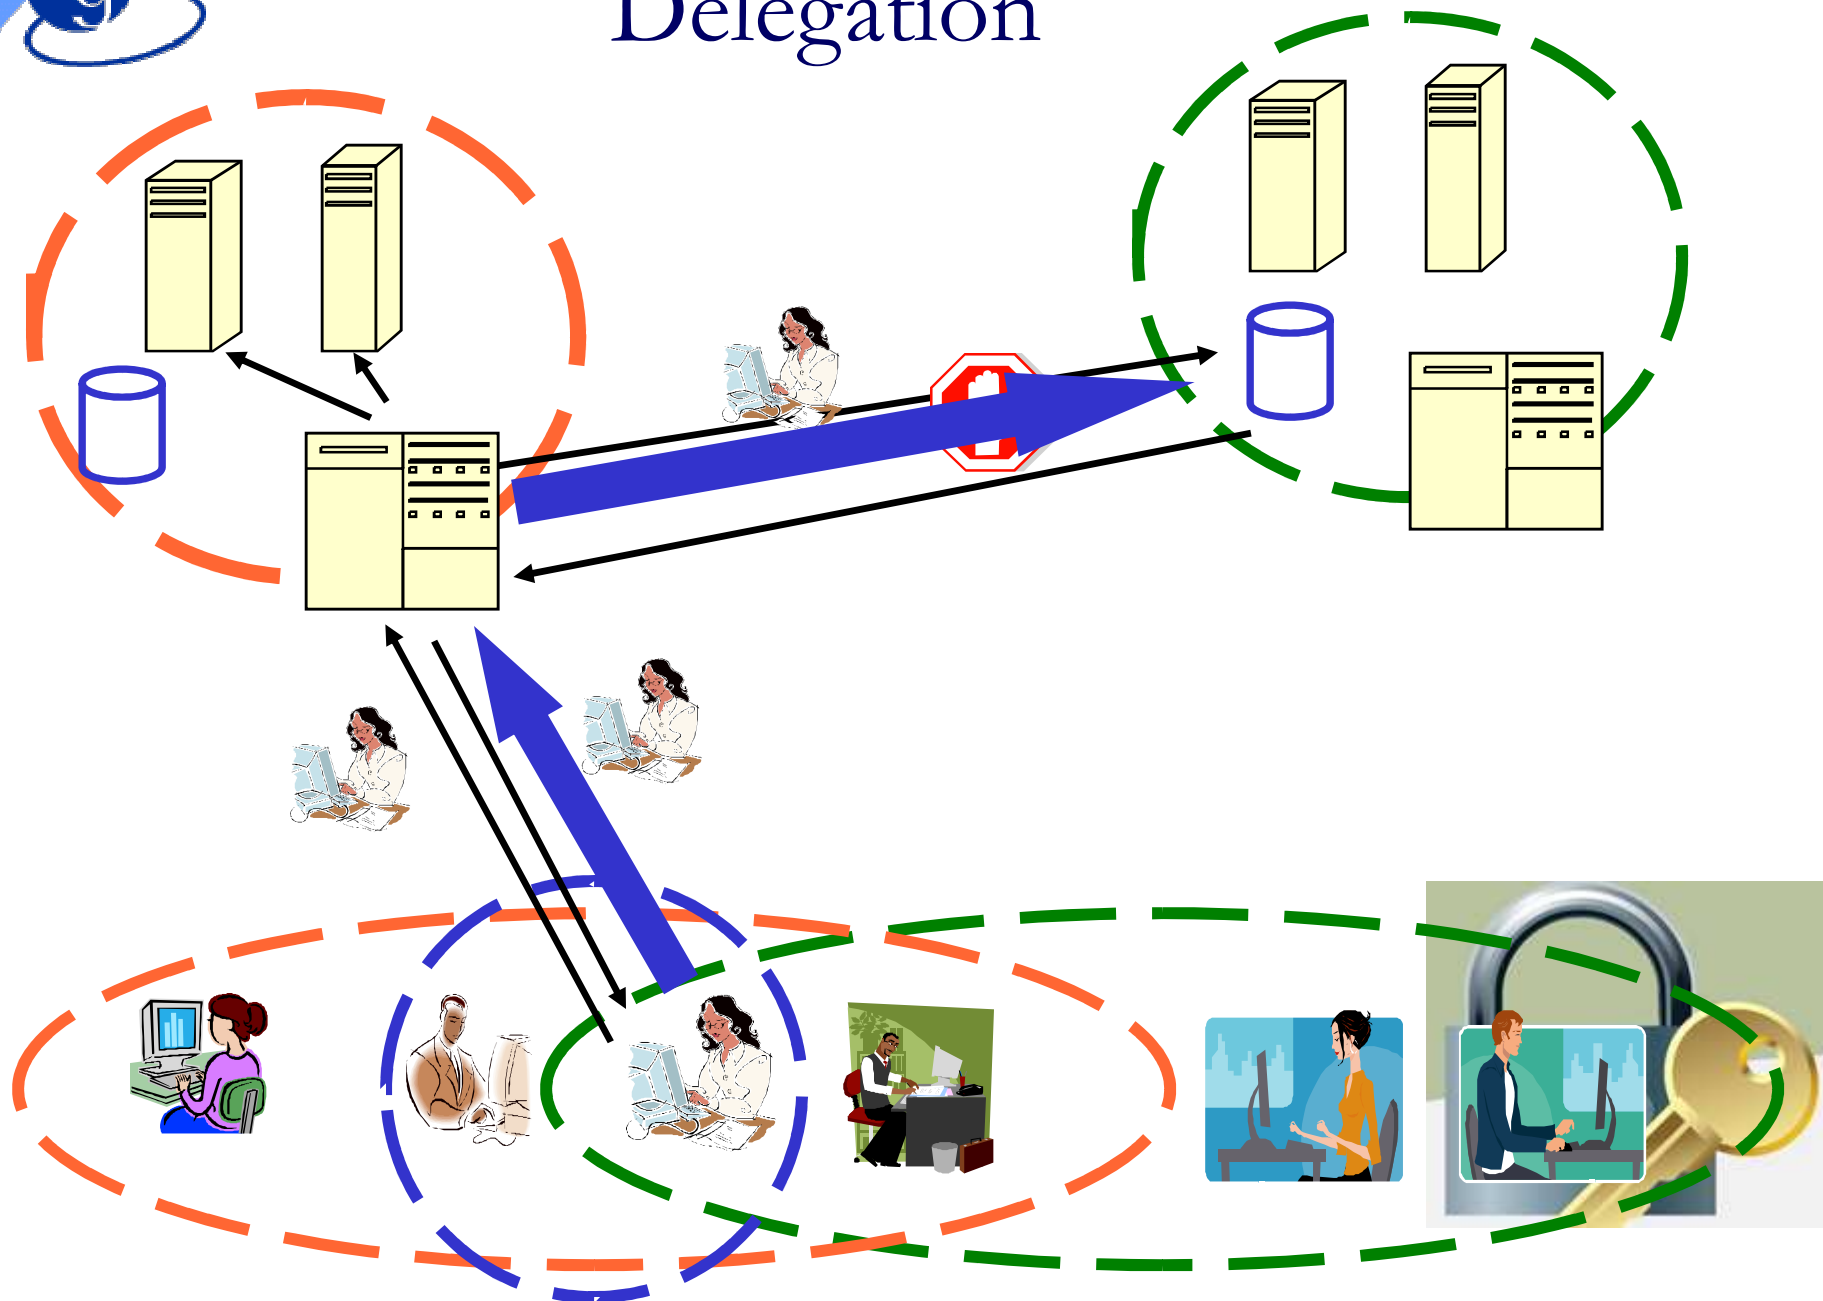
- Resources on the grid can act as you
- Example: Execution jobs can transfer files
- Delegation can be restricted
 - For example: Delegation only valid for a short period of time





the globus project
www.globus.org

Delegation





Solutions

- **Cryptography Overview**
- Public Key Infrastructure (PKI) Overview
- Secure Socket Layer (SSL) Overview
- Grid Security Infrastructure (GSI) Overview





Cryptography Overview

- Keys, Encryption and Decryption
 - Symmetric and Asymmetric
 - Public and Private keys
- Digital Signatures
 - Secure hashes





Keys


- A key can be thought of as simply a collection of bits
- The more bits, the stronger the key
- Keys are tied to specific encryption algorithms
- Lengths vary depending on the encryption algorithm
 - e.g. 128 bits is long for some algorithms, but short for others





Encryption


- Encryption is the process of taking some data and a key and feeding it into a function and getting encrypted data out
- Encrypted data is, in principal, unreadable unless decrypted





Decryption


- Decryption is the process of taking encrypted data and a key and feeding it into a function and getting out the original data
 - Encryption and decryption functions are linked





Symmetric Encryption

- Encryption and decryption functions that use the same key are called symmetric
 - In this case everyone wanting to read encrypted data must share the same key
- DES is an example of symmetric encryption





Asymmetric Encryption


- Encryption and decryption functions that use a key pair are called asymmetric
 - Keys are mathematically linked
- RSA is an example of asymmetric encryption





Asymmetric Encryption

- When data is encrypted with one key, the other key must be used to decrypt the data
 - And vice versa





Public and Private Keys

- With asymmetric encryption each user can be assigned a key pair: a private and public key



Private key is known only to owner



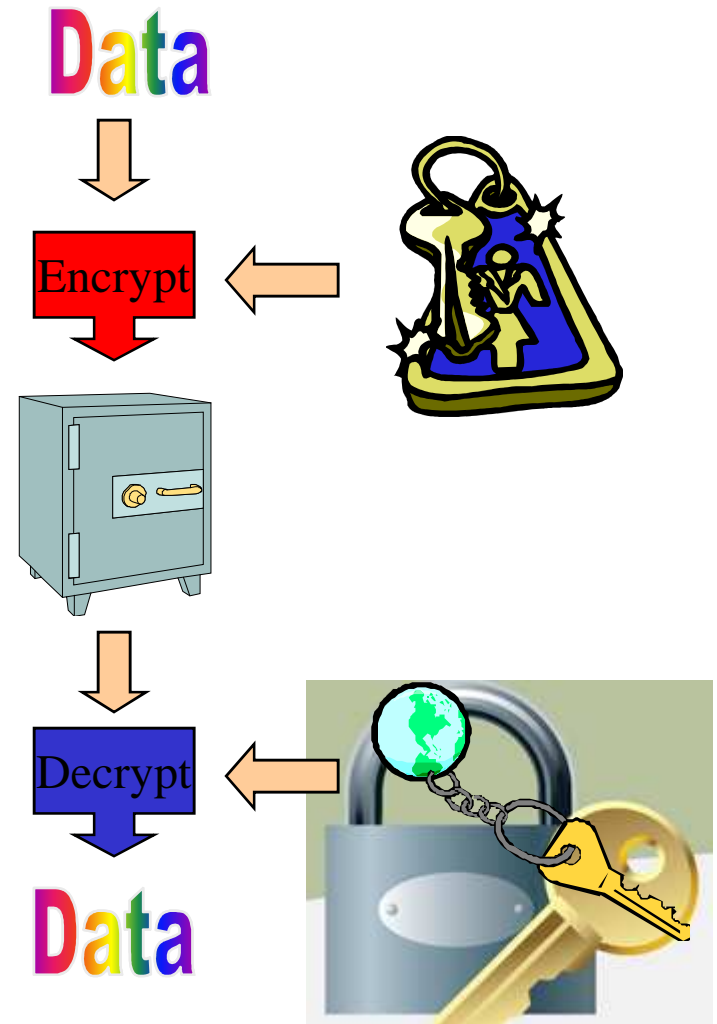
Public key is given away to the world





Public and Private keys


- Anything encrypted with the private key can only be decrypted with the public key
- And vice versa
- Since the private key is known only to the owner, this is very powerful...





Digital Signatures


- Digital signatures allow the world to verify I created a chunk of data
 - e.g. email, code





Digital Signatures


- Digital signatures are created by encrypting a hash of the data with my private key
- The resulting encrypted data is the signature
- This hash can then only be decrypted by my public key





Digital Signature



- Given some data with my signature, if you decrypt a signature with my public key and get the hash of the data, you know it was encrypted with my private key





Digital Signature

- Since I'm the only one with access to my private key, you know I signed the hash and the data associated with it
- But, how do you know that you have my correct public key?
- Answer: A Public Key Infrastructure...





the globus project
www.globus.org

Solutions

- Cryptography Overview
- **Public Key Infrastructure (PKI) Overview**
- Secure Socket Layer (SSL) Overview
- Grid Security Infrastructure (GSI) Overview





Public Key Infrastructure (PKI)

- PKI allows you to know that a given public key belongs to a given user
- PKI builds off of asymmetric encryption:
 - Each entity has two keys: public and private
 - The private key is known only to the entity
- The public key is given to the world encapsulated in a X.509 certificate





the globus project
www.globus.org

Public Key Infrastructure (PKI) Overview


- X.509 Certificates
- Certificate Authorities (CAs)
- Certificate Policies
 - Namespaces
- Requesting a certificate
 - Certificate Request
 - Registration Authority





Certificates

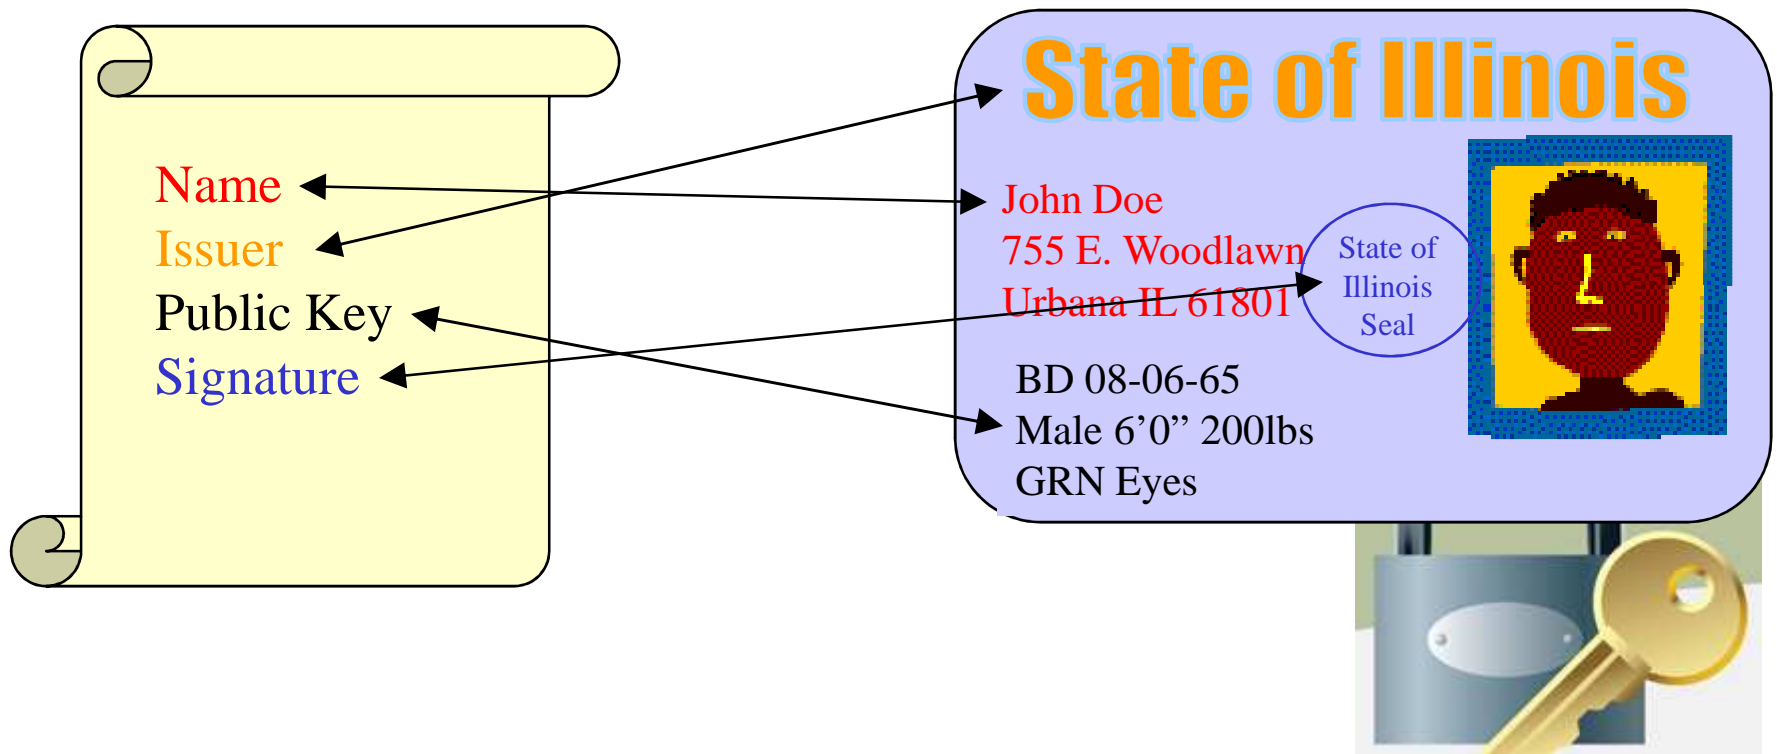
- A X.509 certificate binds a public key to a name
- It includes a name and a public key (among other things) bundled together and signed by a trusted party (Issuer)





Certificates


- Similar to passport or driver's license





Certificates


- By checking the signature, one can determine that a public key belongs to a given user.





Certificates

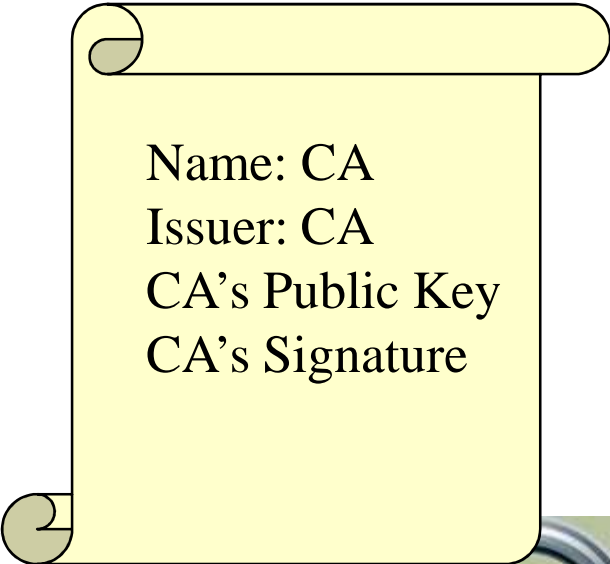
- Question: Who signs certificates?
- Answer: A small set of trusted entities known as Certificate Authorities (CAs)





Certificate Authorities (CAs)

- A Certificate Authority is an entity that exists only to sign user certificates
- The CA signs its own certificate which is distributed in a trusted manner




Name: CA
Issuer: CA
CA's Public Key
CA's Signature





Certificate Authorities (CAs)

- The public key from the CA certificate can then be used to verify other certificates





Certificate Policy (CP)

- Each CA has a Certificate Policy (CA) which states when and how a CA issues certificates.
- It states who it will issue certificates for
 - Just like the State of Illinois only issues driver's licenses' for residents of the state of Illinois
 - A CA for a grid typically only issues certificates for customers who are already approved to use resources on the grid





Certificate Policy (CP)

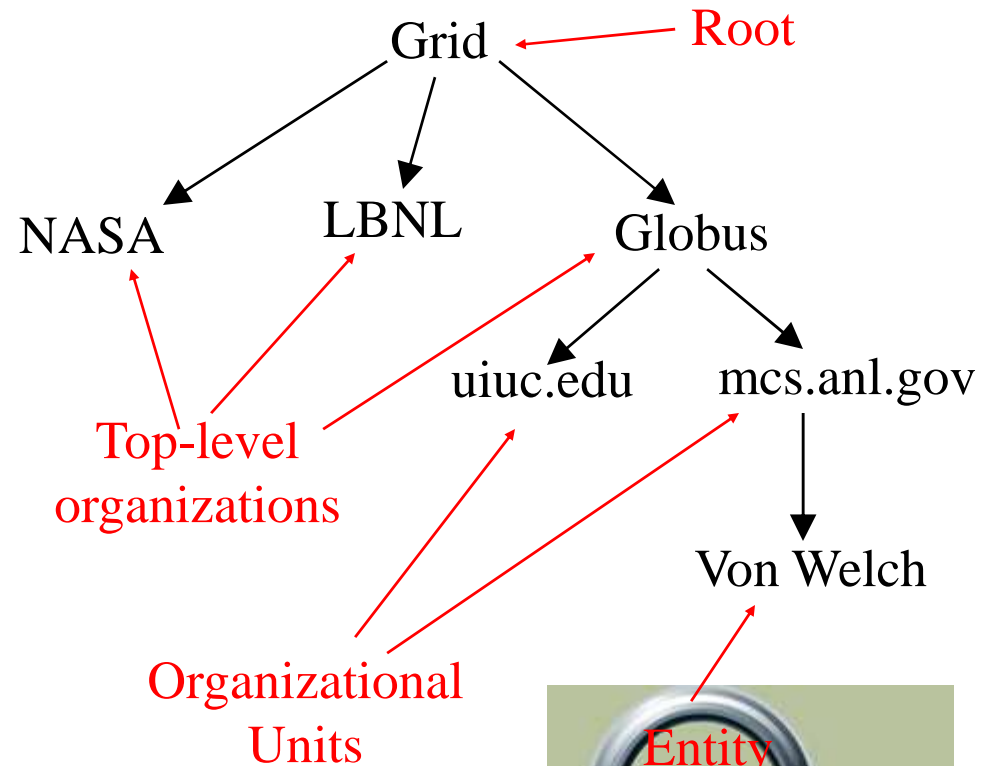
- A CA's CP states how it identifies the people it issues certificates to
 - Similar to having to show a birth certificate to get a driver's license
 - Some CA's are very stringent and require similar proof of identity
 - Others are lenient and only require proof via email





Namespaces


- Each CA's Certificate Policy also states the namespace of certificates issued by the CA
- A namespace is a hierarchy similar to the hierarchy used for Internet hostnames





Namespaces


- Each CA constrains itself to signing certificates that are in a namespace that are a portion of the overall space
 - E.g. the Globus CA signs certificates only under the Globus organization





Requesting a Certificate

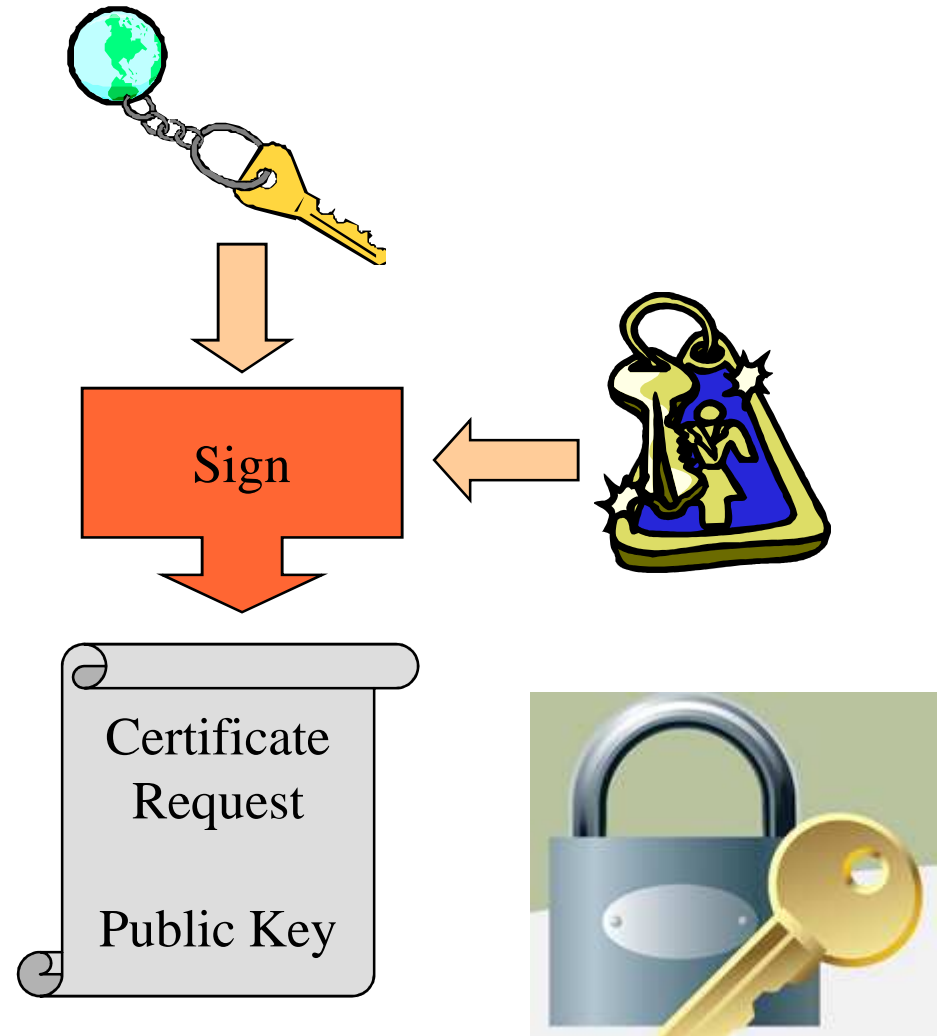
- To request a certificate a user starts by generating a key pair





Certificate Request

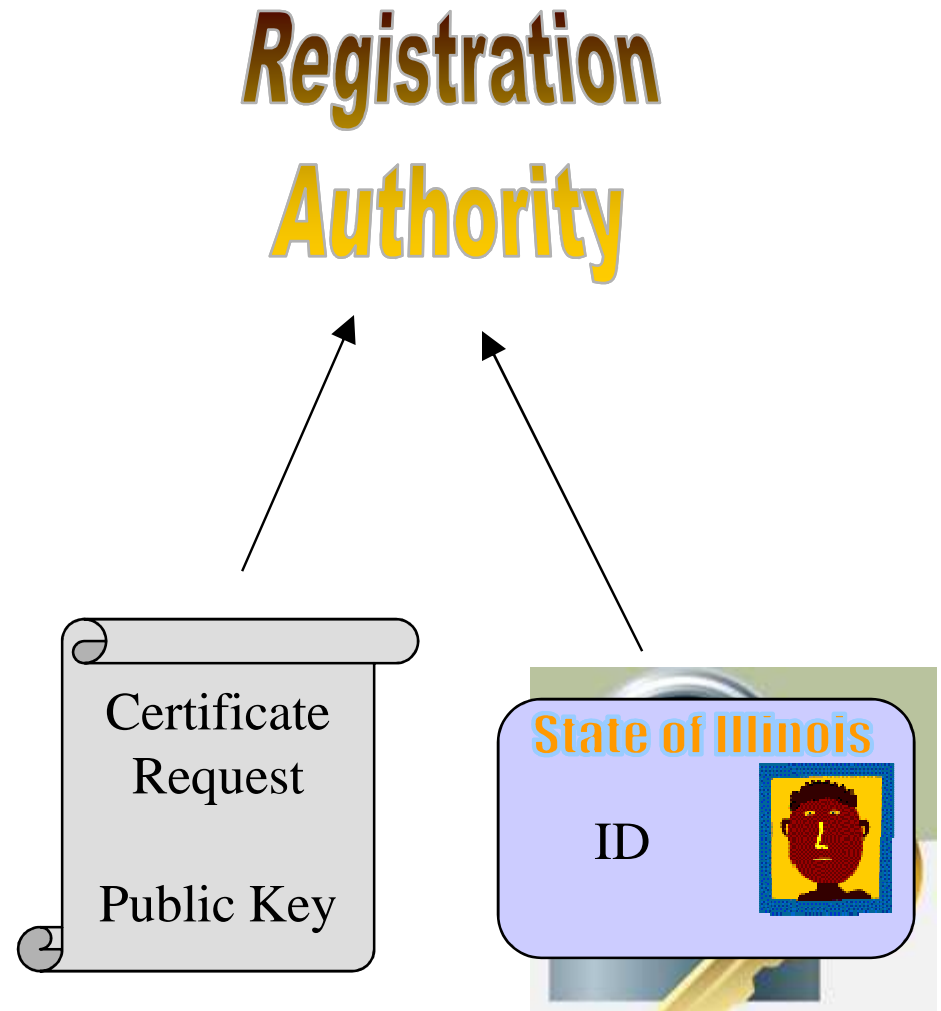
- The user then signs their own public key to form what is called a Certificate Request





Registration Authority (RA)

- The user then takes the certificate to a Registration Authority (RA)
- A RA's responsibility is to verify the user's name
- Often the RA coexists with the CA and is not apparent to the user





Certificate Issuance

- The CA then takes the identity from the RA and the public key from the certificate request
- It then creates, signs and issues a certificate for the user

*Registration
Authority*

Name

CA

Certificate
Request

Public Key

Name
Issuer
Public Key
Signature





Solutions

- Cryptography Overview
- Public Key Infrastructure (PKI) Overview
- **Secure Socket Layer (SSL) Overview**
- Grid Security Infrastructure (GSI) Overview





Secure Socket Layer (SSL)

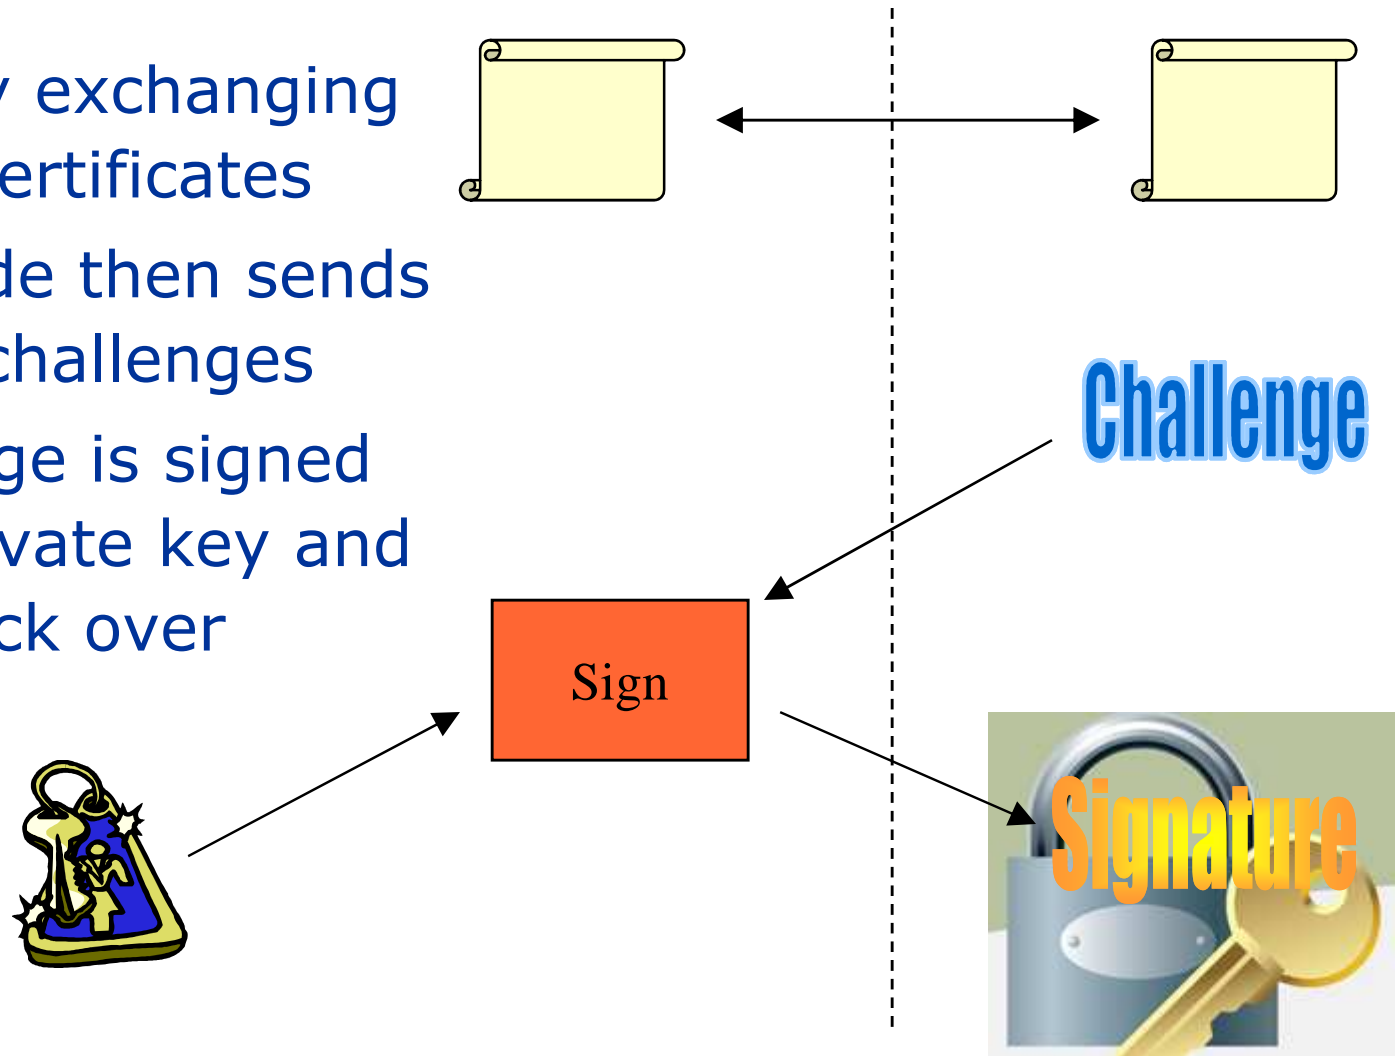
- Protocol above a standard TCP/IP socket to provide security in the forms of:
 - Authentication
 - Message protection
 - > Confidentiality
 - > Integrity





SSL Authentication

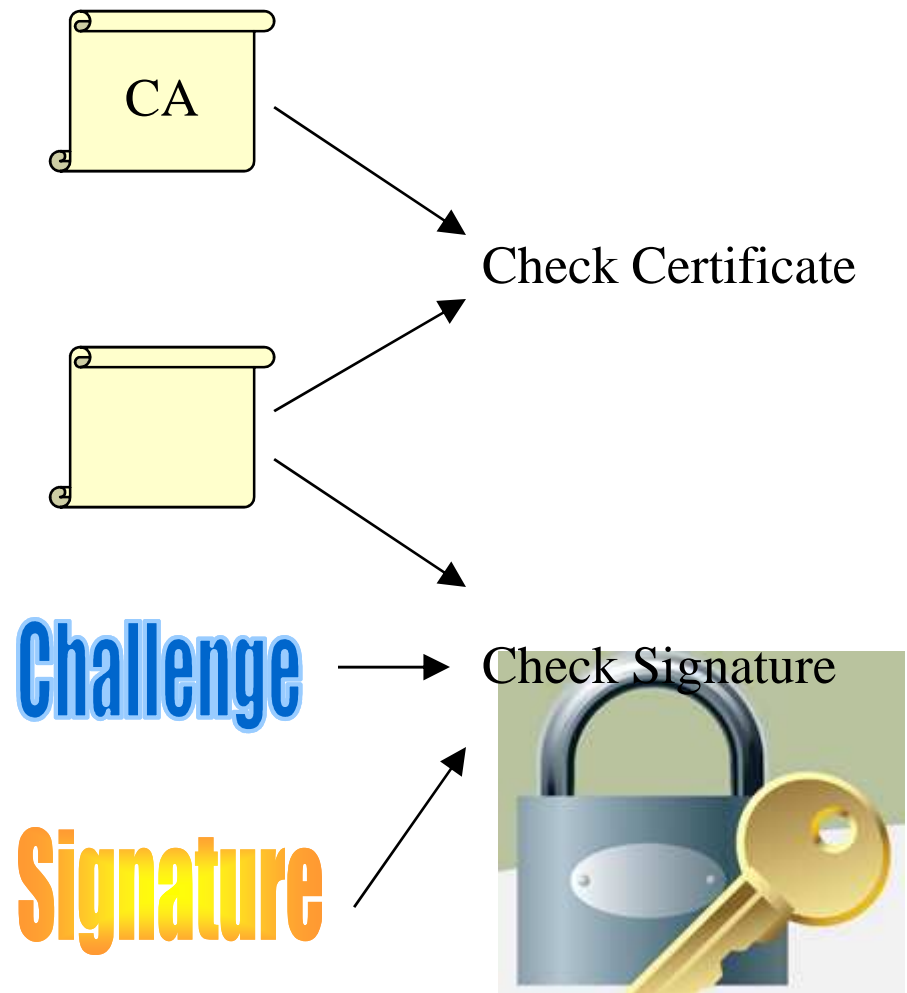
- Start by exchanging X.509 certificates
- Each side then sends over a challenge
- Challenge is signed with private key and sent back over





SSL Authentication


- Each side then verifies certificate using PKI and signature using certificate
- If everything checks then the identity from the certificate can be trusted





SSL Message Protection

- After authentication a shared session key is established to be used for message protection
- Confidentiality == Encryption of messages to prevent eavesdropping
- Integrity == Signing of messages to prevent modification





the globus project
www.globus.org

Solutions

- Cryptography Overview
- Public Key Infrastructure (PKI) Overview
- Secure Socket Layer (SSL) Overview
- **Grid Security Infrastructure (GSI) Overview**





Globus Security: The Grid Security Infrastructure

- The Grid Security Infrastructure (GSI) is a set of tools, libraries and protocols used in Globus to allow users and applications to securely access resources.
- Based on a public key infrastructure, with certificate authorities and X509 certificates



GSI

- Uses SSL for authentication and message protection
- Adds features needed for Single-Sign on
 - Proxy Credentials
 - Delegation





GSI: Credentials

- In the GSI system each user has a set of credentials they use to prove their identity on the grid
 - Consists of a X509 certificate and private key
- Long-term private key is kept encrypted with a pass phrase
 - Good for security, inconvenient for repeated usage





GSI: Single Sign-on

- Single-sign on is important feature for Grid Applications
 - Enables easy coordination of multiple resources
 - User authenticates themselves once, then can perform multiple actions without reauthentication
 - Can allow processes to act on their behalf





the globus project
www.globus.org

GSI: Single Sign-on

- To support single sign-on GSI adds the following functionality to SSL:
 - Proxy credentials
 - Credential delegation





GSI: Proxy Credentials

- Proxy credentials are short-lived credentials created by user
 - Short term binding of user's identity to alternate private key
 - Stored unencrypted for easy repeated access
 - Short lifetime in case of theft
 - Enables user to authenticate once then perform multiple actions without reauthenticating





GSI: Delegation

- GSI enables user to create and delegate proxy credentials to processes running on remote resources
- Allows remote processes and resources to act on user's behalf
- Important for complex applications that need to use Grid resources
 - E.g. jobs that needs to access data storage





Summary

- GSI is:
 - X.509 Certificates for authentication
 - PKI for verifying identities in Certificates
 - SSL as the protocol for authentication, confidentiality and integrity
 - Proxy Credentials and delegation to support single sign-on

