

The background is a deep blue gradient with a subtle pattern of white stars. Overlaid on the left side are several white circular and semi-circular lines, some with arrows indicating a clockwise direction. A large circular scale with numerical markings from 140 to 260 in increments of 10 is positioned on the left, with some numbers appearing to be part of a larger, partially visible scale. The overall aesthetic is scientific and cosmic.

N-BODY SOLVERS

PRESENTED BY
VIDYA NARASIMHAN

WHAT IS A N-BODY SOLVER ?

- **Problem** : Find the positions and velocities of a collection of interacting particles over a period of time
- **Definition** : An n -body solver is a program that finds the solution to an n -body problem by simulating the behavior of the particles.
- **Input to System** : mass, position, and velocity of each particle at the start of the simulation
- **Output of the System** : position and velocity of each particle at a sequence of user-specified times, or simply the position and velocity of each particle at the end of a user-specified time period.

APPLICATIONS

- Astro-Physics : To know the positions and velocities of a collection of stars,
- Chemistry : To know the positions and velocities of a collection of molecules or atoms

THE PROBLEM

- let's write an n -body solver that simulates the motions of planets or stars.
- We'll use Newton's second law of motion and his law of universal gravitation to determine the positions and velocities .

$$\mathbf{f}_{qk}(t) = -\frac{Gm_qm_k}{|\mathbf{s}_q(t) - \mathbf{s}_k(t)|^3} [\mathbf{s}_q(t) - \mathbf{s}_k(t)].$$

EXPANSION

- $f_{qk}(t)$: force on particle 'q' exerted by particle 'k' .
- $s_q(t)$: Position of particle 'q' at time 't' .
- $s_k(t)$: Position of particle 'k' at time 't' .
- G : Gravitational constant ($6.673 \times 10^{-11} \text{m}^3/(\text{kg} \cdot \text{s}^2)$)
- m_q and m_k : Mass of particles 'q' and 'k'
- $|s_q(t) - s_k(t)|$: Distance from particle k to particle q

TOTAL FORCE ON A PARTICLE

$$\mathbf{F}_q(t) = \sum_{\substack{k=0 \\ k \neq q}}^{n-1} \mathbf{f}_{qk} = -Gm_q \sum_{\substack{k=0 \\ k \neq q}}^{n-1} \frac{m_k}{|\mathbf{s}_q(t) - \mathbf{s}_k(t)|^3} [\mathbf{s}_q(t) - \mathbf{s}_k(t)].$$

- $F_q(t) = m_q(t) * a_q(t) = m_q(t) * s_q''(t)$

$$s_q''(t) = -G \sum_{\substack{j=0 \\ j \neq q}}^{n-1} \frac{m_j}{|\mathbf{s}_q(t) - \mathbf{s}_j(t)|^3} [\mathbf{s}_q(t) - \mathbf{s}_j(t)].$$

INPUT AND OUTPUT

- Required Output : find the positions and velocities at the times $t = 0, \Delta t, 2\Delta t, \dots, T\Delta t$,
- Given Input : n , the number of particles, Δt , T , and, for each particle, its mass, its initial position, and its initial velocity.
- Assumption : In a fully general solver, the positions and velocities would be three-dimensional vectors, but in order to keep things simple, we'll assume that the particles will move in a plane, and we'll use two-dimensional vectors instead.
-

SERIAL SOLUTION

In outline, a serial n -body solver can be based on the following pseudocode:

```
1   Get input data;
2   for each timestep {
3       if (timestep output) Print positions and velocities of
        particles;
4       for each particle q
5           Compute total force on q;
6       for each particle q
7           Compute position and velocity of q;
8   }
9   Print positions and velocities of particles;
```


SERIAL SOLUTION

```
for each particle q {  
  for each particle k != q {  
    x_diff = pos[q][X] - pos[k][X];  
    y_diff = pos[q][Y] - pos[k][Y];  
    dist = sqrt(x_diff*x_diff + y_diff*y_diff);  
    dist_cubed = dist*dist*dist;  
    forces[q][X] -= G*masses[q]*masses[k]/dist_cubed * x_diff;  
    forces[q][Y] -= G*masses[q]*masses[k]/dist_cubed * y_diff;  
  }  
}
```

$$\begin{bmatrix} 0 & \mathbf{f}_{01} & \mathbf{f}_{02} & \cdots & \mathbf{f}_{0,n-1} \\ -\mathbf{f}_{01} & 0 & \mathbf{f}_{12} & \cdots & \mathbf{f}_{1,n-1} \\ -\mathbf{f}_{02} & -\mathbf{f}_{12} & 0 & \cdots & \mathbf{f}_{2,n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ -\mathbf{f}_{0,n-1} & -\mathbf{f}_{1,n-1} & -\mathbf{f}_{2,n-1} & \cdots & 0 \end{bmatrix}.$$

```

for each particle q
    forces[q] = 0;
for each particle q {
    for each particle k > q {
        x_diff = pos[q][X] - pos[k][X];
        y_diff = pos[q][Y] - pos[k][Y];
        dist = sqrt(x_diff*x_diff + y_diff*y_diff);
        dist_cubed = dist*dist*dist;
        force_qk[X] = G*masses[q]*masses[k]/dist_cubed * x_diff;
        force_qk[Y] = G*masses[q]*masses[k]/dist_cubed * y_diff

        forces[q][X] += force_qk[X];
        forces[q][Y] += force_qk[Y];
        forces[k][X] -= force_qk[X];
        forces[k][Y] -= force_qk[Y];
    }
}

```

Program 6.1: A reduced algorithm for computing n -body forces

USING TANGENT TO APPROXIMATE A FUNCTION




FIGURE 6.1

Using the tangent line to approximate a function

$$g(t + \Delta t) \approx g(t_0) + g'(t_0)(t + \Delta t - t) = g(t_0) + \Delta t g'(t_0).$$

$$\mathbf{s}_q(\Delta t) \approx \mathbf{s}_q(0) + \Delta t \mathbf{s}'_q(0) = \mathbf{s}_q(0) + \Delta t \mathbf{v}_q(0),$$

$$\mathbf{v}_q(\Delta t) \approx \mathbf{v}_q(0) + \Delta t \mathbf{v}'_q(0) = \mathbf{v}_q(0) + \Delta t \mathbf{a}_q(0) = \mathbf{v}_q(0) + \Delta t \frac{1}{m_q} \mathbf{F}_q(0).$$

```
pos[q][X] += delta_t*vel[q][X];  
pos[q][Y] += delta_t*vel[q][Y];  
vel[q][X] += delta_t/masses[q]*forces[q][X];  
vel[q][Y] += delta_t/masses[q]*forces[q][Y];
```


DATA STRUCTURE

- For each particle we need to know ,
 - Mass
 - Position
 - Initial velocity
 - Acceleration
 - Total force acting on it
- For each particle it suffices to store its mass and the current value of its position, velocity, and force.
- We could store these four variables as a struct and use an array of structs to store the data for all the particles.

PARALLELIZING THE N-BODY SOLVERS

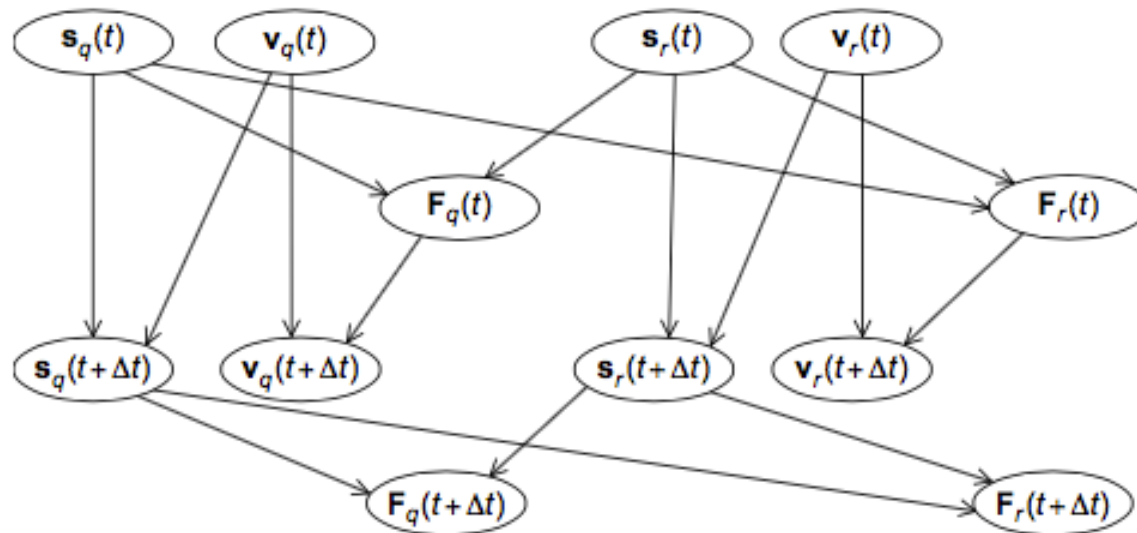



FIGURE 6.3

Communications among tasks in the basic n -body solver



HOW TO SPLIT THE TASK BETWEEN THREADS ?

- If we have n particles and T timesteps, then there will be nT tasks in both the basic and the reduced algorithm.
- Astrophysical n -body problems typically involve thousands or even millions of particles, so n is likely to be several orders of magnitude greater than the number of available cores.
- However, T may also be much larger than the number of available cores. So, in principle, we have two “dimensions” to work with when we map tasks to cores
- Attempting to assign tasks associated with a single particle at different timesteps to different cores won't work very well. Before estimating $\mathbf{s}q(t+1t)$ and $\mathbf{v}q(t+1t)$, Euler's method must “know” $\mathbf{s}q(t)$, $\mathbf{v}q(t)$, and $\mathbf{a}q(t)$.

- Before estimating $\mathbf{s}q(t+\Delta t)$ and $\mathbf{v}q(t+\Delta t)$, Euler's method must "know" $\mathbf{s}q(t)$, $\mathbf{v}q(t)$, and $\mathbf{a}q(t)$. Thus, if we assign particle q at time t to core c_0 , and we assign particle q at time $t + \Delta t$ to core $c_1 \neq c_0$, then we'll have to communicate $\mathbf{s}q(t)$, $\mathbf{v}q(t)$, and $\mathbf{F}q(t)$ from c_0 to c_1 .
- Of course, if particle q at time t and particle q at time $t + \Delta t$ are mapped to the same core, this communication won't be necessary, so once we've mapped the task consisting of the calculations for particle q at the first timestep to core c_0 , we may as well map the subsequent computations for particle q to the same cores, since we can't simultaneously execute the computations for particle q at two different timesteps.
- Thus, mapping tasks to cores will, in effect, be an assignment of particles to cores.
- Assignment of particles to cores that assigns roughly $n/\text{thread count}$ particles to each core will do a good job of balancing the workload among the cores

n-body solver using OpenMP

The n-body problem

- Find the positions and velocities of a collection of interacting particles over a period of time
- For example, **an astrophysicist** might want to know the positions and velocities of a collection of stars, while **a chemist** might want to know the positions and velocities of a collection of molecules or atoms

n-body solvers

An n-body solver is a program that finds the solution to an n-body problem by simulating the behavior of the particles.




Figure: n-body solver

n-body solvers example

- An n-body solver that simulates the motions of planets or stars
- We use Newton's second law of motion and law of universal gravitation

$$\mathbf{f}_{qk}(t) = -\frac{Gm_qm_k}{|\mathbf{s}_q(t) - \mathbf{s}_k(t)|^3} [\mathbf{s}_q(t) - \mathbf{s}_k(t)].$$

Figure: Force on a particle q exerted by a particle k

G- gravitational constant

n-body solvers serial implementation

```
Get input data;
for each timestep {
    if (timestep output) Print positions and velocities of
        particles;
    for each particle q
        Compute total force on q;
    for each particle q
        Compute position and velocity of q;
}
Print positions and velocities of particles;
```

Figure: Serial pseudocode

Total force on a particle

The total force on a particle can be found by adding the forces due to all the particles

$$\mathbf{F}_q(t) = \sum_{\substack{k=0 \\ k \neq q}}^{n-1} \mathbf{f}_{qk} = -Gm_q \sum_{\substack{k=0 \\ k \neq q}}^{n-1} \frac{m_k}{|\mathbf{s}_q(t) - \mathbf{s}_k(t)|^3} [\mathbf{s}_q(t) - \mathbf{s}_k(t)].$$

```
for each particle q {  
    for each particle k != q {  
        x_diff = pos[q][X] - pos[k][X];  
        y_diff = pos[q][Y] - pos[k][Y];  
        dist = sqrt(x_diff*x_diff + y_diff*y_diff);  
        dist_cubed = dist*dist*dist;  
        forces[q][X] -= G*masses[q]*masses[k]/dist_cubed * x_diff;  
        forces[q][Y] -= G*masses[q]*masses[k]/dist_cubed * y_diff;  
    }  
}
```

Newton's Third law


$$\begin{bmatrix} 0 & \mathbf{f}_{01} & \mathbf{f}_{02} & \cdots & \mathbf{f}_{0,n-1} \\ -\mathbf{f}_{01} & 0 & \mathbf{f}_{12} & \cdots & \mathbf{f}_{1,n-1} \\ -\mathbf{f}_{02} & -\mathbf{f}_{12} & 0 & \cdots & \mathbf{f}_{2,n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ -\mathbf{f}_{0,n-1} & -\mathbf{f}_{1,n-1} & -\mathbf{f}_{2,n-1} & \cdots & 0 \end{bmatrix}.$$

Reduced algorithm

```
for each particle q
    forces[q] = 0;
for each particle q {
    for each particle k > q {
        x_diff = pos[q][X] - pos[k][X];
        y_diff = pos[q][Y] - pos[k][Y];
        dist = sqrt(x_diff*x_diff + y_diff*y_diff);
        dist_cubed = dist*dist*dist;
        force_qk[X] = G*masses[q]*masses[k]/dist_cubed * x_diff;
        force_qk[Y] = G*masses[q]*masses[k]/dist_cubed * y_diff

        forces[q][X] += force_qk[X];
        forces[q][Y] += force_qk[Y];
        forces[k][X] -= force_qk[X];
        forces[k][Y] -= force_qk[Y];
    }
}
```


Euler's Method



$$y = g(t_0) + g'(t_0)(t - t_0).$$

Since we're interested in the time $t = t_0 + \Delta t$, we get

$$g(t + \Delta t) \approx g(t_0) + g'(t_0)(t + \Delta t - t) = g(t_0) + \Delta t g'(t_0).$$

$$\mathbf{s}_q(\Delta t) \approx \mathbf{s}_q(0) + \Delta t \mathbf{s}'_q(0) = \mathbf{s}_q(0) + \Delta t \mathbf{v}_q(0),$$

$$\mathbf{v}_q(\Delta t) \approx \mathbf{v}_q(0) + \Delta t \mathbf{v}'_q(0) = \mathbf{v}_q(0) + \Delta t \mathbf{a}_q(0) = \mathbf{v}_q(0) + \Delta t \frac{1}{m_q} \mathbf{F}_q(0).$$

Pseudocode to compute the position and velocity

```
pos[q][X] += delta_t*vel[q][X];  
pos[q][Y] += delta_t*vel[q][Y];  
vel[q][X] += delta_t/masses[q]*forces[q][X];  
vel[q][Y] += delta_t/masses[q]*forces[q][Y];
```

Parallelizing the basic solver using OpenMP

```
for each timestep {  
    if (timestep output) Print positions and velocities of  
        particles;  
#   pragma omp parallel for  
    for each particle q  
        Compute total force on q;  
#   pragma omp parallel for  
    for each particle q  
        Compute position and velocity of q;  
}
```



```
# pragma omp parallel for
  for each particle q {
    forces[q][X] = forces[q][Y] = 0;
    for each particle k != q {
      x_diff = pos[q][X] - pos[k][X];
      y_diff = pos[q][Y] - pos[k][Y];

      dist = sqrt(x_diff*x_diff + y_diff*y_diff);
      dist_cubed = dist*dist*dist;
      forces[q][X] -= G*masses[q]*masses[k]/dist_cubed * x_diff;
      forces[q][Y] -= G*masses[q]*masses[k]/dist_cubed * y_diff;
    }
  }
```

Parallelizing the second loop

```
# pragma omp parallel for
for each particle q {
    pos[q][X] += delta_t*vel[q][X];
    pos[q][Y] += delta_t*vel[q][Y];
    vel[q][X] += delta_t/masses[q]*forces[q][X];
    vel[q][Y] += delta_t/masses[q]*forces[q][Y];
}
```

parallel directive for the outermost loop

```
# pragma omp parallel
  for each timestep {
    if (timestep output) Print positions and velocities of
      particles;
#   pragma omp for
    for each particle q

        Compute total force on q;
#   pragma omp for
    for each particle q
      Compute position and velocity of q;
  }
```


Adding the single directive

```
# pragma omp parallel
  for each timestep {
    if (timestep output) {
#       pragma omp single
        Print positions and velocities of particles;
    }
#   pragma omp for
    for each particle q
        Compute total force on q;
#   pragma omp for
    for each particle q
        Compute position and velocity of q;
  }
```

Parallelizing the reduced solver using OpenMP

```
# pragma omp parallel
  for each timestep {
    if (timestep output) {
#       pragma omp single
        Print positions and velocities of particles;
    }
#   pragma omp for
    for each particle q
      forces[q] = 0.0;
#   pragma omp for
    for each particle q
      Compute total force on q;
#   pragma omp for
    for each particle q
      Compute position and velocity of q;
  }
```

Critical Directive

```
# pragma omp critical
{
    forces[q][X] += force_qk[X];
    forces[q][Y] += force_qk[Y];
    forces[k][X] -= force_qk[X];
    forces[k][Y] -= force_qk[Y];
}
```

```
omp_set_lock(&locks[q]);  
forces[q][X] += force_qk[X];  
forces[q][Y] += force_qk[Y];  
omp_unset_lock(&locks[q]);
```

```
omp_set_lock(&locks[k]);  
forces[k][X] -= force_qk[X];  
forces[k][Y] -= force_qk[Y];  
omp_unset_lock(&locks[k]);
```