

SIMD systems

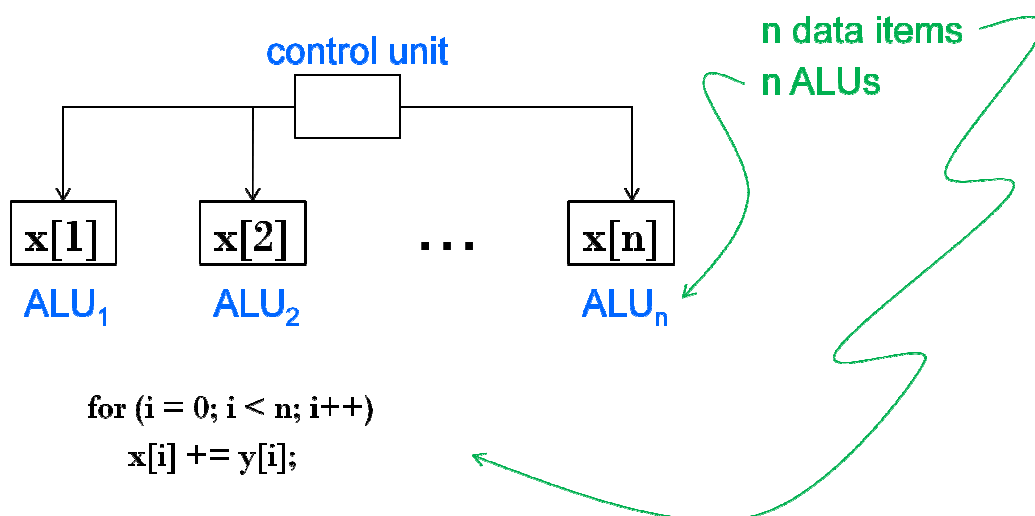
- Single instruction, multiple data, or SIMD, systems are parallel systems. As the name suggests, SIMD systems operate on multiple data streams by applying the same instruction to multiple data items.
- SIMD system can be thought of as having a single control unit and multiple ALUs. An instruction is broadcast from the control unit to the ALUs, and each ALU either applies the instruction to the current data item, or it is idle. As an example, suppose we want to carry out a “vector addition.”
- That is, suppose we have two arrays x and y , each with n elements, and we want to add the elements of y to the elements of x :

```
for (i = 0; i < n; i++)  
    x[i] += y[i];
```

- Suppose SIMD system has n ALUs. Then we could load $x[i]$ and $y[i]$ into the i th ALU, have the i th ALU add $y[i]$ to $x[i]$, and store the result in $x[i]$. If the system has m ALUs and $m < n$, we can simply execute the additions in blocks of m elements at a time.
 - For example, if $m = 4$ and $n = 15$, we can first add elements 0 to 3, then elements 4 to 7, then elements 8 to 11, and finally elements 12 to 14. Note that in the last group of elements in our example—elements 12 to 14—we’re only operating on three elements of x and y , so one of the four ALUs will be idle. The requirement that all the ALUs execute the same instruction or are idle can seriously degrade the overall performance of a SIMD system.
- SIMD Architectures have significant DLP. Single Instruction can launch many data ops
- SIMD is more energy efficient than MIMD
 - MIMD needs to fetch and execute one instruction per data ops.
 - SIMD is more attractive for PMDs.
 - Advantage of SIMD over MIMD

- Programmer thinks sequential execution yet achieves parallel speedup by having parallel data operations
- Parallelism achieved by dividing data among the processors.
- Applies the same instruction to multiple data items.
- Called data parallelism.
- SIMD has 3 variations:
- Vector Architectures
- Allows pipelined execution of many data operations
- SIMD MMX
- Allows simultaneous parallel data operations that support Multimedia applications.
- GPU Architectures
- They offer higher performance than traditional multicore
- They have system processor, system memory & graphics memory.

SIMD example



- What if we don't have as many ALUs as data items?
- Divide the work and process iteratively.
- Ex. $m = 4$ ALUs and $n = 15$ data items.

| Round | ALU ₁ | ALU ₂ | ALU ₃ | ALU ₄ |
|-------|------------------|------------------|------------------|------------------|
| 1 | X[0] | X[1] | X[2] | X[3] |
| 2 | X[4] | X[5] | X[6] | X[7] |
| 3 | X[8] | X[9] | X[10] | X[11] |
| 4 | X[12] | X[13] | X[14] | |

- All ALUs are required to execute the same instruction, or remain idle.
- In classic design, they must also operate synchronously.
- The ALUs have no instruction storage.
- Efficient for large data parallel problems, but not other types of more complex parallel problems.
- SIMD drawbacks