

Classification Methods

1. Classification- Basic Concepts
2. Classification- Two step process
3. Example of Classification
4. Decision tree Induction
5. Attribute selection Measures

Classification –Basic Concepts

Classification:

- Classification is a form of data analysis task where a model or classifier is constructed to predict class labels.
- loan application data : “safe” or “risky” and Marketing data yes” or “no”
- Labels are categorical and represented as discrete values and unordered.
- Analysis provide better understanding of the data at large.
- Classification is a two step process
- **Learning step** - Building of classification model based on training data
- **Classification step**- If the model’s accuracy is acceptable, use the model to classify the test data.



Classification –Basic Concepts

- Typical applications
 - Credit/loan approval
 - Medical diagnosis: if a tumor is cancerous or benign
 - Fraud detection: if a transaction is fraudulent
 - Web page categorization: which category it is
- In each of these examples the data analysis task is **classification** where a **model or classifier** is constructed to **predict class**

Classification—A Two-Step Process

- Data classification is a two step process consisting of **Learning and Classification** step
- **Learning step:** The classification algorithm built a classifier by analyzing or “learning from” a training set.
- Training set made of database tuples and their associated class label.
- Tuple X is represented by n dimensional vector $n(x_1, x_2, x_3, x_4 \dots x_n)$ depicting n measurements.
- Each X belong to a predefined class and the class label attribute is discrete-valued and unordered.
- The individual tuples making up the training set are referred to as **training tuples**



Classification—A Two-Step Process

- Training tuples are randomly sampled from the database under analysis.
- Classification process learns mapping or function $y=f(X)$ that predict the class label y of a given tuple X .
- The mapping or function separates the data classes.
- Mapping represented as classification rules, decision trees, or mathematical formula



Supervised vs. Unsupervised Learning

- **Supervised learning (classification):**
 - The learning of the classifier is “supervised” when the class labels of each training tuple is known.
 - New data is classified based on the training set.
- **Unsupervised learning (clustering)**
 - The class labels of training data is unknown.
 - Given a set of measurements, observations, etc. with the aim of establishing the existence of classes or clusters in the data

Classification—A Two-Step Process

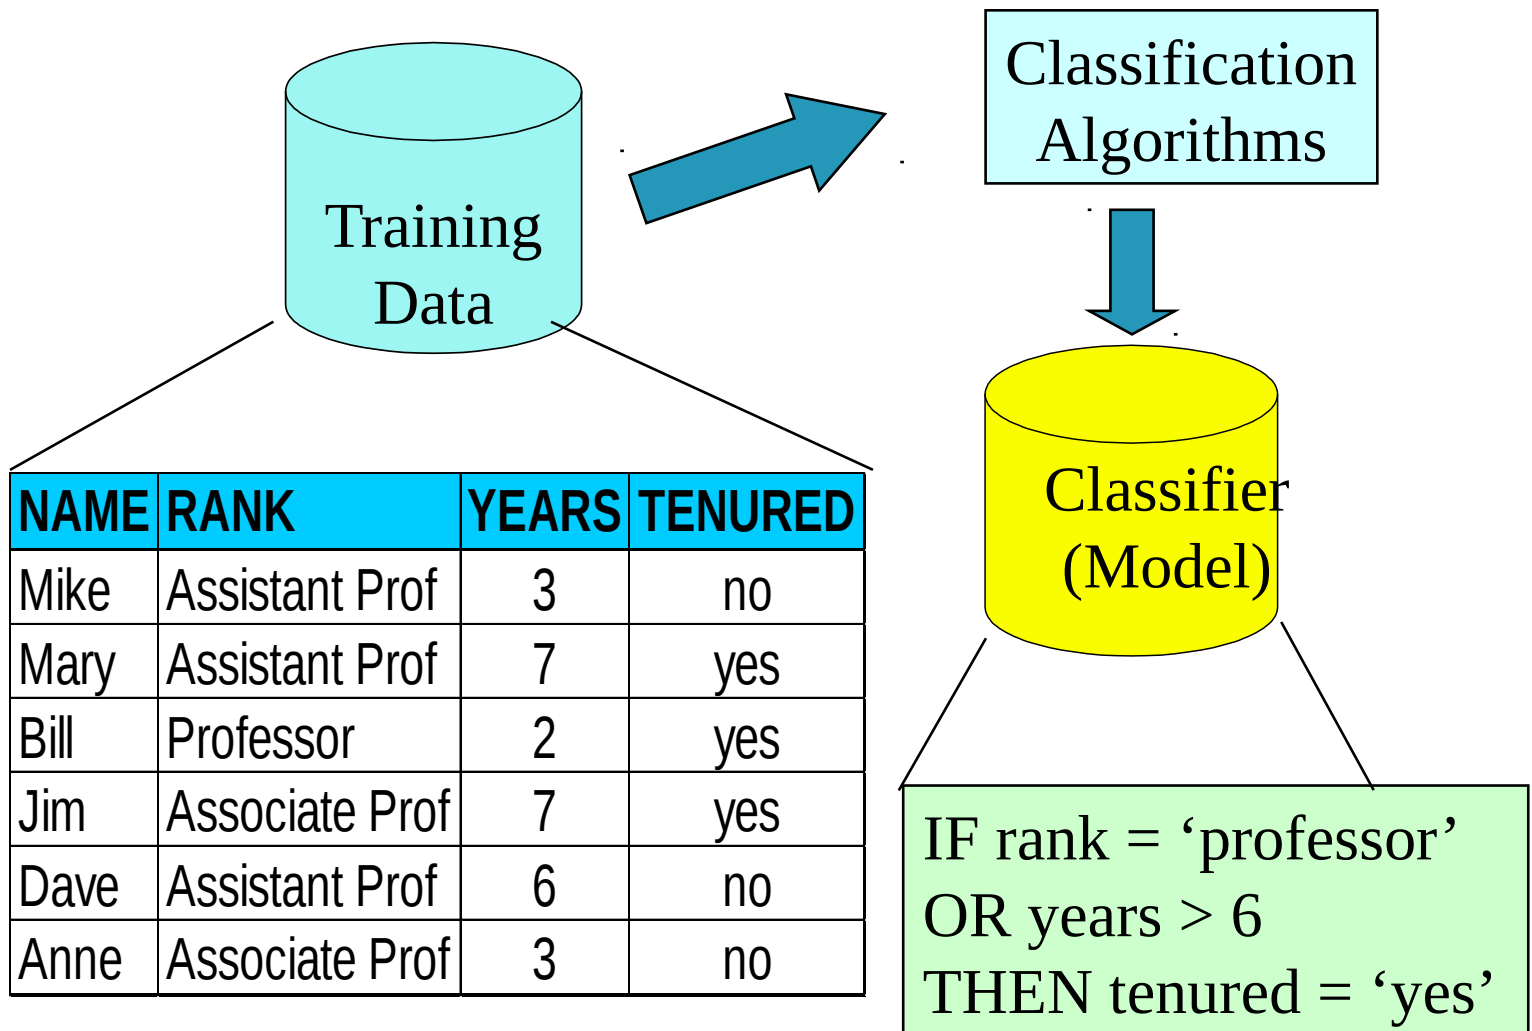
- **Classification step:** For classifying future or unknown objects use the estimated model (rules).
- The test data are used to estimate the accuracy of classification rules.
- Test set made of test tuples is independent of training set (otherwise overfitting)
 - If the accuracy is acceptable, use the model to classify new data
- The known label of test sample is compared with the classified result from the model.
- **Accuracy:** % of test set samples that are correctly classified by the model
- Note: If the test set is used to select/refine models, it is called **validation (test) set or development test set**



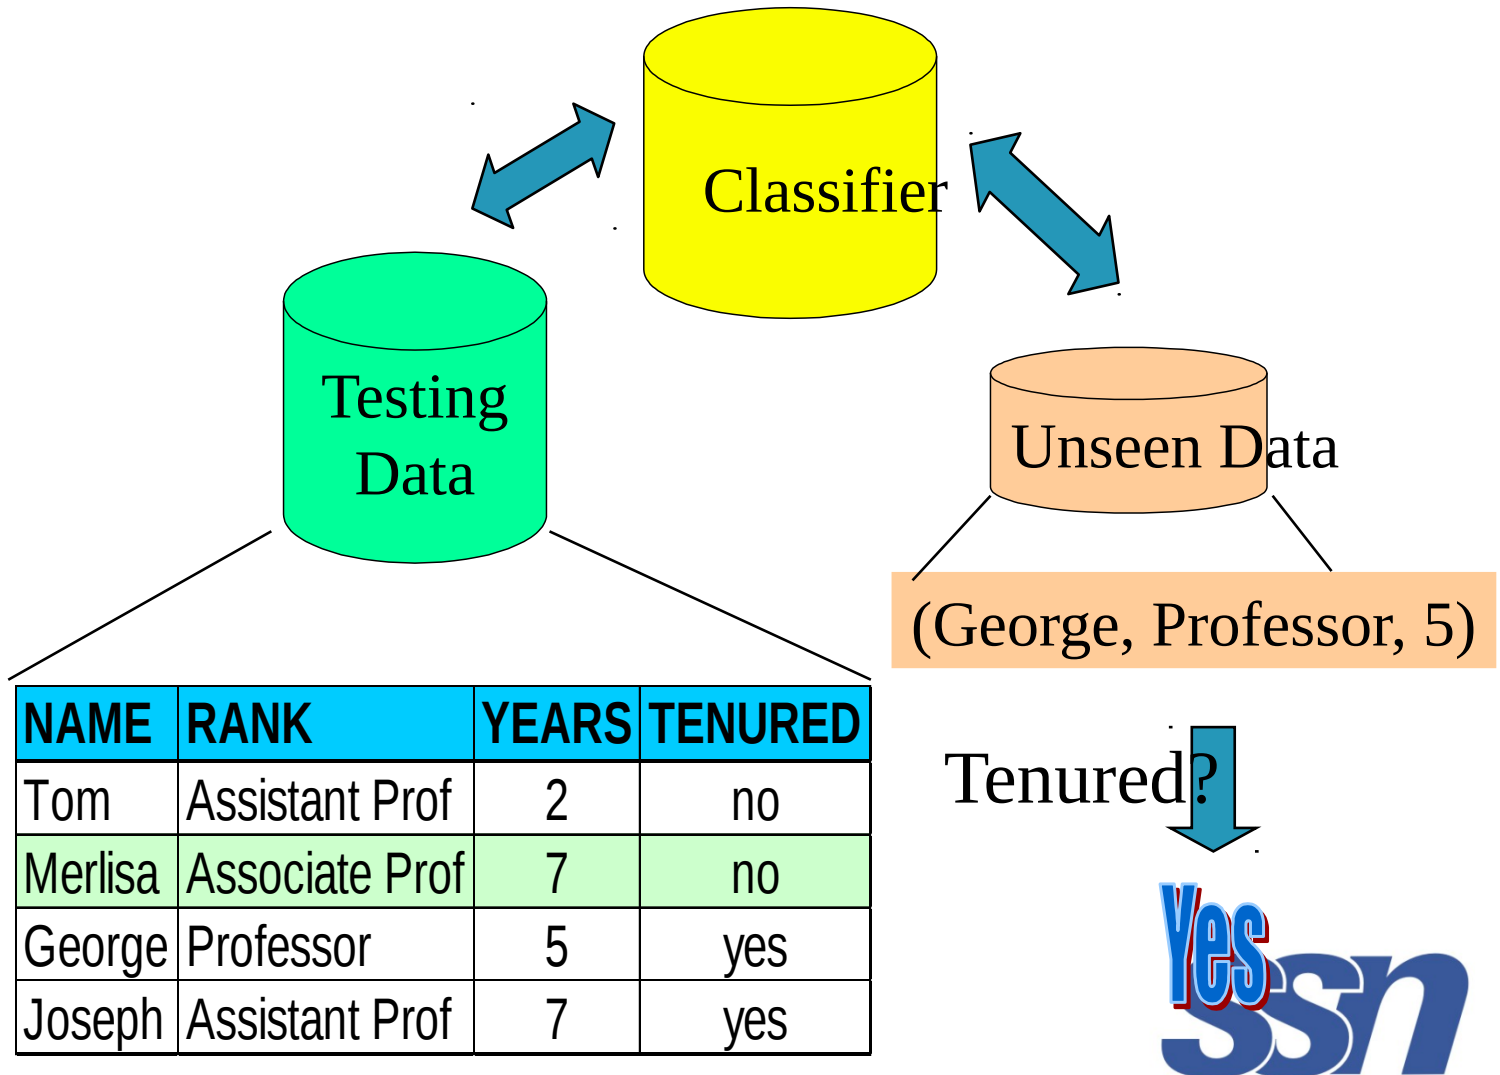
Prediction

- **Numeric Prediction**
 - Eg: The marketing manager wants to predict how much a given customer will spend.
 - The model constructed (predictor) predicts a continuous-valued function or ordered value.
 - The data analysis task is called as numeric prediction.
 - Regression analysis is a statistical methodology often used for numeric prediction

Process (1): Learning Step



Process (2): Classification Step



Decision Tree Induction

- Decision tree induction is the learning of decision trees from class labelled training tuples.

- Decision tree algorithm known as ID3

- Decision tree is a flowchart-like structure ,

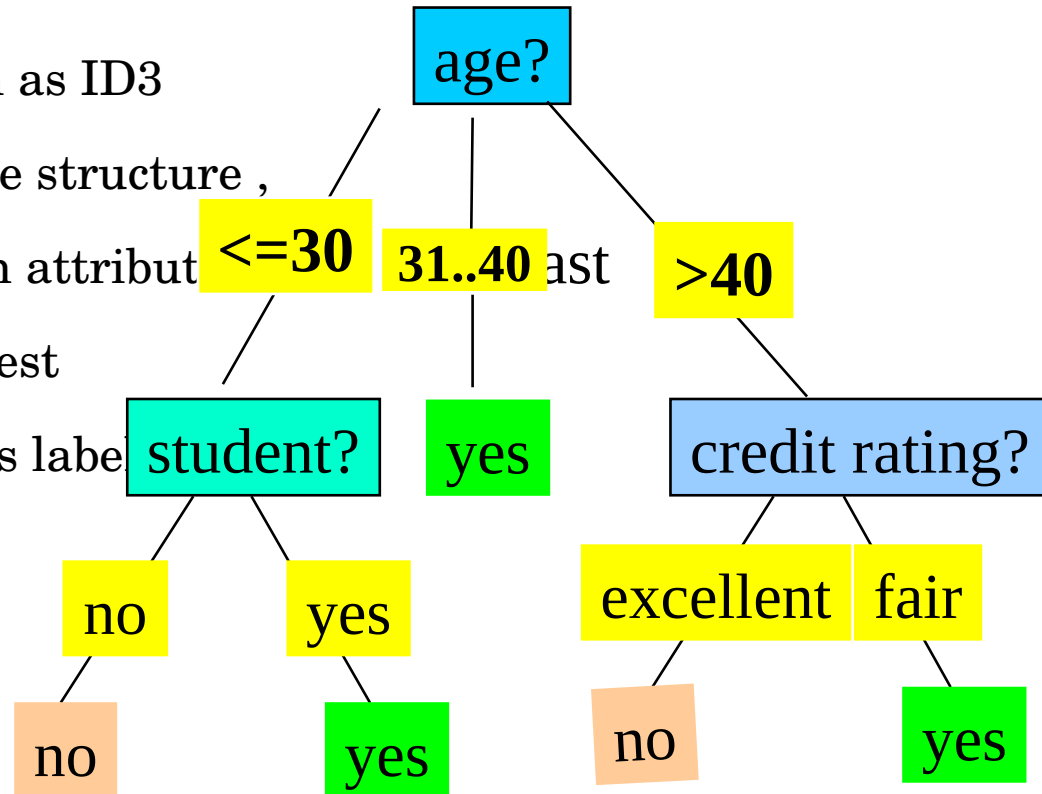
- Internal node=> test on an attribute
- Branch=> outcome of the test
- Leaf node=> holds the class label

- Eg: a decision tree for concept

buys_computer, internal node

represents a test on attribute

and each leaf node represents a class



Algorithm for Decision Tree Induction

- **How are decision trees used for classification?**
 - Given a tuple X for which the associated class label is unknown
 - The attribute values of the tuple are tested against the decision tree
 - A path is traced from the root to the leaf node which holds the class prediction of X
 - Decision trees are easily converted to classification rules.
- **Why decision trees?**
 - Does not require domain knowledge
 - Handles multidimensional data
 - Learning and classification steps are simple and fast
 - Easy to understand and attains good accuracy



Decision Tree Induction

- **Basic algorithm (a greedy algorithm)**
 - Tree is constructed in a top-down recursive divide-and-conquer manner
 - Starts with a training set of tuples and their associated class labels.
 - Training set is recursively partitioned into smaller subsets as the tree is being built.
 - Attribute selection measures are used to select the attribute that best partitions the tuples into distinct classes.
 - The tree nodes are created and the partition is labeled with the splitting criterion, branches are grown out for each outcome of the criteria



Attribute Selection Measures

- Attribute selection measure is a heuristic for selecting the splitting criterion that best separates a given data partition data D of training tuples into individual classes.
- Attribute selection measure are called as “**splitting rules**” because they determine how the tuples at a given node are to be split.
- Three popular attribute selection measures are:
 - Information gain
 - Gain Ratio
 - Gini Index

Attribute Selection Measure: Information Gain (ID3/C4.5)

- ❑ Select the attribute with the highest information gain
- ❑ Let p_i be the probability that an arbitrary tuple in D belongs to class C_i , estimated by $|C_{i,D}|/|D|$
- ❑ Expected information (entropy) needed to classify a tuple in D :

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$

- ❑ Partition the tuples in D on some attribute A having v distinct values $\{a_1, a_2, a_3 \dots a_v\}$ observed from the training set.
- ❑ Attribute A can be used to split D into v partitions and the partitions corresponds to the branches of the node N
- ❑ Information needed to do exact classification :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

$Gain(A) = Info(D) - Info_A(D)$

- ❑ Information gain is defined as the difference between original and new requirement
- $$Gain(A) = Info(D) - Info_A(D)$$



Attribute Selection: Information Gain

Class P: buys_computer = "yes"

Class N: buys_computer = "no"

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

age	p _i	n _i	I(p _i , n _i)
<=30	2	3	0.971
31...40	4	0	0
>40	3	2	0.971

$\frac{5}{14} I(2,3)$ means "age <=30" has 5 out of 14 samples, with 2 yes'es and 3 no's.

Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly:


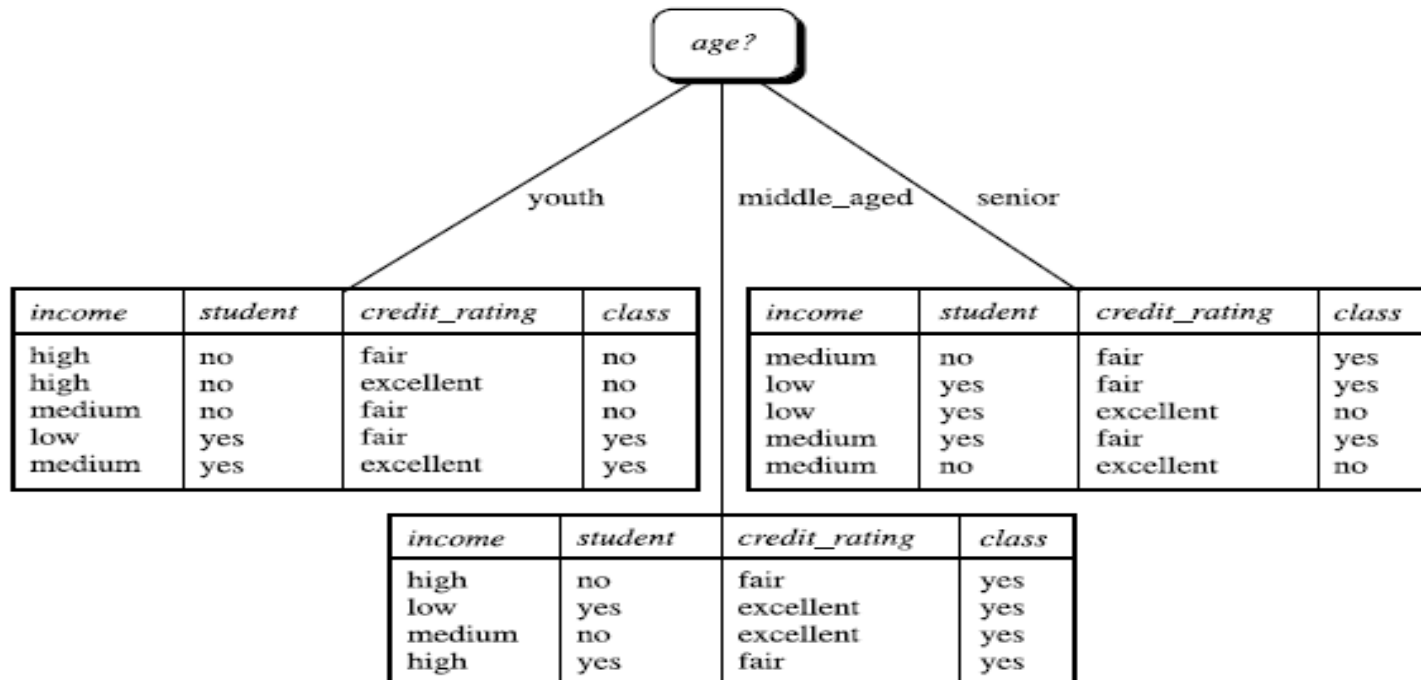
$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit_{rating}) = 0.048$$

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Attribute selection-Information gain



Attribute Selection: Information Gain

The mutual information is $I(S_{\text{Yes}}, S_{\text{No}}) = I(2,3) = -2/5 \log_2(2/5) - 3/5 \log_2(3/5) = 0.97$

- For Income we have three values $\text{income}_{\text{high}}$ (0 yes and 2 no), $\text{income}_{\text{medium}}$ (1 yes and 1 no) and $\text{income}_{\text{low}}$ (1 yes and 0 no)

$$\begin{aligned}\text{Entropy}(\text{income}) &= 2/5(0) + 2/5 (-1/2\log(1/2) - 1/2\log(1/2)) + 1/5 (0) \\ &= 2/5 (1) = 0.4\end{aligned}$$

$$\text{Gain}(\text{income}) = 0.97 - 0.4 = 0.57$$

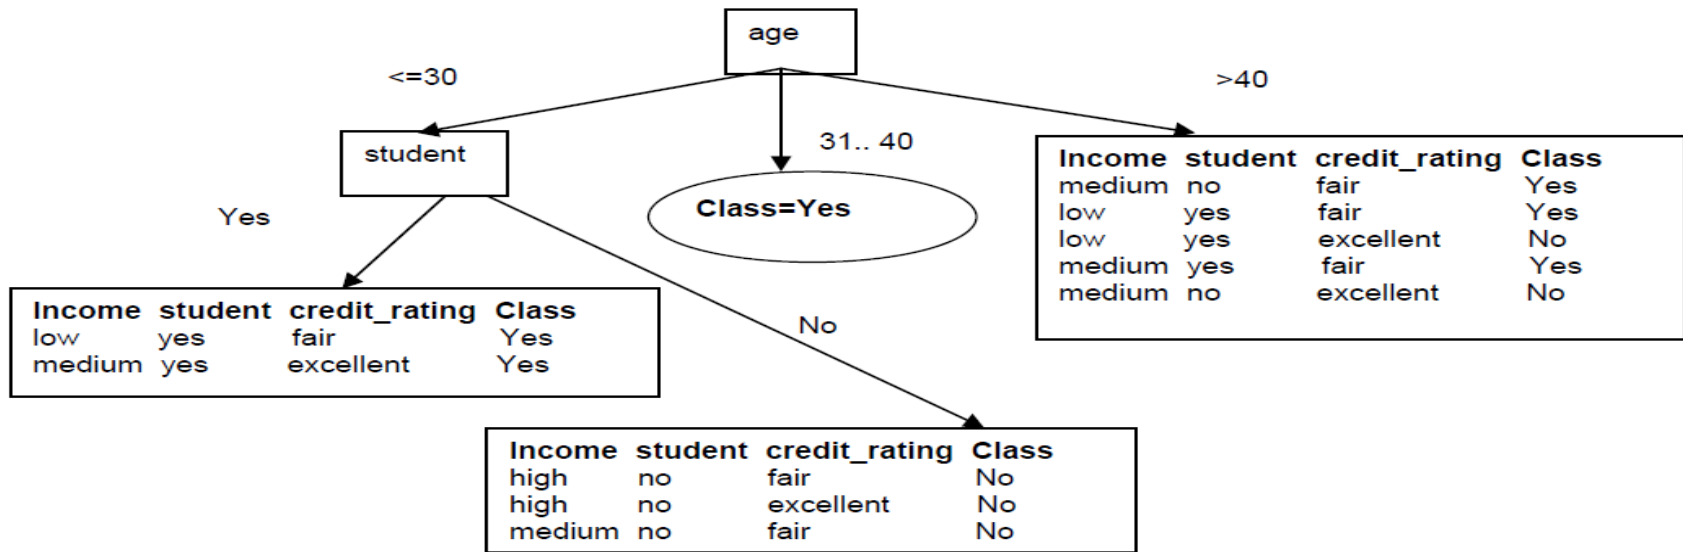
- For Student we have two values $\text{student}_{\text{yes}}$ (2 yes and 0 no) and $\text{student}_{\text{no}}$ (0 yes 3 no)

$$\text{Entropy}(\text{student}) = 2/5(0) + 3/5(0) = 0$$

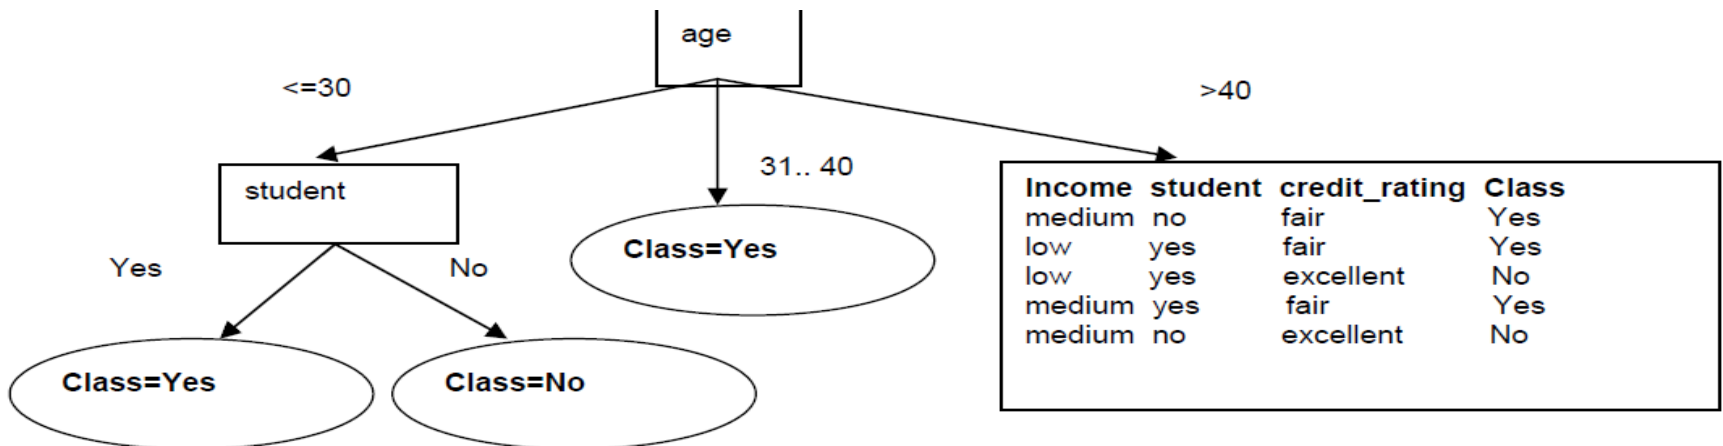
$$\text{Gain}(\text{student}) = 0.97 - 0 = 0.97$$

We can then safely split on attribute student without checking the other attributes since the information gain is maximized.

Attribute Selection: Information Gain



Since these two new branches are from distinct classes, we make them into leaf nodes with their respective class as label:



Attribute Selection: Information Gain

Again the same process is needed for the other branch of age.

The mutual information is $I(S_{\text{Yes}}, S_{\text{No}}) = I(3,2) = -3/5 \log_2(3/5) - 2/5 \log_2(2/5) = 0.97$

- For Income we have two values $\text{income}_{\text{medium}}$ (2 yes and 1 no) and $\text{income}_{\text{low}}$ (1 yes and 1 no)

$$\begin{aligned}\text{Entropy}(\text{income}) &= 3/5(-2/3\log(2/3)-1/3\log(1/3)) + 2/5(-1/2\log(1/2)-1/2\log(1/2)) \\ &= 3/5(0.9182)+2/5(1) = 0.55+0.4 = 0.95\end{aligned}$$

$$\text{Gain}(\text{income}) = 0.97 - 0.95 = 0.02$$

- For Student we have two values $\text{student}_{\text{yes}}$ (2 yes and 1 no) and $\text{student}_{\text{no}}$ (1 yes and 1 no)

$$\text{Entropy}(\text{student}) = 3/5(-2/3\log(2/3)-1/3\log(1/3)) + 2/5(-1/2\log(1/2)-1/2\log(1/2)) = 0.95$$

$$\text{Gain}(\text{student}) = 0.97 - 0.95 = 0.02$$

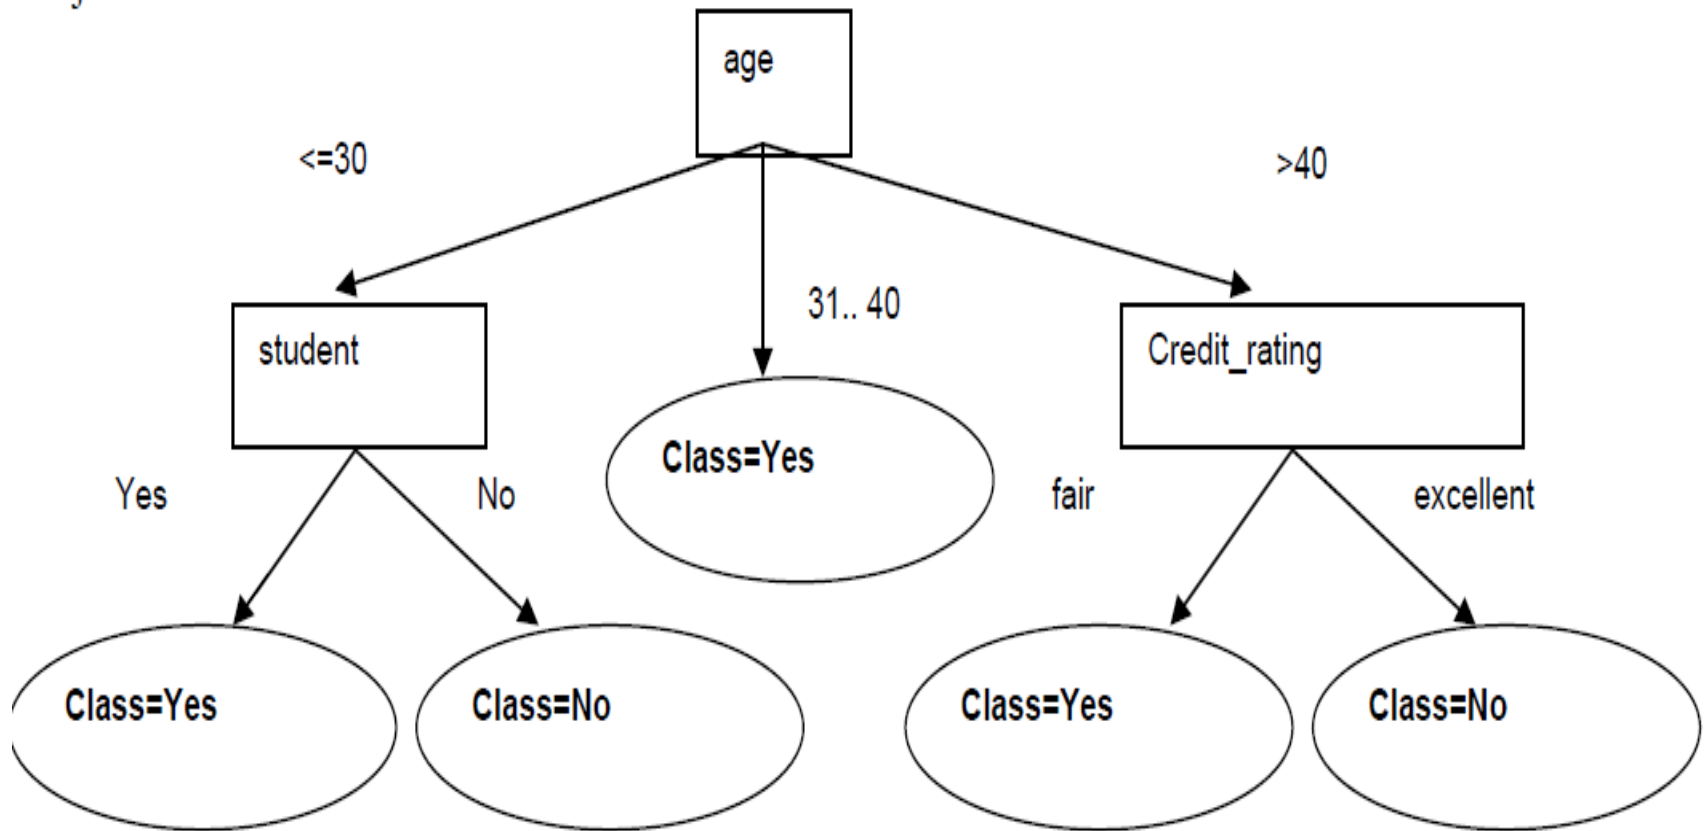
- For Credit_Rating we have two values $\text{credit_rating}_{\text{fair}}$ (3 yes and 0 no) and $\text{credit_rating}_{\text{excellent}}$ (0 yes and 2 no)

$$\text{Entropy}(\text{credit_rating}) = 0$$

$$\text{Gain}(\text{credit_rating}) = 0.97 - 0 = 0.97$$

We then split based on credit_rating. These splits give partitions each with records from the same class. We just need to make these into leaf nodes with their class label attached:

Output



New example: age<=30, income=medium, student=yes, credit-rating=fair

Follow branch(age<=30) then student=yes we predict Class=yes → Buys_computer = yes

Computing Information-Gain for Continuous-Valued Attributes

- Let attribute A be a continuous-valued attribute
- Must determine the *best split point* for A
 - Sort the value A in increasing order
 - Typically, the midpoint between each pair of adjacent values is considered as a possible *split point*
 - $(a_i + a_{i+1})/2$ is the midpoint between the values of a_i and a_{i+1}
 - The point with the *minimum expected information requirement* for A is selected as the split-point for A
- Split:
 - D1 is the set of tuples in D satisfying $A \leq \text{split-point}$, and D2 is the set of tuples in D satisfying $A > \text{split-point}$

Algorithm for Decision Tree Induction

- Algorithm is called with three parameters D , `attribute_list` and `attribute_selection_method`
- $D \Rightarrow$ complete set of training tuples and associated class labels
- `attribute_list` \Rightarrow list of attribute describing the tuples.
- `attribute_selection_method`:
 - heuristics procedure to determine the splitting criterion.
 - Determines which attribute to test at node N by determining the “best” way to separate the tuples D into classes
 - Determines which branch to grow from node N w.r to outcome of the chosen test.
- Attribute selection measure: Information gain and Gini index



Generation of Decision Tree

Algorithm: Generate_decision_tree. Generate a decision tree from the training tuples of data partition, D .

Input:

- Data partition, D , which is a set of training tuples and their associated class labels;
- *attribute_list*, the set of candidate attributes;
- *Attribute_selection_method*, a procedure to determine the splitting criterion that “best” partitions the data tuples into individual classes. This criterion consists of a *splitting_attribute* and, possibly, either a *split-point* or *splitting_subset*.

Output: A decision tree.

Method:

- (1) create a node N ;
- (2) if tuples in D are all of the same class, C , then
- (3) return N as a leaf node labeled with the class C ;
- (4) if *attribute_list* is empty then
- (5) return N as a leaf node labeled with the majority class in D ; // majority voting
- (6) apply *Attribute_selection_method*(D , *attribute_list*) to find the “best” *splitting_criterion*;
- (7) label node N with *splitting_criterion*;
- (8) if *splitting_attribute* is discrete-valued and
 multiway splits allowed then // not restricted to binary trees
- (9) *attribute_list* \leftarrow *attribute_list* $-$ *splitting_attribute*; // remove *splitting_attribute*
- (10) for each outcome j of *splitting_criterion*
 // partition the tuples and grow subtrees for each partition
- (11) let D_j be the set of data tuples in D satisfying outcome j ; // a partition
- (12) if D_j is empty then
- (13) attach a leaf labeled with the majority class in D to node N ;
- (14) else attach the node returned by *Generate_decision_tree*(D_j , *attribute_list*) to node N ;
- endfor
- (15) return N ;

Algorithm for Decision Tree Induction

- Conditions for stopping partitioning
 - All samples for a given node belong to the same class
 - There are no remaining attributes for further partitioning
 - majority voting** is employed for classifying the leaf
 - There are no samples left

Gain Ratio for Attribute Selection (C4.5)

- Information gain measure is biased towards attributes with a large number of values
- C4.5 (a successor of ID3) uses gain ratio to overcome the problem (normalization to information gain)

$$SplitInfo_A(D) = - \sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2 \left(\frac{|D_j|}{|D|} \right)$$

$$\text{GainRatio}(A) = \text{Gain}(A) / \text{SplitInfo}(A)$$

$$SplitInfo_{income}(D) = -\frac{4}{14} \times \log_2 \left(\frac{4}{14} \right) - \frac{6}{14} \times \log_2 \left(\frac{6}{14} \right) - \frac{4}{14} \times \log_2 \left(\frac{4}{14} \right) = 1.557$$

$$\text{gain_ratio}(\text{income}) = 0.029 / 1.557 = 0.019$$

- The attribute with the maximum gain ratio is selected as the splitting attribute

Gini Index (CART, IBM IntelligentMiner)

- If a data set D contains examples from n classes, gini index, $gini(D)$ is defined as

$$gini(D) = 1 - \sum_{j=1}^n p_j^2$$

where p_j is the probability that a tuple in D belongs to class C_j

- If a data set D is split on A into two subsets D_1 and D_2 , the gini index $gini_A(D)$ is defined as

$$gini_A(D) = \frac{|D_1|}{|D|} gini(D_1) + \frac{|D_2|}{|D|} gini(D_2)$$

- The subset with minimum gini index is selected as splitting subset
- Reduction in Impurity:

$$\Delta gini(A) = gini(D) - gini_A(D)$$

Computation of Gini Index

- Ex. D has 9 tuples in buys_computer = “yes” and 5 in “no”

$$gini(D) = 1 - \left(\frac{9}{14}\right)^2 - \left(\frac{5}{14}\right)^2 = 0.459$$

- Suppose the attribute income partitions D into 10 in D_1 : {low, medium} and 4 in D_2

$$gini_{income \in \{low, medium\}}(D) = \left(\frac{10}{14}\right) Gini(D_1) + \left(\frac{4}{14}\right) Gini(D_2)$$

$$= \frac{10}{14} \left(1 - \left(\frac{7}{10}\right)^2 - \left(\frac{3}{10}\right)^2\right) + \frac{4}{14} \left(1 - \left(\frac{2}{4}\right)^2 - \left(\frac{2}{4}\right)^2\right)$$

$$= 0.443$$

$$= Gini_{income \in \{high\}}(D).$$

$Gini_{\{low, high\}}$ is 0.458; $Gini_{\{medium, high\}}$ is 0.450. Thus, split on the {low, medium} (and {high}) since it has the lowest Gini index

Comparing Attribute Selection Measures

- The three measures, in general, return good results but
 - **Information gain:**
 - biased towards multivalued attributes
 - **Gain ratio:**
 - tends to prefer unbalanced splits in which one partition is much smaller than the others
 - **Gini index:**
 - biased to multivalued attributes
 - has difficulty when # of classes is large
 - tends to favor tests that result in equal-sized partitions and purity in both partitions

Overfitting and Tree Pruning

- Overfitting: An induced tree may overfit the training data
 - Too many branches, some may reflect anomalies due to noise or outliers
 - Poor accuracy for unseen samples
- Two approaches to avoid overfitting
 - Prepruning: *Halt tree construction early*- do not split a node if this would result in the goodness measure falling below a threshold (Difficult to choose best threshold)
 - Postpruning: *Remove branches* from a “fully grown” tree—get a sequence of progressively pruned trees
 - Use a set of data different from the training data to decide which is the “best pruned tree”



Classification in Large Databases

- Classification—a classical problem extensively studied by statisticians and machine learning researchers
- Scalability: Classifying data sets with millions of examples and hundreds of attributes with reasonable speed
- Why is decision tree induction popular?
 - relatively faster learning speed (than other classification methods)
 - convertible to simple and easy to understand classification rules
 - can use SQL queries for accessing databases
 - comparable classification accuracy with other methods

