


Process Management

Reference: Pradeep K Sinha,
"Distributed Operating Systems: Concepts and
Design", Prentice Hall of India, 2007



Overview

- Process Migration
 - Features
 - Mechanisms




Introduction

- In conventional (centralized) operating system, process management deals with mechanisms and policies for **sharing the processor of the system among all processes**.
- In a **distributed operating system**, the main goal of process management is to make the best possible use of the **processing resources of the entire system by sharing them among all processes**.
- **Three important concepts are used to achieve this goal:**
 - Processor allocation
 - Process migration
 - Threads




Introduction

- **Process allocation** deals with the mechanism of deciding which process should be assigned to which processor.
- **Process migration** deals with the movement of a process from its current location to the processor to which it has been assigned.
- **Threads** deals with fine-grained parallelism for better utilization of the processing capability of the system.

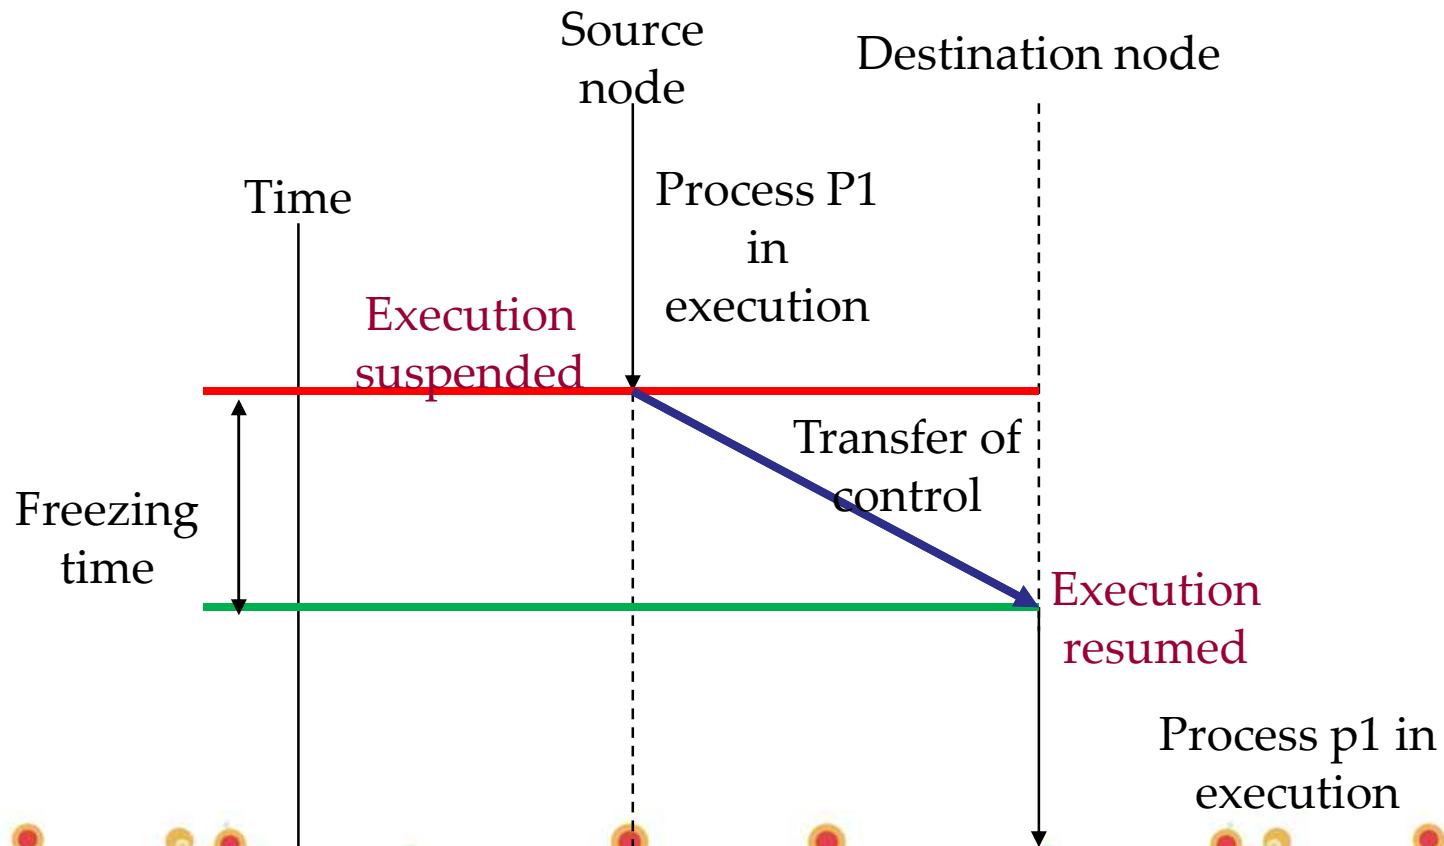


Process Migration

- Relocation of a process from its current location to another node.
- Process may be migrated
 - either before it start executing on its source node
 - Known as **non-preemptive process migration**.
 - or during the course of its execution
 - Known as **preemptive process migration**.



Flow of execution of a migration process



Process Migration

- Preemptive process migration is costlier.
- **Process Migration Policy**
 - Selection of a process that should be migrated.
 - Selection of the destination node to which the selected process should be migrated.
- **Process Migration Mechanism**
 - Actual transfer of the selected process to the destination node.



Desirable features of a good process migration mechanism


- Transparency
- Minimal Interference
- Minimal Residual Dependencies
- Efficiency
- Robustness
- Communication between co-processes of a job



Transparency

Level of transparency:

- Object Access Level
- System Call & Interprocess Communication level



Object Access level Transparency

- Minimum requirement for a system to support **non-preemptive process migration** facility.
- Access to objects such as **files** and **devices** can be **location independent**.
- Allows free initiation of **program** at **arbitrary node**.
- Requires **transparent** object **naming** and **locating**.



System Call & IPC level

- For migrated process, **system calls** should be **location independent**.
- For **transparent redirection of messages** during the **transient** state of a process.
- Transparency must be provided to support **preemptive process migration** facility.




Minimal Interference

- Migration of a process should cause **minimal interference** of **progress** of the **process** involved.
- Achieve by **minimizing** the **freezing time** of the process being migrated.
 - Freezing time is the time period for which the **execution of the process is stopped** for transferring its **information** to the **destination node**.



Minimal Residual Dependencies

- No residual dependency should be left on previous node.




Efficiency

- **Minimum time** required for migrating a process.
- **Minimum cost** of locating an object.
- **Minimum cost** of supporting **remote execution** once the process is migrated.




Robustness

- The failure of a node other than the one on which a process is currently running should not in any way affect the accessibility or execution of that process.



Communication between coprocessors of a job

- To **reduce communication cost**, it is necessary that **coprocesses** are able to **directly communicate** with each other **irrespective** of their **locations**.



Process Migration Mechanisms

Four major activities

- Freezing the process on its source node and restarting it on its destination node.
- Transferring the process's address space from its source node to its destination node.
- Forwarding messages meant for the migrant process.
- Handling communication between cooperating processes that have been separated as a result of process migration.



Mechanisms for Freezing and Restarting a Process

- **Freezing the process:**
 - The **execution** of the **process** is **suspended** and all **external interactions** with the **process** are **deferred**.
- **Issues:**
 - Immediate and delayed blocking of the process
 - Fast and slow I/O operations
 - Information about open files
 - Reinstating the process on its Destination node



Immediate and delayed blocking of the process

- If the process is **not executing a system call**, it can be **immediately blocked** from further execution.
- If the process is **executing a system call** but is sleeping at an **interruptible priority** waiting for a kernel event to occur, it can be **immediately blocked** from further execution.
- If the process is **executing a system call** but is sleeping at an **non-interruptible priority** waiting for a kernel event to occur, it **cannot be blocked immediately**.



Fast and Slow I/O Operations

- Process is frozen after the completion of all fast I/O operations like disk access.
- Slow I/O operation (pipe or terminal) is done after process migration and when process is executed on destination node.



Information about open files

- Includes **name** and **identifier** of the **file**, their **access modes** and the **current positions** of their **file pointers**.
- OS returns a file descriptor to process that is used for all I/O.
- It is necessary to somehow **preserve** a **pointer** to the **file** so that migrated process could continue to access it.



Information about open files

- **Approaches :**


- **Link** is created to the **file and the pathname** of the link is used as an **access point** to the file after the process migrates.
- An open file's **complete pathname** is **reconstructed** when required by modifying the kernel.

- Keeping **Track of file replicas.**



Reinstating the process on its Destination Node

- On the **destination node**, an **empty process state** is created.
- **Newly allocated process** may or may not have the **same process identifier** as the **migrating process**.



Reinstating the process on its Destination Node


- Once all the state of the migrating process has been transferred from the source to destination node and copied into the empty state, new copy of the process is unfrozen and old copy is deleted.
- The process is restarted on its destination node in whatever state it was in before being migrated.



Address Space Transfer Mechanisms

- The migration of a process involves the transfer of
 - Process's state
 - Process's address space

from the source node to the destination node.



Address Space Transfer Mechanisms

- **Process State consists of**
 - Execution Status – Register Contents
 - Memory Tables
 - I/O State : I/O Queue, I/O buffers, Interrupts
 - Capability list
 - Process's Identifier
 - Process's user and group identifier
 - Information about Open Files




Address Space Transfer Mechanisms

- **Process address space**

- code,
- data and
- program stack

- The size of the process's address space (several megabytes) overshadows the size of the process's state information (few kilobytes).



Address Space Transfer Mechanisms

- Mechanisms for address space transfer:
 - Total freezing
 - Pretransferring
 - Transfer on reference

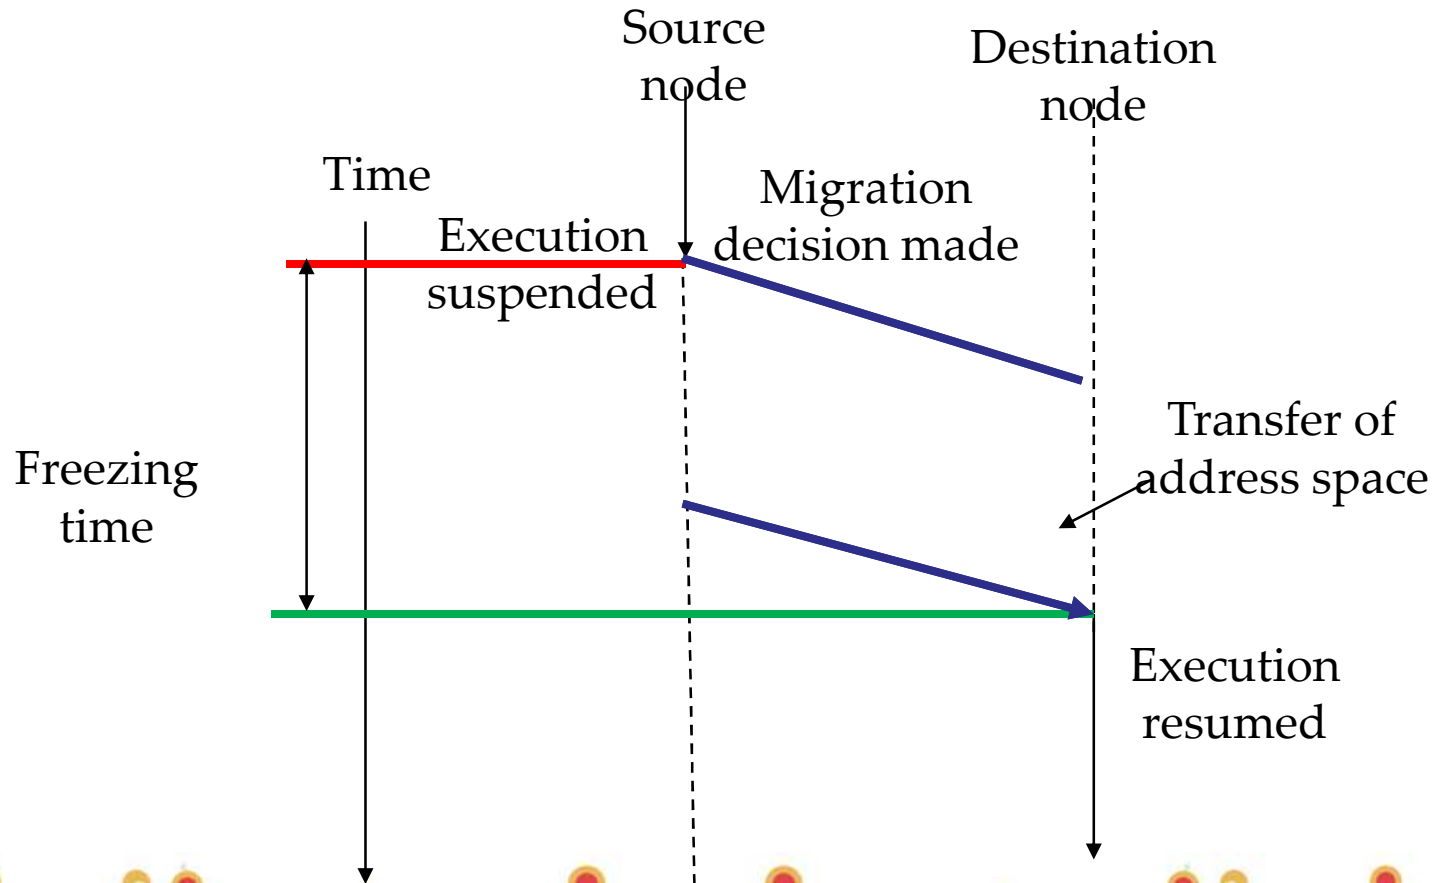


Total Freezing

- A process's **execution** is **stopped** while its address space is being transferred.
- **Disadvantage:**
 - Process is **suspended** for **long time** during migration, timeouts may occur, and if process is interactive, the delay will be noticed by the user.



Total Freezing




Pretransferring

- Also known as **precopying**.
- The address space is transferred while the **process** is **still running** on the source node.
- It is done as an initial transfer of the complete address space followed by **repeated transfers** of the **page modified** during the **previous transfer**.




Pretransferring

- The **pretransfer** operation is executed at a **higher priority** than all other programs on the source node.
- **Reduces the freezing time** of the process but it may increase the **total time for migrating** due to the possibility of **redundant page transfers**.



Pretransferring



Transfer on Reference

- The process address space is left behind on its source node, and as the relocated process executes on its destination node.
- Attempts to reference memory page results in the generation of requests to copy in the desired blocks from their remote location.
- A page is transferred from its source node to its destination node only when referenced.

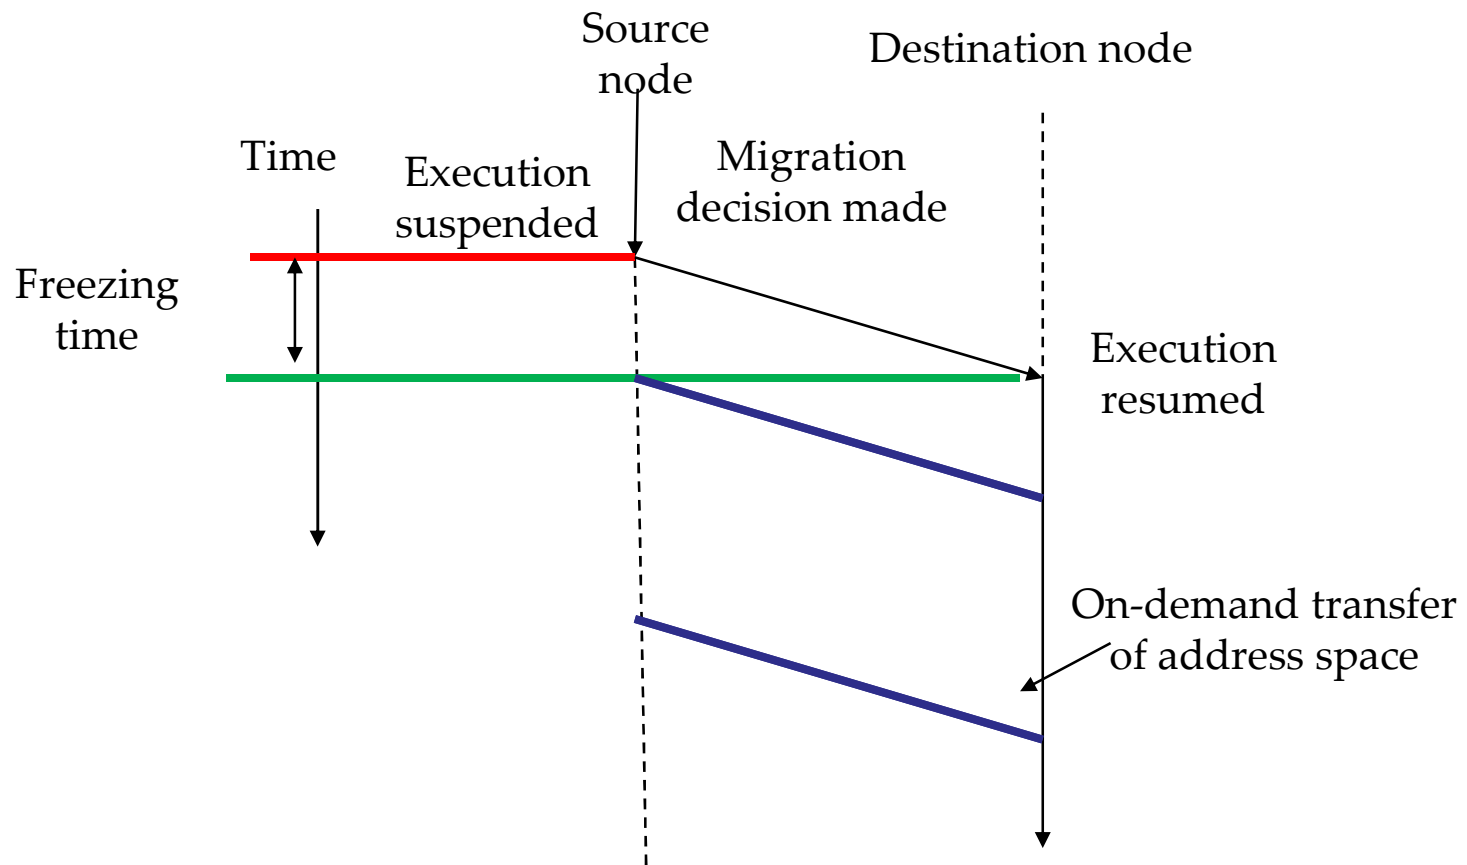


Transfer on Reference

- Very short **switching time** of the process from its source node to its destination node.
- Imposes a continued **load on the process's source node** and **results** in the **process** if **source node fails** or is **rebooted**.



Transfer on Reference



Message-forwarding Mechanisms

- In **moving** a **process**, it must be ensured that all **pending**, **re-route**, and **future messages** arrive at the **process's new location**.
- **Types of messages:**
 1. **Messages received** at the **source node** after the **process's execution** has been **stopped** on its **source node** and the **process's execution** has not yet been started on its **destination node**.
 2. **Messages received** at the **source node** after the **process's execution** has started on its **destination node**.
 3. **Messages** that are to be sent to the **migrant process** from any **other node** after it has started **executing** on the **destination node**.




Message-forwarding Mechanisms

- Mechanisms:
 - Mechanism of resending the message
 - Origin site mechanism
 - Link traversal mechanism
 - Link update mechanism



Mechanisms of resending the message

- Messages of the **types 1 and 2 are returned** to the **sender** as **not deliverable** or are simply **dropped**, with the assurance that the sender of the message is storing a copy of the data and is prepared to retransmit it.
- **Does not** require any **process state** to be **left** behind on the process's **source** node.




Mechanisms of resending the message

- Disadvantage:
 - The message forwarding mechanism of process migration operation is non-transparent to the processes interacting with the migrant process.




Origin Site Mechanism

- Each site is responsible for keeping information about the current locations of all the processes created on it.
- Messages for a particular process are always first sent to its origin site.
- The origin site then forwards the message to the process's current location.



Origin Site Mechanism

- Disadvantage:
 - Failure of the origin site will disrupt the message forwarding mechanism.
 - Continuous load on the migrant process's origin site even after the process has migrated from that node.



Link Traversal Mechanism

- To redirect message **type 1**, a **message queue** for the **migrant process** is created on its **source** node.
- For **2 and 3 type message**, a forwarding **address** known as **link** is **left at the source node** pointing to the **destination** of the **migrant process**.
- Two component of link, one is **unique process identifier** and second is **last known location**.



Link Traversal Mechanism

- Disadvantage:
 - Several link may have to be traversed to locate a process from a node and if any node in chain of link fails, the process cannot be located.



Link Update Mechanisms

- During the transfer phase of the migrant process, the **source node sends link-update messages** to the kernels controlling all of the **migrant process's communication partners**.
- Link update message
 - Tells the **new address of each link** held by the **migrant process**.
 - **Acknowledged** for **synchronization** purposes.



Mechanisms for Handling Coprocesses

- To provide efficient **communication** between a **process** and its **sub processes** which might have been **migrated** on **different nodes**.
- **Mechanisms :**
 - Disallowing separation of coprocesses.
 - Home node or origin site concept.




Disallowing Separation of coprocesses

- Easiest method of handling communication between coprocesses is to **disallow their separation**.
- **Methods :**
 - By **disallowing** the **migration** of **processes** that **wait** for **one** or **more** of their **children** to **complete**.
 - By ensuring that when a **parent process migrates**, its **children process** will be **migrated along** with it.




Disallowing Separation of coprocesses

- Disadvantage:
 - It does **not allow** the use of **parallelism** within **jobs**.



Home node or Origin Sites Concept

- Used for **communication** between a **process** and its **sub process** when the **two** are **running** on **different nodes**.
- Allows the **complete freedom** of **migrating** a **process** or its **sub process independently** and **executing** them on **different nodes** of the system.




Home node or Origin Sites Concept

- Disadvantage:
 - All communication between a parent process and its children processes take place via the home node.
 - The message traffic and the communication cost increase considerably.

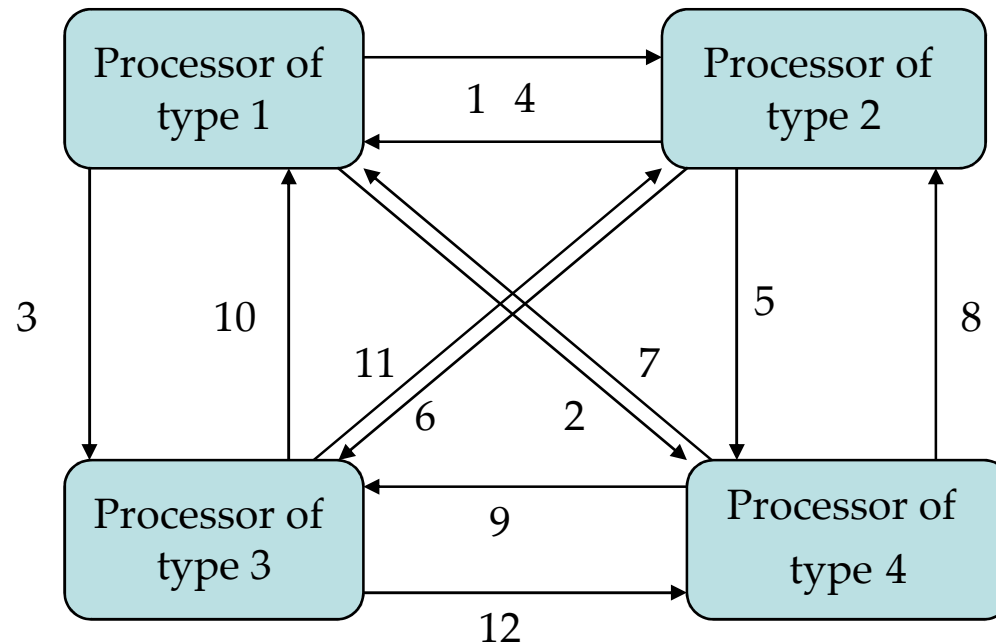


Process Migration in Heterogeneous Systems

- All the concerned **data** must be **translated** from the **source CPU format** to the **destination CPU format** before it can be executed on the destination node.
- A heterogeneous system having n CPU types must have $n(n-1)$ pieces of translation software.
- Handles problem of **different data representations** such as characters, integers and floating-point numbers.

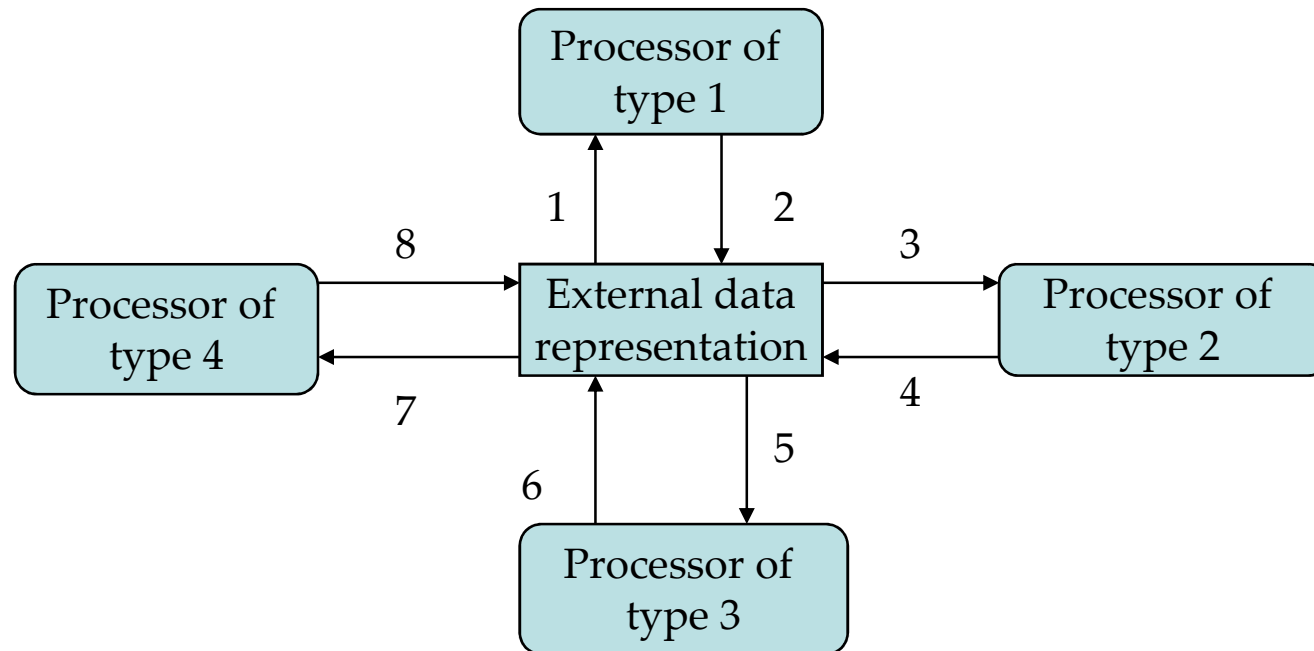


Process Migration in Heterogeneous Systems



Example: The need for 12 pieces of translation software required in a heterogeneous system having 4 types of processors

Process Migration in Heterogeneous Systems



Example: The need for only 8 pieces of translation software in a heterogeneous system having 4 types of processors when the EDR mechanism is used.

Handling Exponent

- The **number of bits** for the **exponent** varies from **processor to processor**.
- The EDR have at least as many **bits** in the exponent as the **largest exponent**.
- **Overflow** or **underflow** can be managed upon conversion from the **external data representation**.



Handling Exponent

- Some solutions:
 - Ensuring that **numbers** used by the programs that migrate have a **smaller exponent** value **than** the **smallest processor's exponent** value in the system.
 - Emulating the **larger processor's value**.
 - Restricting the migration of the process to **only those nodes** whose **processor's exponent** representation is at **least as large as** that of the **source node's processor**.



Handling the Mantissa

- Suppose processor **A** uses **32 bits** and processor **B** uses **64 bits**.
 - No problem from A to B migration.
- **Problem:**
 - **Half-Precision** problem from B to A migration.
 - **Loss of precision** due to multiple migrations
- **Solution:**
 - The **external data representation** must have sufficient precision to handle the **largest mantissa**.



Advantages of Process Migration

- *Reducing average response time of processes*
 - To reduce the average response time of the processes, processes of a **heavily loaded node** are **migrated** to **idle** or **underutilized nodes**.
- *Speeding up individual jobs*
 - A migration of job to different node is done and execute them **concurrently**.
 - Migrate a job to a node having a **faster CPU** or to a node at which it has **minimum turnaround time**.
 - More speed up more migration cost involved.



Advantages of Process Migration

- *Gaining higher throughput*
 - Process migration facility may also be used properly to mix **I/O and CPU-bound processes** on a global basis for **increasing the throughput** of the system.
- *Utilizing resources effectively*
 - Depending upon the **nature of a process**, it can be migrated to **suitable node** to **utilize** the system **resource** in the most **efficient** manner.



Advantages of Process Migration

- *Reducing network traffic*
 - Migrating a process **closer** to the **resources** it is using most heavily.
- *Improving system reliability*
 - Migrate a **copy** of a **critical process** to some other node and to **execute both** the original and copied processes **concurrently** on **different nodes**.
- *Improving system security*
 - A **sensitive process** may be migrated and run on a **secure node**.



Thank You

