

INTERPROCESS COMMUNICATION

**Reference: George Coulouris, Jean Dollimore and Tim Kindberg,
“Distributed Systems Concepts and Design”, Fifth Edition, Pearson
Education, 2012**

Topics

- Introduction
- The API For The INTERNET PROTOCOLS
- External Data Representation
- Client-server Communication
- Group Communication

Introduction

- The java API for interprocess communication in the internet provides both datagram and stream communication.
- The two communication patterns that are most commonly used in distributed programs:
 - **Client-Server communication**
 - ❖ The request and reply messages provide the basis for remote method invocation (RMI) or remote procedure call (RPC).

Introduction

- Group communication
 - ❖ The same message is sent to several processes.

Introduction

- This chapter is concerned with middleware.

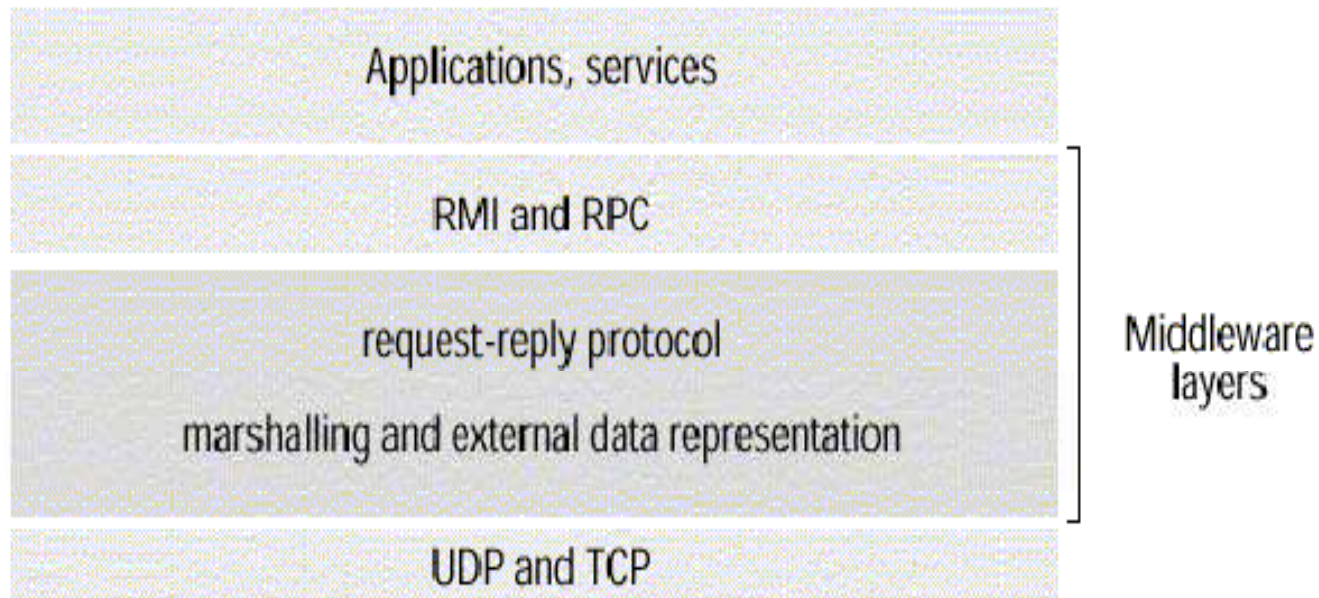


Figure 1. Middleware layers

Introduction

- **Remote Method Invocation (RMI)**
 - It allows an object to invoke a method in an object in a remote process.
 - ❖ E.g. CORBA and Java RMI
- **Remote Procedure Call (RPC)**
 - It allows a client to call a procedure in a remote server.

Introduction

- The application program interface (API) to UDP provides a **message passing** abstraction.
 - Message passing is the simplest form of interprocess communication.
 - API enables a sending process to transmit a single message to a receiving process.
 - The independent packets containing these messages are called **datagrams**.
 - In the Java and UNIX APIs, the sender specifies the destination using a **socket**.

Introduction

- **Socket** is an indirect reference to a particular port used by the destination process at a destination computer.
- The application program interface (API) to TCP provides the abstraction of a two-way stream between pairs of processes.
- The information communicated consists of a stream of data items with no message boundaries.

Introduction

- Request-reply protocols are designed to support client-server communication in the form of either RMI or RPC.
- Group multicast protocols are designed to support group communication.

Introduction

- **Group multicast** is a form of interprocess communication in which one process in a group of processes transmits the same message to all members of the group.

The API for the Internet Protocols

- The CHARACTERISTICS of INTERPROCESS COMMUNICATION
- SOCKET
- UDP DATAGRAM COMMUNICATION
- TCP STREAM COMMUNICATION

The Characteristics of Interprocess Communication

- Synchronous and asynchronous communication
 - In the synchronous form, both send and receive are blocking operations.
 - In the asynchronous form, the use of the send operation is non-blocking and the receive operation can have blocking and non-blocking variants.

The Characteristics of Interprocess Communication

- Message destinations
 - A local port is a message destination within a computer, specified as an integer.
 - A port has an exactly one receiver but can have many senders.

The Characteristics of Interprocess Communication

- **Reliability**
 - A reliable communication is defined in terms of validity and integrity.
 - A point-to-point message service is described as reliable if messages are guaranteed to be delivered despite a reasonable number of packets being dropped or lost.
 - For integrity, messages must arrive uncorrupted and without duplication.

The Characteristics of Interprocess Communication

- Ordering
 - Some applications require that messages be delivered in sender order.

Sockets

- Internet IPC mechanism of Unix and other operating systems (BSD Unix, Solaris, Linux, Windows NT, Macintosh OS)
- Processes in the above OS can send and receive messages via a socket.
- Sockets need to be bound to a port number and an internet address in order to send and receive messages.
- Each socket has a transport protocol (TCP or UDP).

Sockets

- Messages sent to some internet address and port number can only be received by a process using a socket that is bound to this address and port number.
- Processes cannot share ports (exception: TCP multicast).
- 2^{16} of possible port numbers for local processes to receive messages.

Sockets

- Both forms of communication, UDP and TCP, use the socket abstraction, which provides an endpoint for communication between processes.
- Interprocess communication consists of transmitting a message between a socket in one process and a socket in another process.

(Figure 2)

Sockets




Figure 2. Sockets and ports

UDP Datagram Communication

- **UDP datagram properties**
 - No guarantee of order preservation
 - Message loss and duplications are possible
- **Necessary steps**
 - **Creating** a socket
 - **Binding** a socket to a port and local Internet address
 - ❖ A client binds to any free local port
 - ❖ A server binds to a server port

UDP Datagram Communication

- Receive method
 - It returns Internet address and port of sender, plus message to receipt to check from where message came.

UDP Datagram Communication

- Issues related to datagram communications are:
 - Message size
 - ❖ IP allows for messages of up to 216 bytes.
 - ❖ Most implementations restrict this to around 8 kbytes.
 - ❖ Any application requiring messages larger than the maximum must fragment.
 - ❖ If arriving message is too big for array allocated to receive message content, truncation occurs.

UDP Datagram Communication

➤ Blocking

❖ Send: non-blocking

- upon arrival, message is placed in a queue for the socket that is bound to the destination port.

❖ Receive: blocking

- Pre-emption by timeout possible
- If process wishes to continue while waiting for packet, use separate thread

UDP Datagram Communication

- Timeout
- Receive from any Origin

UDP Datagram Communication

- UDP datagrams suffer from following failures:
 - Omission failure
 - Messages may be dropped occasionally,
 - Ordering

Java API for UDP Datagrams

- The Java API provides datagram communication by two classes:

- DatagramPacket

- ❖ It provides a constructor to make an array of bytes comprising:
 - Message content
 - Length of message
 - Internet address
 - Local port number
 - ❖ It provides another similar constructor for receiving a message.

array of bytes containing message | length of message| Internet address | port number|

Java API for UDP Datagrams

➤ DatagramSocket

- ❖ This class supports sockets for sending and receiving UDP datagram.
- ❖ It provides a constructor with port number as argument.
- ❖ No-argument constructor is used to choose a free local port.
- ❖ DatagramSocket methods are:
 - `send` and `receive`
 - `setSoTimeout`
 - `connect`

Java API for UDP Datagrams

```
import java.net.*;
import java.io.*;
public class UDPClient{
public static void main(String args[]){
// args give message contents and destination hostname
try {
DatagramSocket aSocket = new DatagramSocket(); // create socket
byte [] m = args[0].getBytes();
InetAddress aHost = InetAddress.getByName(args[1]); // DNS lookup
int serverPort = 6789;
DatagramPacket request =
new DatagramPacket(m, args[0].length(), aHost, serverPort);
aSocket.send(request); //send message
byte[] buffer = new byte[1000];
DatagramPacket reply = new DatagramPacket(buffer, buffer.length);
aSocket.receive(reply); //wait for reply
System.out.println("Reply: " + new String(reply.getData()));
aSocket.close();
}catch (SocketException e){System.out.println("Socket: " + e.getMessage());
}catch (IOException e){System.out.println("IO: " + e.getMessage());}
} finally{if (aSocket !=null)aSocket.close()}
}
}
```

Figure 3. UDP client sends a message to the server and gets a reply

Couloris, Dollimore and Kindberg Distributed Systems: Concepts & Design Edn. 4, Pearson Education 2005

Java API for UDP datagrams

```
import java.net.*;
import java.io.*;
public class UDPServer{
    public static void main(String args[]){
        DatagramSocket aSocket = null;
        try {
            aSocket = new DatagramSocket(6789);
            byte []buffer = new byte[1000];
            While(true){
                DatagramPacket request =new DatagramPacket(buffer, buffer.length);
                aSocket.receive(request);
                DatagramPacket reply = new DatagramPacket(request.getData();
                    request.getLength(),request.getAddress(), request.getPort());
                aSocket.send(reply);
            }
        }catch (SocketException e){System.out.println("Socket: " + e.getMessage());}
        }catch (IOException e){System.out.println("IO: " + e.getMessage());}
        }finally{if (aSocket !=null)aSocket.close()}
    }
}
```

Figure 4. UDP server repeatedly receives a request and sends it back to the client

Java API for UDP Datagrams

Example

- The process creates a socket, sends a message to a server at port 6789 and waits to receive a reply.

Java API for UDP Datagrams

Example

- The process creates a socket, bound to its server port 6789 and waits to receive a request message from a client.

TCP Stream Communication

- The API to the TCP protocol provides the abstraction of a stream of bytes to be written to or read from.
 - Characteristics of the stream abstraction:
 - ❖ Message sizes
 - ❖ Lost messages
 - ❖ Flow control
 - ❖ Message Duplication and Ordering
 - ❖ Message destinations

TCP Stream Communication

- Issues related to stream communication:
 - Matching of data items
 - Blocking
 - Threads

TCP Stream Communication

■ Use of TCP

➤ Many services that run over TCP connections, with reserved port number are:

- ❖ HTTP (Hypertext Transfer Protocol)
- ❖ FTP (File Transfer Protocol)
- ❖ Telnet
- ❖ SMTP (Simple Mail Transfer Protocol)

TCP Stream Communication

■ Java API for TCP streams

➤ The Java interface to TCP streams is provided in the classes:

❖ ServerSocket

- It is used by a server to create a socket at server port to listen for connect requests from clients.

❖ Socket

- It is used by a pair of processes with a connection.
- The client uses a constructor to create a socket and connect it to the remote host and port of a server.
- It provides methods for accessing input and output streams associated with a socket.

Java API for UDP Datagrams

Example

- The client process creates a socket, bound to the hostname and server port 6789.

Java API for UDP Datagrams

Example

- The server process opens a server socket to its server port 6789 and listens for connect requests.

TCP Stream Communication

```
import java.net.*;
import java.io.*;
public class TCPServer {
    public static void main (String args[]) {
        try{
            int serverPort = 7896;
            ServerSocket listenSocket = new ServerSocket(serverPort);
            while(true) {
                Socket clientSocket = listenSocket.accept();
                Connection c = new Connection(clientSocket);
            }
        } catch(IOException e) {System.out.println("Listen socket:"+e.getMessage());}
    }
}
```

Figure 6. TCP server makes a connection for each client and then echoes the client's request

Couloris, Dollimore and Kindberg Distributed Systems: Concepts & Design Edn. 4, Pearson Education 2005

TCP Stream Communication

```
class Connection extends Thread {
    DataInputStream in;
    DataOutputStream out;
    Socket clientSocket;
    public Connection (Socket aClientSocket) {
        try {
            clientSocket = aClientSocket;
            in = new DataInputStream( clientSocket.getInputStream());
            out =new DataOutputStream( clientSocket.getOutputStream());
            this.start();
        } catch(IOException e){System.out.println("Connection:"+e.getMessage());}
    }
    public void run(){
        try {
            // an echo server
            String data = in.readUTF();
            out.writeUTF(data);
        } catch (EOFException e){System.out.println("EOF:"+e.getMessage());}
        } catch (IOException e) {System.out.println("readline:"+e.getMessage());}
    } finally {try{clientSocket.close();}catch(IOException e){/*close failed*/}}
    }
}
```

Figure 7. TCP server makes a connection for each client and then echoes the client's request

Couloris, Dollimore and Kindberg Distributed Systems: Concepts & Design Edn. 4, Pearson Education 2005

External Data Representation

- The information stored in running programs is represented as data structures, whereas the information in messages consists of sequences of bytes.
- Irrespective of the form of communication used, the data structure must be converted to a sequence of bytes before transmission and rebuilt on arrival.

External Data Representation

- External Data Representation is an agreed standard for the representation of data structures and primitive values.

- Data representation problems are:
 - Using agreed external representation, two conversions necessary
 - Using sender's or receiver's format and convert at the other end

External Data Representation

■ Marshalling

- Marshalling is the process of taking a collection of data items and assembling them into a form suitable for transmission of a message.

■ Unmarshalling

- Unmarshalling is the process of disassembling a collection of data on arrival to produce an equivalent collection of data items at the destination.

External Data Representation

- Three approaches to external data representation and marshalling are:
 - CORBA
 - Java's object serialization
 - XML

External Data Representation

- Marshalling and unmarshalling activities is usually performed automatically by middleware layer.
- Marshalling is likely error-prone if carried out by hand.

CORBA Common Data Representation (CDR)

- CORBA Common Data Representation (CDR)
 - CORBA CDR is the external data representation for the structured and primitive types.
 - It can be used by a variety of programming languages
 - It consists 15 primitive types:
 - Short (16 bit)
 - Long (32 bit)
 - Unsigned short
 - Unsigned long
 - Float(32 bit)
 - Double(64 bit)
 - Char
 - Boolean(TRUE,FALSE)
 - Octet(8 bit)
 - Any(can represent any basic or constructed type)
 - Composite type are shown in Figure 8.

CORBA Common Data Representation (CDR)

<i>Type</i>	<i>Representation</i>
<i>sequence</i>	length (unsigned long) followed by elements in order
<i>string</i>	length (unsigned long) followed by characters in order (can also can have wide characters)
<i>array</i>	array elements in order (no length specified because it is fixed)
<i>struct</i>	in the order of declaration of the components
<i>enumerated</i>	unsigned long (the values are specified by the order declared)
<i>union</i>	type tag followed by the selected member

Figure 8. CORBA CDR for constructed types

CORBA Common Data Representation (CDR)

- **Constructed types:** The primitive values that comprise each constructed type are added to a sequence of bytes in a particular order, as shown in **Figure 8**.

CORBA Common Data Representation (CDR)

- Figure 9 shows a message in CORBA CDR that contains the three fields of a **struct** whose respective types are **string**, **string**, and **unsigned long**.

CORBA Common Data Representation (CDR)

example: struct with value {'Smith', 'London', 1934}

<i>index in sequence of bytes</i>	<i>4 bytes</i>	<i>notes on representation</i>
0-3	5	<i>length of string</i>
4-7	"Smit"	<i>'Smith'</i>
8-11	"h____"	
12-15	6	<i>length of string</i>
16-19	"Lond"	<i>'London'</i>
20-23	"on____"	
24-27	1934	<i>unsigned long</i>

Figure 9. CORBA CDR message

Java object serialization

- Java's object serialization, which is concerned with the flattening and external data representation of any single object or tree of objects that may need to be transmitted in a message or stored on a disk. It is for use only by Java.
- In Java RMI, both object and primitive data values may be passed as arguments and results of method invocation.

Java object serialization

- An object is an instance of a Java class.
 - Example, the Java class equivalent to the Person struct

```
Public class Person implements Serializable {  
    Private String name;  
    Private String place;  
    Private int year;  
    Public Person(String aName ,String aPlace, int aYear) {  
        name = aName;  
        place = aPlace;  
        year = aYear;  
    }  
    //followed by methods for accessing the instance variables  
}
```

To serialize the *Person* object, create an instance of the class **ObjectOutputStream** and invoke its **writeObject** method, passing the *Person* object as its argument.

To deserialize an object from a stream of data, open an **ObjectInputStream** on the stream and use its **readObject** method to reconstruct the original object.

Java object serialization

The serialized form is illustrated in **Figure 10**.

<i>Serialized values</i>				<i>Explanation</i>
Person	8-byte version number		h0	class name, version number
3	int year	java.lang.String name	java.lang.String place	number, type and name of instance variables
1934	5 Smith	6 London	h1	values of instance variables

Figure 10. Indication of Java serialization form

XML

- ▶ XML (Extensible Markup Language), which defines a textual format for representing structured data.
- ▶ It was originally intended for documents containing textual self-describing structured data.
- ▶ Documents accessible on the Web – but it is now also used to represent the data sent in messages exchanged by clients and servers in web services

Remote Object References

- Remote object references are needed when a client invokes an object that is located on a remote server.
- A remote object reference is passed in the invocation message to specify which object is to be invoked.
- Remote object references must be unique over space and time.

Remote Object References

- In general, may be many processes hosting remote objects, so remote object referencing must be unique among all of the processes in the various computers in a distributed system.
- generic format for remote object references is shown in Figure 11.




Figure 11. Representation of a remote object references

Couloris, Dollimore and Kindberg Distributed Systems: Concepts & Design Edn. 4, Pearson Education 2005

Remote Object References

- internet address/port number: process which created object
- time: creation time
- object number: local counter, incremented each time an object is created in the creating process
- interface: how to access the remote object (if object reference is passed from one client to another)