

Threads

Reference: Pradeep K Sinha,
"Distributed Operating Systems: Concepts
and Design", Prentice Hall of India, 2007




Overview

- Threads
- Organizing Threads
- Issues in Designing Threads
- Thread Scheduling
- Thread Implementation




Process

- A basic unit of **CPU utilization** in traditional OS.
- A **program** in **execution**.
- Defines a **data space**.
- Has its own program counter, its own register states, its own stack and its own address space.
- Has at **least one** associated **thread**.
- Unit of distribution.



Threads

- It is a **basic unit** of **CPU utilization** used for improving a system performance through **parallelism**.
- A process consist of an **address space** and **one** or **more threads** of control.
- Threads share **same address space** but having its own **program counter, register states and its own stack**.
- Less protection due to the sharing of address space.




Threads


- On a **uniprocessor**, threads run in quasi-parallel (**time sharing**), whereas on a **shared-memory multiprocessor**, as many threads can run **simultaneously** as there are processors.
- **States of Threads:**
 - Running, blocked, ready, or terminated.

- Threads are viewed as **mini-processes**.


Threads



(a) Single-threaded




(b) Multithreaded processes




Motivations for using Threads

- Overheads involved in creating a new process is more than creating a new thread.
- Context switching between threads is cheaper than processes due to their same address space.
- Threads allow parallelism to be combined with sequential execution and blocking system calls.
- Resource sharing can be achieved more efficiently and naturally between threads due to same address space.



Different models to construct a server process: As a single-thread process

- Use **blocking system calls** but **without** any **parallelism**.
- If a dedicated machine is used for the file server, the **CPU remains idle** while the file server is waiting for a **reply** from the **disk space**.
- No **parallelism** is achieved in this method and **fewer client requests** are **processed** per unit of time.




As a finite state machine

- Model support parallelism but with **nonblocking** system calls.
- Implemented as a single threaded process and is operated like a **finite state machine**.
- An **event queue** is maintained for request & reply.
- During time of a disk access, it records **current state** in a **table** & **fetches** next **request** from queue.
- When a disk operation completes, the appropriate piece of **client state** must be **retrieved** to find out how to continue **carrying** out the request.




As a group of threads

- Supports **parallelism** with **blocking** system calls.
- Server process is comprised of a **single dispatcher thread** and **multiple worker** threads.
- Dispatcher thread keeps waiting in a loop for **request** from the **clients**.
- A server process designed in this way has **good performance** and is also **easy to program**.




Models for organizing threads

- Dispatcher-workers model
- Team model
- Pipeline model




Dispatcher-worker model

- Single dispatcher thread and multiple worker threads.
- Dispatcher thread accepts requests from clients and after examining the request, dispatches the request to one of the free worker threads for further processing of the request.
- Each worker thread works on a different client request.



Dispatcher-worker model




Team Model

- There is no dispatcher-worker relationship for processing clients requests.
- Each thread gets and processes clients' requests on its own.
- Each thread of the process is specialized in servicing a specific type of request.
- Multiple types of requests can be simultaneously handled by the process.




Team Model




A Server process for processing incoming requests that may be of three different types, each types of request being handled by a different thread.

Pipeline Model

- Useful for application based on the **producer-consumer model**.
- The **threads** of a **process** are organized as a **pipeline** so that the output data generated by the **first thread** is **used** for **processing** by the **second thread** and so on.




Pipeline Model



A Server process for processing incoming requests, each request processed in three steps, each step handled by a different thread and output of one step as input to the next step

Issues in designing a Thread Package

1. Threads creation
2. Thread termination
3. Threads synchronization
4. Threads scheduling
5. Signal handling




Threads Creation

- Created either **statically** or **dynamically**.
- **Static approach:**
 - **Threads** remain **fixed** for its entire lifetime.
 - **Fixed stack** is allocated to each thread.
 - **No. of threads** of a process is **decided** at the time of **writing** a program or during **compilation**.
- **Dynamic approach:**
 - **Number of threads** changes **dynamically**.




Threads Termination

- A thread may either **destroy itself** when it finishes its job by making an exit call
or
- be **killed** from **outside** by using the **kill command** and specifying the **thread identifier** as its parameter.




Threads Synchronization

- Threads share a common address space, so it is necessary to prevent multiple threads from trying to access the same data simultaneously.
- Mechanism for threads synchronization:
 - Use of global variable within the process.
 - mutex variable and condition variable.




Threads Synchronization

- **Segment** of **code** in which a thread may be accessing some **shared** variable is called a **critical region**.
- Two commonly used mutual exclusion techniques in a threads package are **mutex** variables and **condition** variables.

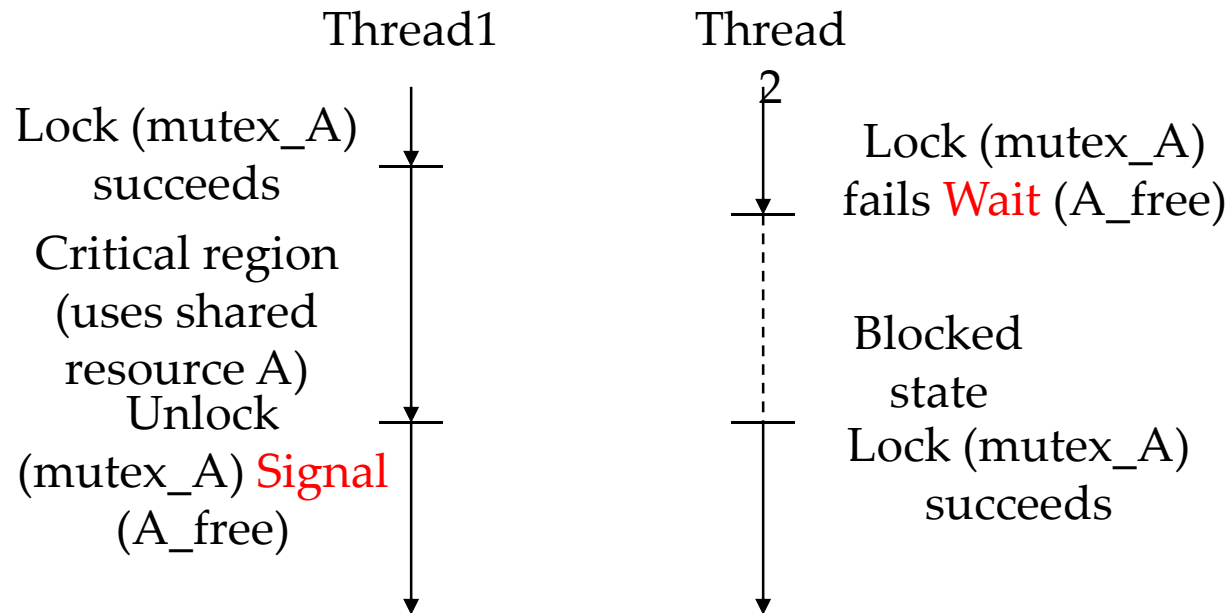


Threads Synchronization

- **Mutex** have only 2 states, hence simple to implement.
- **Condition variable** associated with mutex variable & reflects Boolean state of that variable.
- Condition – **wait and signal** operation.




Threads Synchronization



Mutex_A is a mutex variable for exclusive use of shared resource A.
A_free is a condition variable for resource A to become free.


Threads Scheduling

- Special features for threads scheduling :
 1. Priority assignment facility
 2. Flexibility to vary quantum size dynamically
 3. Handoff scheduling
 4. Affinity scheduling




Priority assignment facility

- Threads are scheduled on the first-in, first-out basis or the round-robin policy is used.
- Used to timeshares the CPU cycles.
- To provide more flexibility priority is assigned to the various threads of the applications.
- Priority thread maybe non-preemptive or preemptive.




Flexibility to vary quantum size dynamically

- Use of **round-robin** scheduling scheme.
- Varies the size of **fixed-length time quantum** to **timeshare the CPU cycle** among the **threads**.
- Not suitable for **multiprocessor** system.
- Gives **good response time** to **short requests**, even on heavily loaded systems.
- Provides **high efficiency** on **lightly loaded** systems.




Handoff Scheduling

- Allows a **thread** to name its **successor** if it wants to.
- Provides flexibility to **bypass** the queue of **runnable threads** and **directly switch** the **CPU** to the **thread** specified by the **currently running thread**.




Affinity scheduling

- A **thread** is **scheduled** on the **CPU** it last ran on, in hopes that part of its **address space** is still in that CPU's cache.
- Gives better **performance** on a **multiprocessor** system.




Signal handling

- Signals provide software-generated **interrupts** and **exceptions**.
- Issues :
 - A **signal** must be **handled properly** no matter which **thread** of the **process** receives it.
 - Signals must be **prevented** from getting **lost**.




Signal handling

- Approach:
 - Create a **separate exception handler thread** in each process.
 - Assign each thread its own **private global variables** for **signaling** conditions.




Thread Implementation

- Often provided in form of a **thread package**.
- Two important approaches:
 - First is to construct **thread library** that is entirely **executed** in **user mode**.
 - Second is to have the **kernel** be **aware** of **threads** and **schedule** them.




User level threads

- It is **cheap** to create and destroy threads.
- Cost of creating or destroying thread is determined by the cost for **allocating memory** to set up a thread stack.
- Switching thread context is done in few instructions.
- Only CPU registers need to be **stored** & subsequently **reloaded** with the previously stored values of the **thread** to which it is being **switched**.
- There is **no** need to change **memory maps**, **flush** the TLB, do CPU accounting etc.




User level threads

- Drawback:
 - Invocation of a **blocking system call** will **immediately block** the entire process to which the thread belongs.




User and Kernel Level

- **Threads** are useful to **structure** large application into parts that could be **logically executed** at the **same time**.
- In user level thread, **blocking** on **I/O** does prevent other parts to be **executed** in the meantime.
- This can be circumvented by implementing **threads** in the **OS Kernel**.



User and Kernel Level


- Price:
 - Every **thread** operation will have to be carried out by **kernel** requiring **system calls**.
 - Hence benefits of using threads disappear.




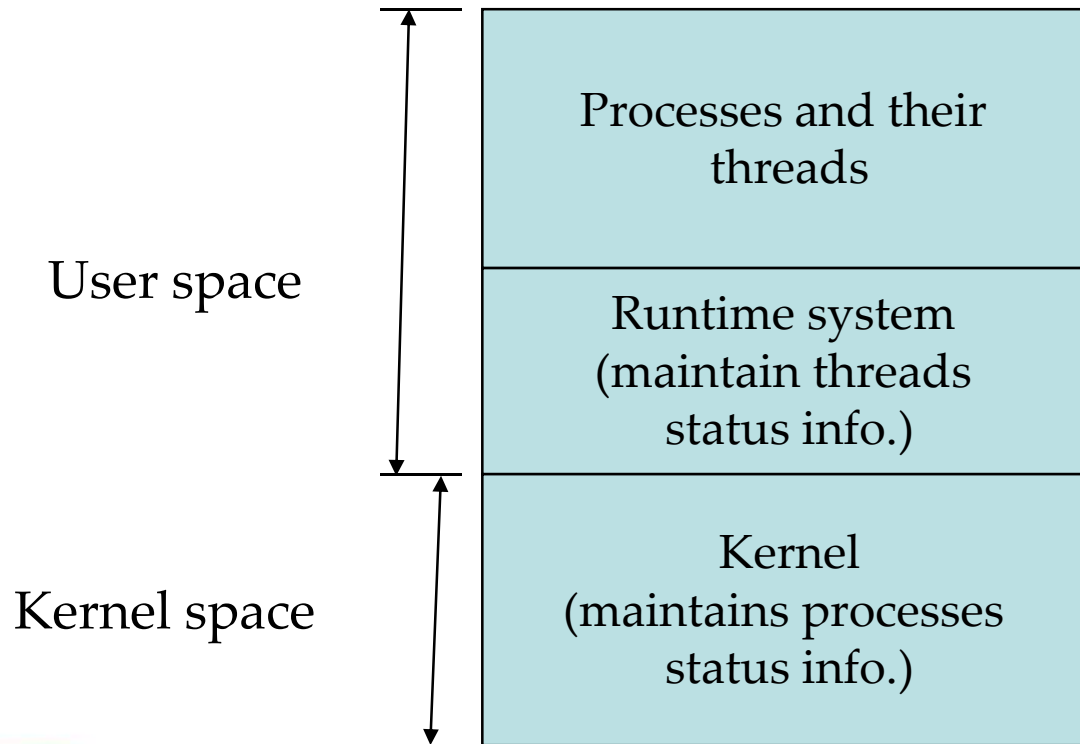
Implementing a Threads Package

User-level approach

- Consists of a **runtime** system that is a **collection** of **threads management** routines.
- Threads run in the **user space** on the top of the **runtime system** and are managed by it.
- System maintains a **status information table** having **one entry per thread**.
- **Two-level scheduling** is performed in this approach.




User level




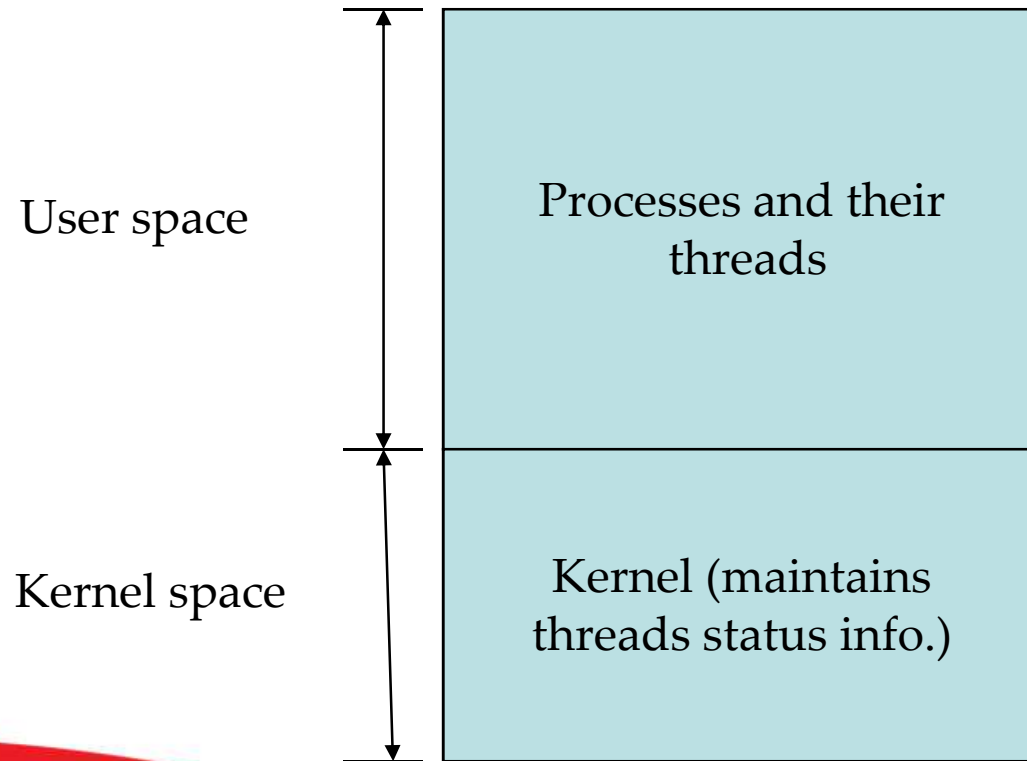
Implementing a Threads Package

Kernel-level approach

- No runtime system is used and threads are managed by the kernel.
- Threads status info. table maintained within kernel.
- Single-level scheduling is used in this approach.
- All calls that might block a thread are implemented as system calls that trap to the kernel.




Kernel level



Implementing a Threads Package

Advantages of the approaches


- User-level approach can be implemented on top of an existing OS that does not support threads.
- In user-level approach due to use of two-level scheduling, users have the flexibility to use their own customized algorithm to schedule the threads of a process.
- Switching the context from one thread to another is faster in the user-level approach than in the kernel approach.
- No status information cause poor scalability in the kernel as compared to the user-level approach.



Implementing a Threads Package

Disadvantages of the approaches:

- Since there is **no clock interrupt** within a single process, so **once CPU** is given to a **thread** to run, there is **no way to interrupt** it.
- To solve this **runtime system** can request a **clock interrupt** after every fixed **unit of time**



Thank You

