

Peer-to-Peer Systems

Case Study: Pastry & Tapestry

**Reference: George Coulouris, Jean Dollimore and Tim Kindberg,
“Distributed Systems Concepts and Design”, Fifth Edition, Pearson
Education, 2012**

Case Studies: Pastry

- ❑ Nodes & objects assigned 128 bit GUID computed by applying Secure Hash Algorithm to node's public key or object's name or stored state.
- ❑
 - ❑ GUIDs randomly distributed in range 0 to $2^{128} - 1$
 - ❑ Provide no clue to value from which computed
 - ❑ Clashes between GUIDs for different nodes are unlikely
- ❑ If GUID identifies node currently active, message delivered to it else to node whose GUID is numerically closest
- ❑ Delivery in $O(\log N)$ steps
- ❑ Routing uses underlying transport to transfer message to node closer to destination (may involve many IP hops)
- ❑ Pastry uses locality metric based on hop count or delay in underlying network to select appropriate neighbors when setting up routing tables



Distributed Hash Table API in Pastry

put(*GUID*, *data*)

- The *data* is stored in replicas at all nodes responsible for the object identified by *GUID*.


remove(*GUID*)

- Deletes all references to *GUID* and the associated data.

value = *get*(*GUID*)

- The data associated with *GUID* is retrieved from one of the nodes responsible for it.

Simple Pastry Routing Algorithm



- Each node stores leaf set – vector L (of size $2l$) containing GUIDs and IP addresses of l nearest nodes above and l below its GUID
- GUID space is circular: 0's neighbour is $2^{128}-1$
- The dots depict live nodes.
- The diagram illustrates the routing of a message from node 65A1FC to D46A1C using leaf set information alone, assuming leaf sets of size 8 ($l = 4$).
- This is a degenerate type of routing that would scale very poorly; it is not used in practice.



Simple Pastry Routing Algorithm

- Each Pastry node maintains a **tree** structured **routing table** giving **GUIDs** and **IP addresses** for a **set** of **nodes** spread throughout the entire range of 2^{128} possible values with increased density of coverage for **GUIDs numerically close** to its own.
- For GUIDs represented as **hexadecimal** numbers, **routing table** has as many rows as hex digits in a **GUID** = $128/4 = 32$ **rows**
- Any row has **15 entries** - **one** for **each possible value** of the **nth hex digit** excluding the value in the **local node's GUID**
- Each **entry** in **table** points to **one** of the **potentially many nodes** whose **GUIDs** have **relevant prefix**

First 4 Rows of a Pastry Routing Table

$p =$	<i>GUID prefixes and corresponding nodehandles n</i>															
0	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
	n	n	n	n	n	n		n	n	n	n	n	n	n	n	n
1	60	61	62	63	64	65	66	67	68	69	6A	6B	6C	6D	6E	6F
	n	n	n	n	n		n	n	n	n	n	n	n	n	n	n
2	650	651	652	653	654	655	656	657	658	659	65A	65B	65C	65D	65E	65F
	n	n	n	n	n	n	n	n	n	n		n	n	n	n	n
3	65A0	65A1	65A2	65A3	65A4	65A5	65A6	65A7	65A8	65A9	65AA	65AB	65AC	65AD	65AE	65AF
	n		n	n	n	n	n	n	n	n	n	n	n	n	n	n

The routing table is located at a node whose GUID begins 65A1. Digits are in hexadecimal. The n 's represent [GUID, IP address] pairs specifying the next hop to be taken by messages addressed to GUIDs that match each given prefix. Grey-shaded entries indicate that the prefix matches the current GUID up to the given value of p : the next row down or the leaf set should be examined to find a route. Although there are a maximum of 128 rows in the table, only $\log_{16} N$ rows will be populated on average in a network with N active nodes.



Pastry Routing Algorithm

if (destination is within range of our leaf set)

forward to numerically closest member

else

if (there is a longer prefix match in table)

forward to node with longest match

else

forward to node in table which

(a) has a common prefix of length p and

(b) GUID that is numerically closer.



Tapestry

- Tapestry implements a **distributed hash table** and routes **messages** to **nodes** based on **GUIDs** associated with resources using **prefix routing** in a manner similar to Pastry.
- Tapestry applications give additional **flexibility**:
- They can **place replicas close** (in network distance) to **frequent users** of resources in order to **reduce latency** and
- **Minimize network** load or to ensure **tolerance** of **network** and host **failures**

Distributed Object Location and Routing in Tapestry

publish(GUID)

- *GUID can be computed from the object (or some part of it, e.g. its name).*
- *This function makes the node performing a *publish* operation as host for the object corresponding to *GUID*.*

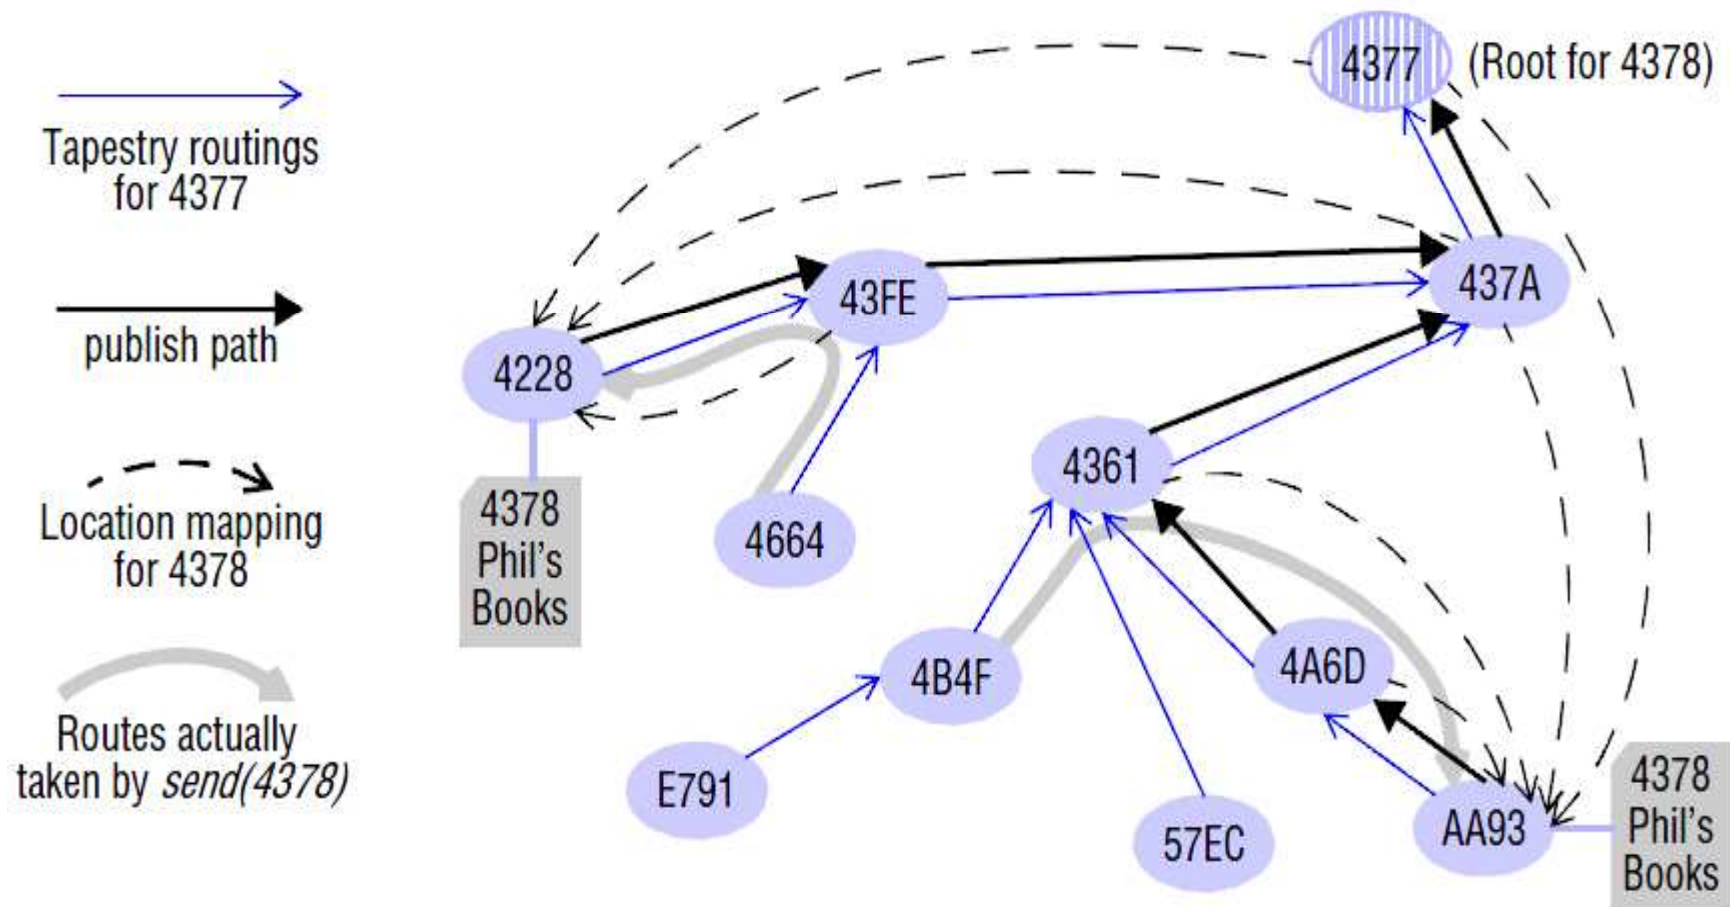
unpublish(GUID)

- *Makes the object corresponding to *GUID* inaccessible.*

sendToObj(msg, GUID, [n])

- *Following the object-oriented paradigm, an invocation message is sent to an object in order to access it. This might be a request to open a TCP connection for data transfer or to return a message containing all or part of the object's state. The final optional parameter *[n]*, if present, requests the delivery of the same message to *n* replicas of the object.*

Distributed Object Location and Routing in Tapestry



Structured Vs Unstructured P2P

	<i>Structured peer-to-peer</i>	<i>Unstructured peer-to-peer</i>
<i>Advantages</i>	Guaranteed to locate objects (assuming they exist) and can offer time and complexity bounds on this operation; relatively low message overhead.	Self-organizing and naturally resilient to node failure.
<i>Disadvantages</i>	Need to maintain often complex overlay structures, which can be difficult and costly to achieve, especially in highly dynamic environments.	Probabilistic and hence cannot offer absolute guarantees on locating objects; prone to excessive messaging overhead which can affect scalability.



Strategies for Search

- ❑ In P2P file sharing, all nodes in the network offer files to other nodes.
- ❑ Searching for a file in unstructured P2P network follows following strategies.
 1. Expanded Ring Search
 2. Random Walks
 3. Gossiping.



Thank You