

# Now back to Distributed Systems: remember?

4

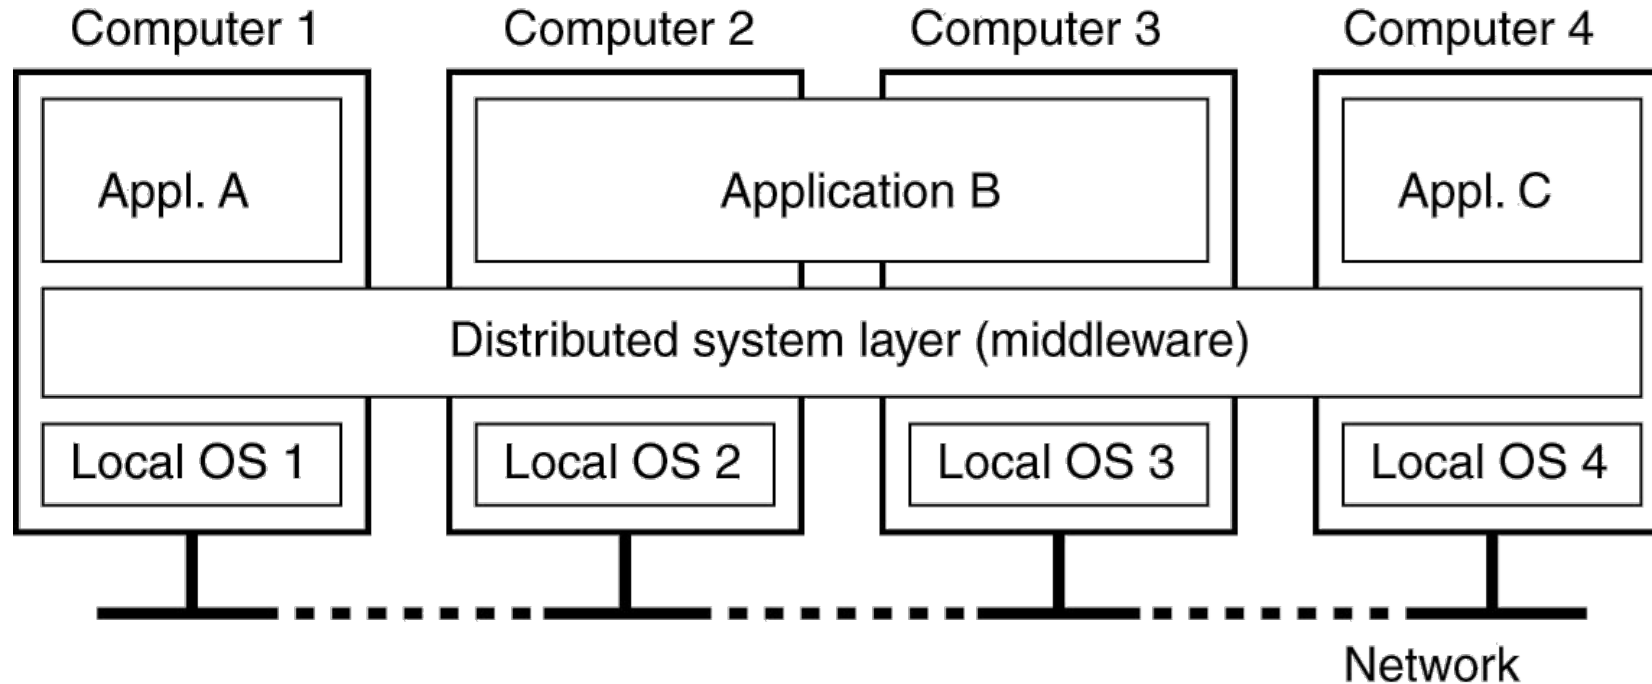
□ A distributed system is:

“A collection of independent computers that appears to its users as a single coherent system”

"A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable." – Leslie Lamport

# Distributed Systems

5



- The middleware layer extends over multiple machines, and offers each application the same interface.

# What does it do?

Hide complexity to programmers/users


Hide the fact that its processes and resources are physically distributed across multiple machines.

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource is replicated
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource

Transparency in a Distributed System

# How?

7



- The middleware layer extends over multiple machines, and offers each application the same interface.

# Starter for Today

---

## □ Splitting computation across the network

What programming abstractions work well to split work among multiple networked computers?

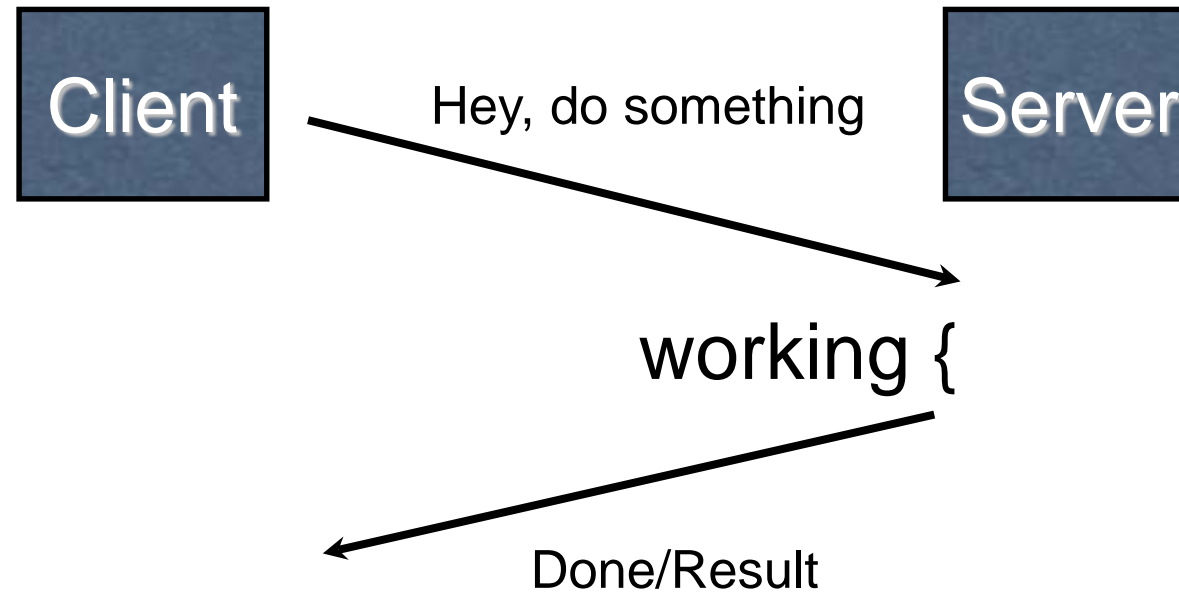
# Many ways

- Request-reply protocols
- Remote procedure calls (RPC)
- Remote method invocation (RMI)

Recommended reading :

**Distributed Systems: Concepts and Design 5<sup>th</sup> edition**, by Coulouris, et al.  
(CDK5) chapter 5

# Request-reply protocols



# Request-reply protocols

- eg, your PA1 (binary protocol)

```
struct foormsg {  
    u_int32_t len;  
}  
  
send_foo(char *contents) {  
    int msglen = sizeof(struct foormsg) + strlen(contents);  
    char buf = malloc(msglen);  
    struct foormsg *fm = (struct foormsg *)buf;  
    fm->len = htonl(strlen(contents));  
    memcpy(buf + sizeof(struct foormsg),  
           contents,  
           strlen(contents));  
    write(outsock, buf, msglen);  
}
```

Then wait for response, handle timeout, etc.



# Request-reply protocols: text protocol

## □ HTTP

- ▣ See the HTTP/1.1 standard:
- ▣ `http://www.w3.org/Protocols/rfc2616/rfc2616.html`
- ▣ Done with your PA2 yet?

# Remote Procedure Call (RPC)

- A type of client/server communication
- Attempts to make remote procedure calls look like local ones




figure from Microsoft MSDN

```
{ ...  
  foo()  
}  
void foo() {  
  invoke_remote_foo()  
}
```

# RPC Goals

- Ease of programming
- Hide complexity
- Automate a lot of task of implementing
- Familiar model for programmers (just make a function call)

Historical note: Seems obvious in retrospect, but RPC was only invented in the '80s. See Birrell & Nelson, "Implementing Remote Procedure Call" ... or Bruce Nelson, Ph.D. Thesis, Carnegie Mellon University: Remote Procedure Call., 1981 :)

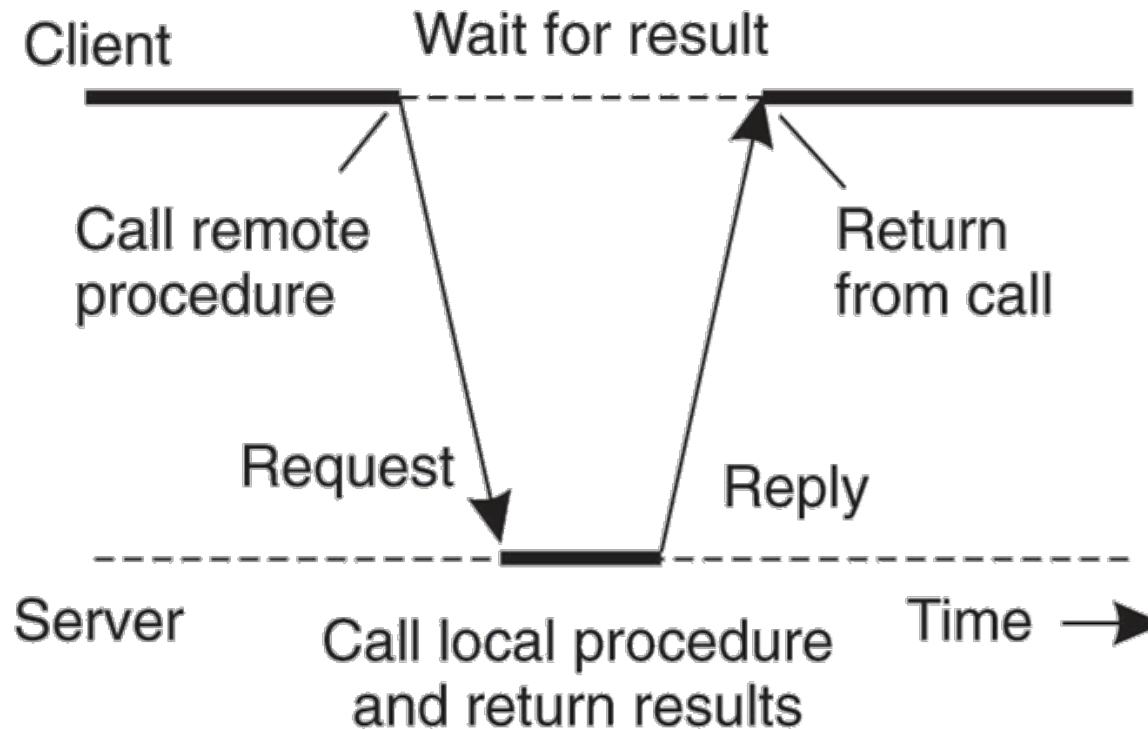
# Remote procedure call

15

- A remote procedure call makes a call to a remote service look like a local call
  - ▣ RPC makes transparent whether server is local or remote
  - ▣ RPC allows applications to become distributed transparently
  - ▣ RPC makes architecture of remote machine transparent

# RPC

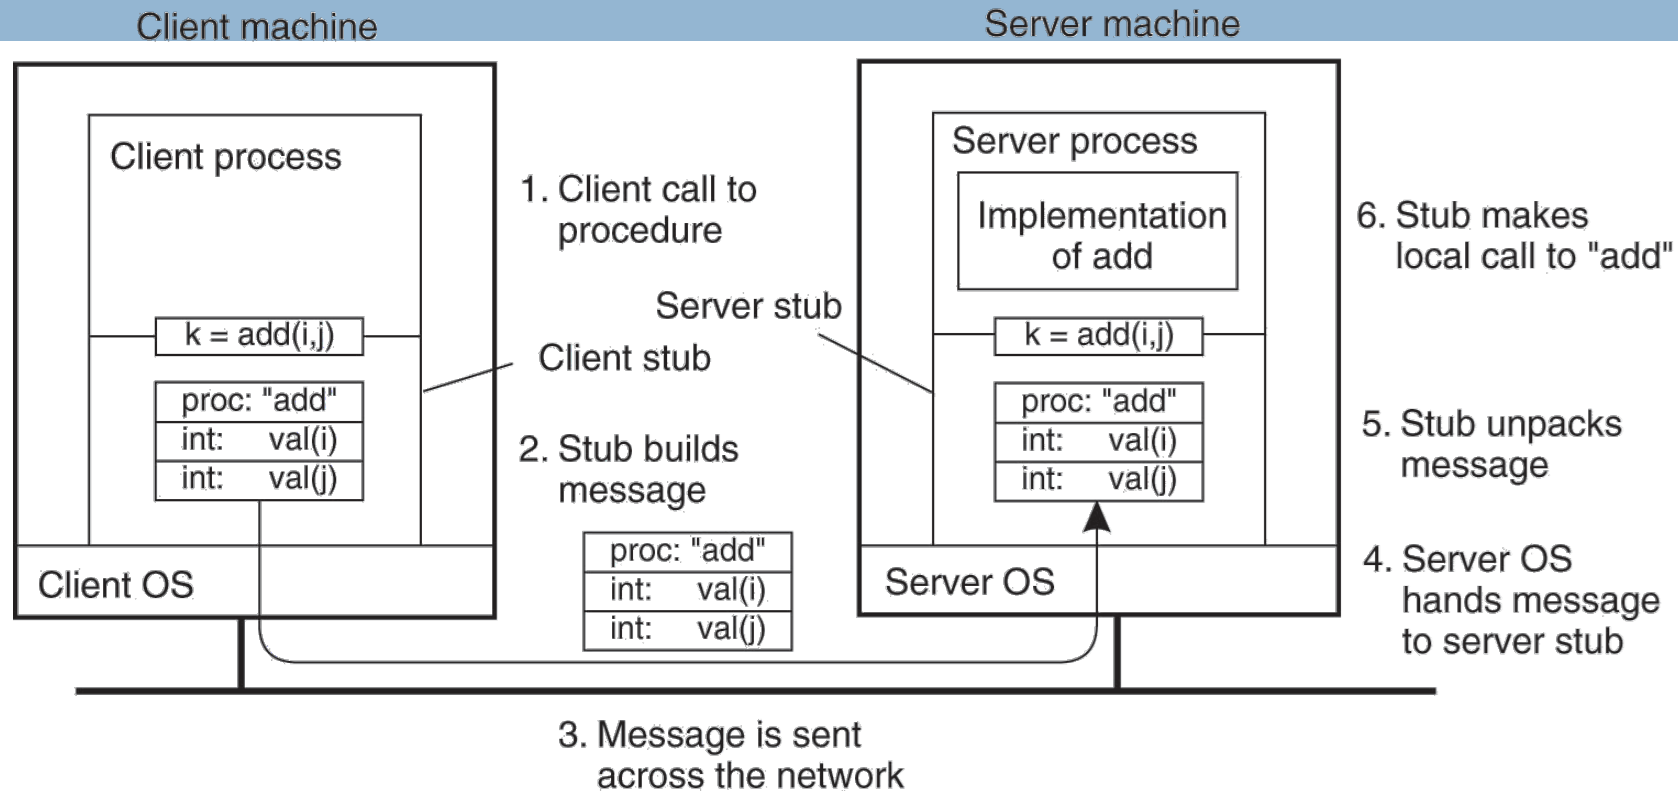
16



- The interaction between client and server in a traditional RPC.

# Passing Value Parameters (1)

17



- The steps involved in a doing a remote computation through RPC.

# But it's not always simple

- Calling and called procedures run on different machines, with different address spaces
  - ▣ And perhaps different environments .. or operating systems ..
- Must convert to local representation of data
- Machines and network can fail

# Marshaling and Unmarshaling

- (From example) `htonl()` -- “host to network-byte-order, long”.
  - ▣ network-byte-order (big-endian) standardized to deal with cross-platform variance
- Note how we arbitrarily decided to send the string by sending its length followed by L bytes of the string? That’s marshalling, too.
- Floating point...
- Nested structures? (Design question for the RPC system - do you support them?)
- Complex datastructures? (Some RPC systems let you send lists and maps as first-order objects)



# “stubs” and IDLs

- RPC stubs do the work of marshaling and unmarshaling data
- But how do they know how to do it?
- Typically: Write a description of the function signature using an *IDL* -- interface definition language.
  - ▣ Lots of these. Some look like C, some look like XML, ... details don't matter much.