

Publish-Subscribe Systems

George Coulouris, Jean Dollimore, Tim Kindberg and Gordon Blair , “**Distributed Systems Concepts and Design**” Edition 5, © Addison-Wesley 2012



Motivations for Pub/Sub model

- **Traditional Client/Server communication model**
 - (Employs RPC, message queue, shared memory etc..)
 - Synchronous, tightly-coupled request invocations.
 - Very restrictive for distributed applications, especially for WAN and mobile environments.
 - When nodes/links fail, system is affected. Fault Tolerance must be built in to support this.
- Require a more flexible and de-coupled communication style that offers anonymous and asynchronous mechanisms.

What is a Publish/Subscribe System?

- **Distributed event-based system**
- **Distributed Pub/Sub System** is a communication paradigm that allows **freedom in the distributed system** by the **decoupling of communication entities** in terms of time, space and synchronization.
- An event service system that is **asynchronous, anonymous and loosely-coupled**.
- Ability to quickly **adapt** in a dynamic environment.

Key components of Pub/Sub System

- **Publishers** : Publishers generate event data and publishes them.
- **Subscribers** : Subscribers submit their subscriptions and process the events received
- **P/S service**: It's the mediator/broker that filters and routes events from publishers to interested subscribers.
- Task of the publish/subscribe system is to match subscriptions against published events and ensure the correct delivery of event notifications.
- A given event will be delivered to potentially many Subscribers.
- Publish-subscribe is fundamentally a **one to-many communications paradigm**.

Publish/Subscribe System

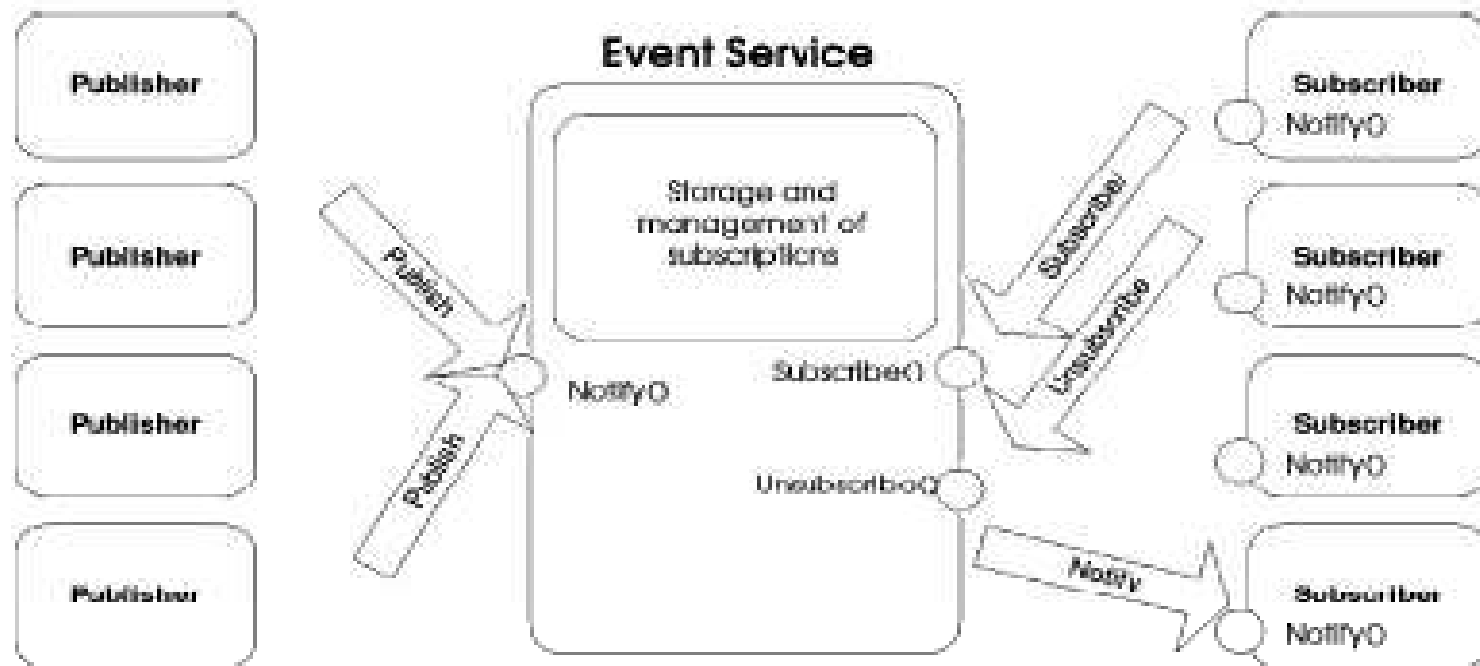


Fig. 1. A simple object-based publish/subscribe system.

Applications

- Financial information systems;
- Other areas with live feeds of real-time data (including RSS feeds);
- Support for cooperative working, where a number of participants need to be informed of events of shared interest;
- Support for ubiquitous computing, including the management of events emanating from the ubiquitous infrastructure (for example, location events);
- A broad set of monitoring applications, including network monitoring in the Internet.

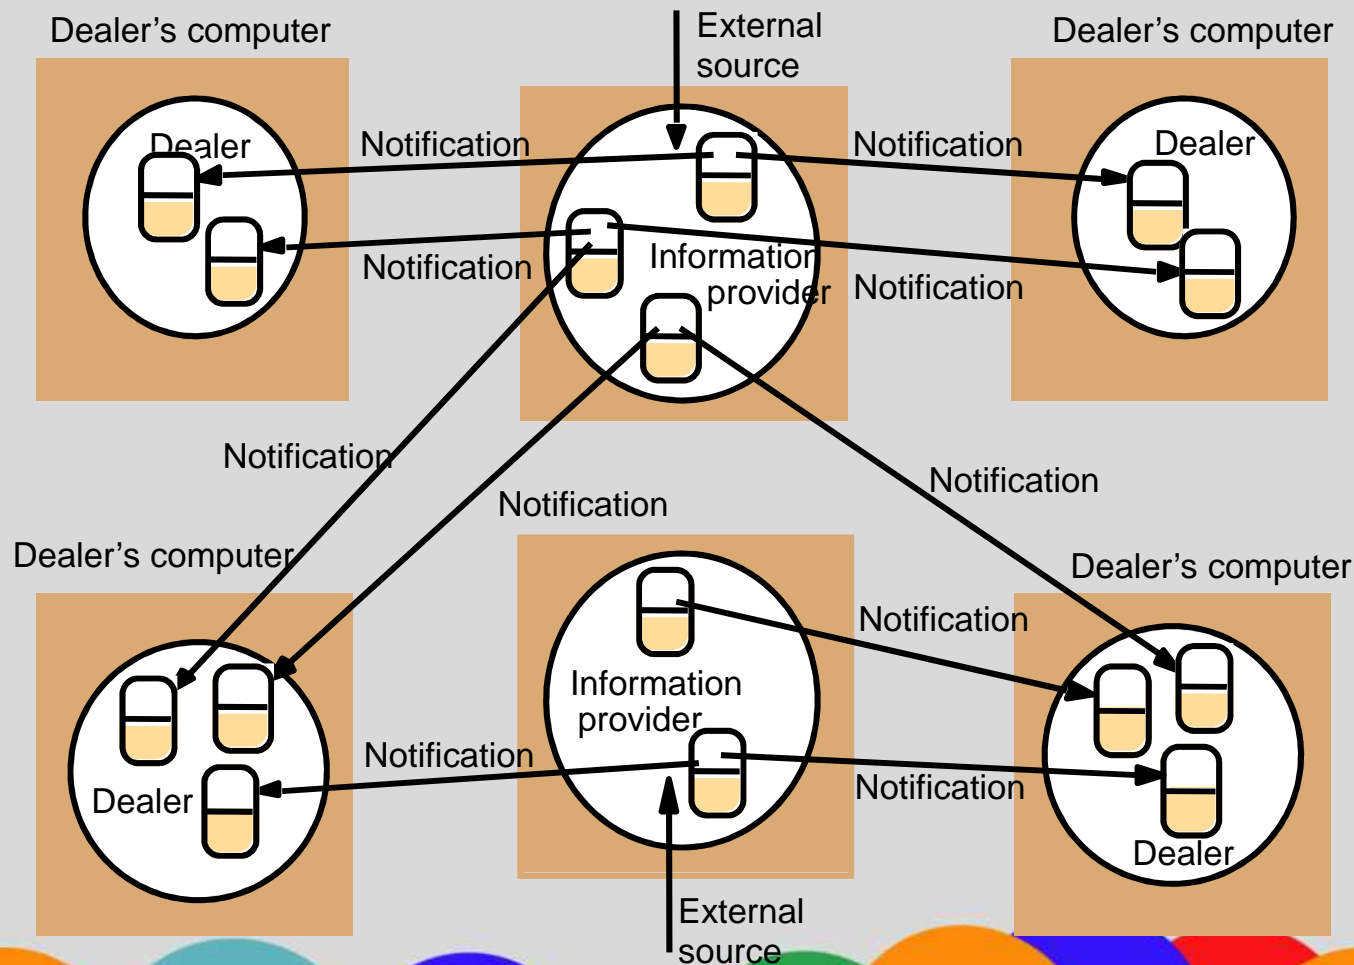


Dealing room system

A dealing room system could be implemented by processes with two different tasks: To know market updates on Stock exchange.

- An **information provider process** continuously receives new trading information from a single external source.
- Each of the updates is regarded as an event.
- The information provider publishes such events to the publish-subscribe system for delivery to all of the dealers who have expressed an interest in the corresponding stock.
- There will be a separate information provider process for each external source.
- A **dealer process** creates a subscription representing each named stock that the user asks to have displayed.
- Each subscription expresses an interest in events related to a given stock at the relevant information provider.
- It then receives all the information sent to it in notifications and displays it to the user.

Dealing room system



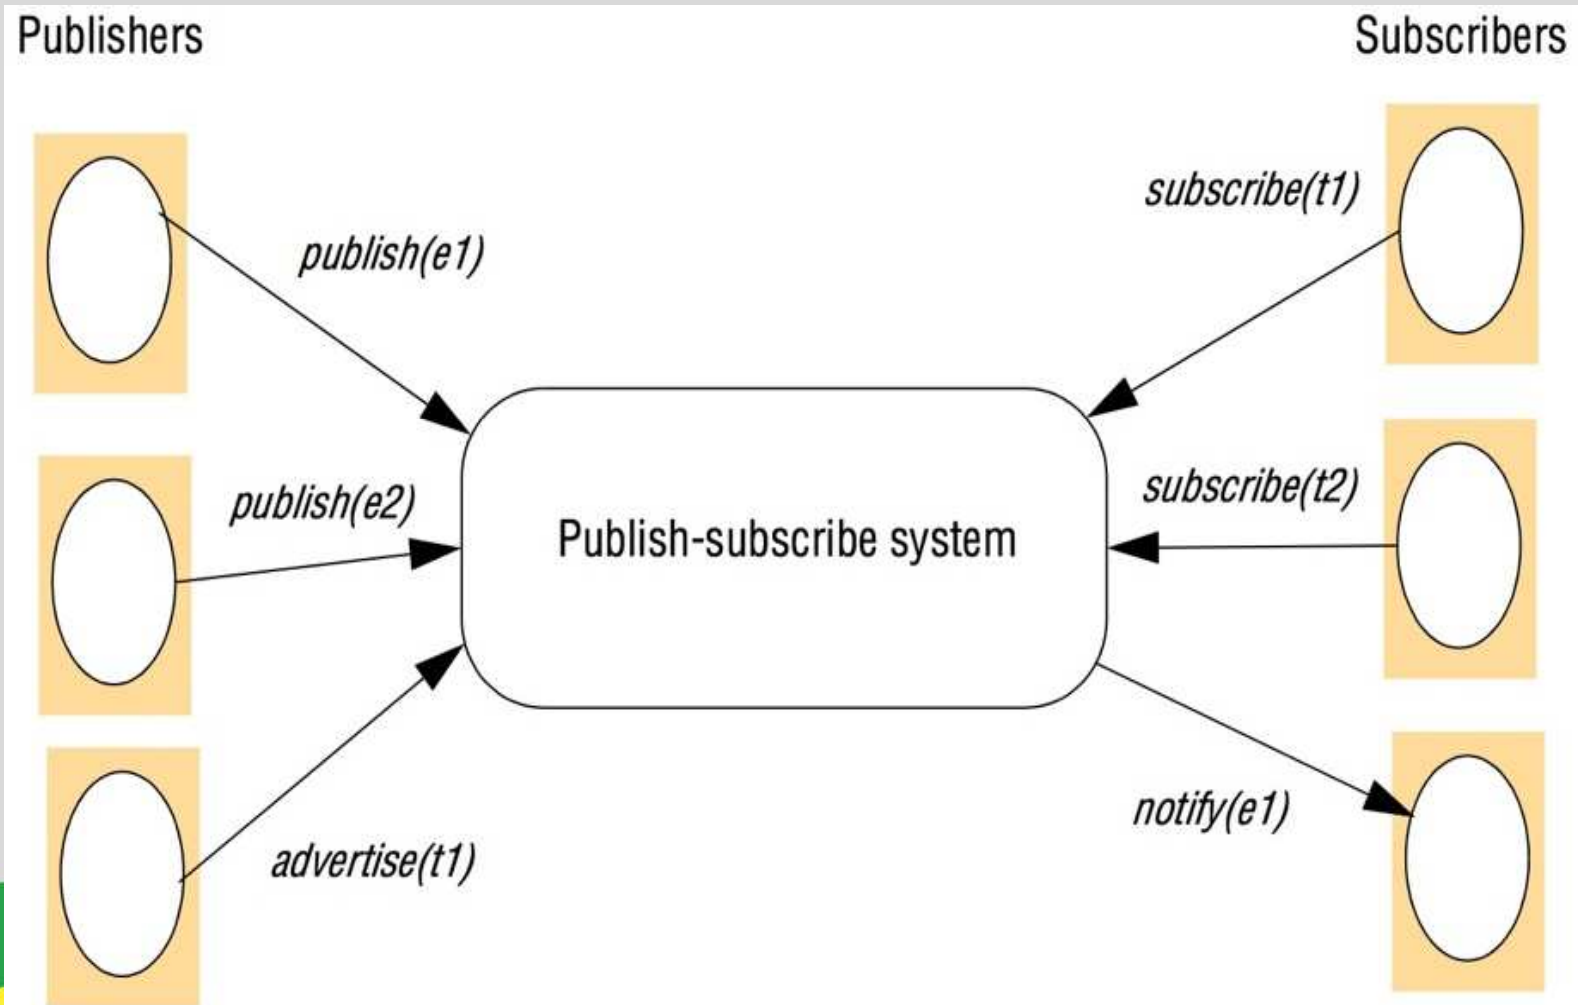
Characteristics of publish-subscribe systems

Publish-subscribe systems have two main characteristics:

Heterogeneity: When event notifications are used as a means of communication, components in a distributed system that were not designed to interoperate can be made to work together.

Asynchronicity: Notifications are sent asynchronously by event-generating publishers to all the subscribers that have expressed an interest in them to prevent publishers needing to synchronize with subscribers – publishers and subscribers need to be decoupled.

The publish-subscribe paradigm



Publish-subscribe systems

The expressiveness of publish-subscribe systems is determined by the subscription (filter) model, with a number of schemes defined and considered in increasing order of sophistication:

Channel-based : publishers publish events to named channels and subscribers then subscribe to one of these named channels to receive all events sent to that channel.

Topic-based (also referred to as subject-based) : each notification is expressed in terms of a number of fields, with one field denoting the topic. Subscriptions are then defined in terms of the topic of interest.

Publish-subscribe systems


Content-based : generalization of topic-based approaches allowing the expression of subscriptions over a range of fields in an event notification.

Type-based : intrinsically linked with object-based approaches where objects have a specified type. In type-based approaches, subscriptions are defined in terms of types of events and matching is defined in terms of types or subtypes of the given filter.

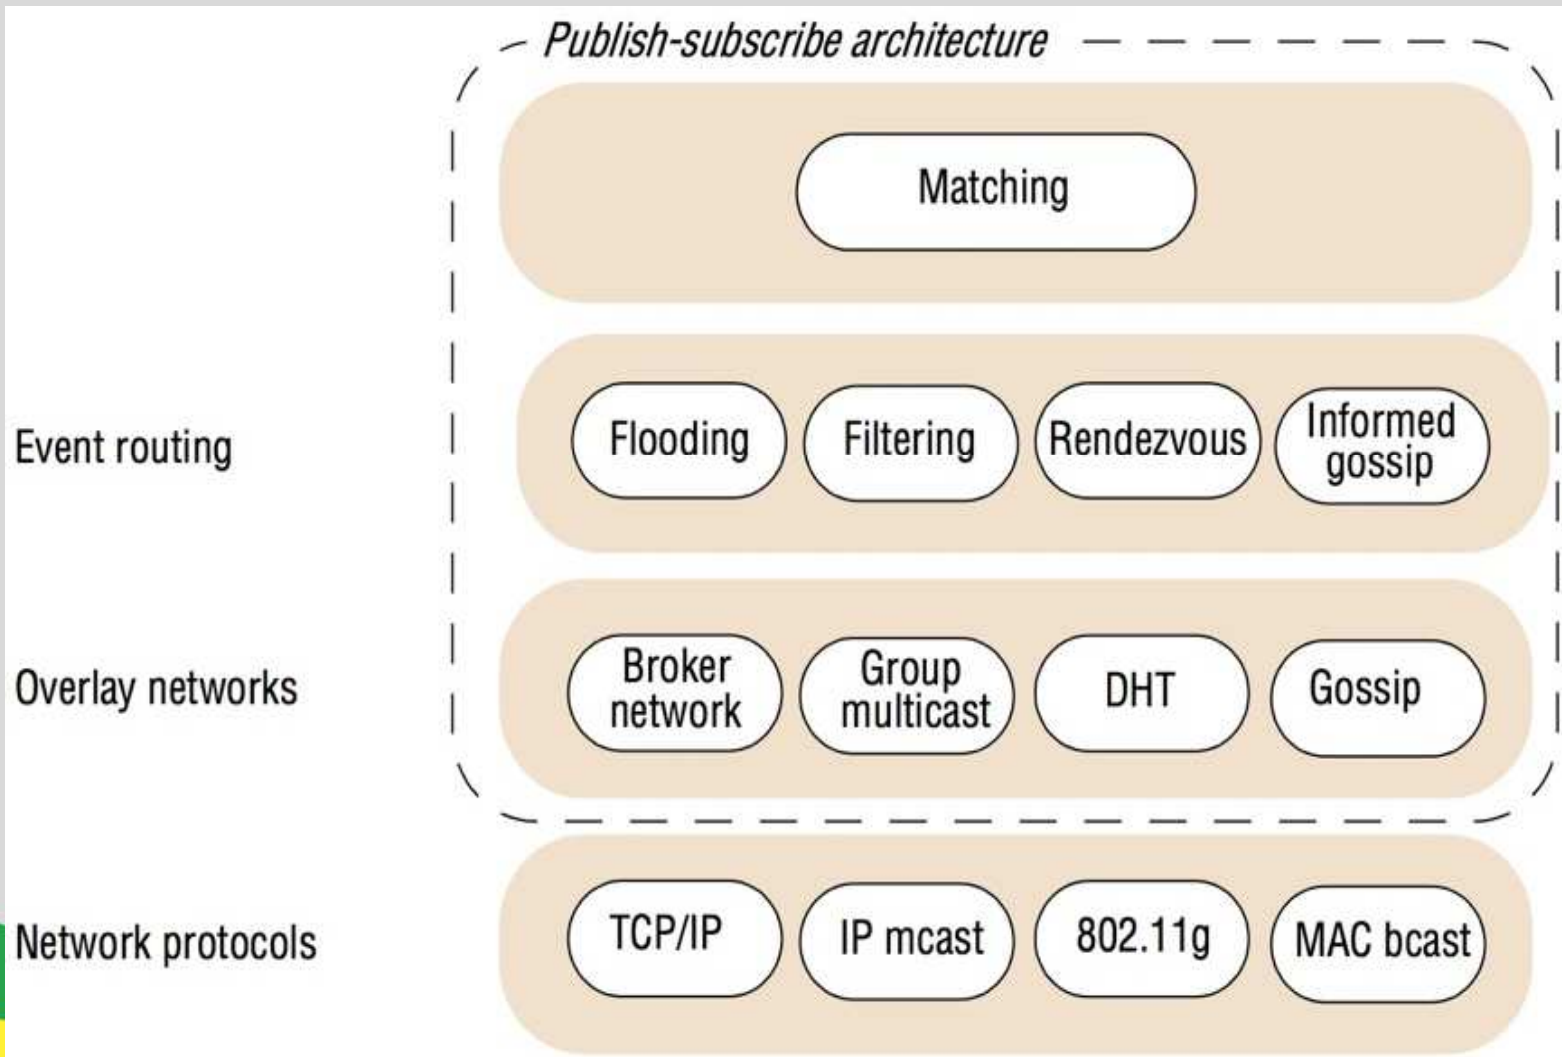
Implementation Issues

- **Centralized Implementation** : A single node with a server on that node acts as an event broker.
 - Publishers then publish events to this broker, subscribers send subscription to the broker and receive notifications in return.
 - Broker is Single point of failure
- **Distributed Implementation**: Centralized broker is replaced by network of brokers that cooperate to offer desired functionality.
- **Peer –to-peer Implementation**:
 - Each node can be publisher, subscriber or broker.
 - Subscribers subscribe to publishers directly and publishers notify subscribers directly. Therefore they must maintain knowledge of each other.

A network of brokers



The architecture of publish-subscribe systems



Architecture

- Event routing performs the task of ensuring that **event notifications** are **routed** as efficiently as possible to appropriate subscribers.
- The **overlay infrastructure** supports this by setting up appropriate networks of **brokers** or **peer-to-peer** structures.
- **Content-based routing (CBR)**, with the goal being to exploit content information to efficiently route events to their required destination

Flooding

- In Flooding: Sending an event notification to all nodes in the network and then carrying out the appropriate matching at the subscriber end.
- As an alternative, flooding can be used to send subscriptions back to all possible publishers, with the matching carried out at the publishing end.
- Matched events sent directly to the relevant subscribers using point-to-point communication.

Filtering-based routing

- Brokers forward notifications through the network only where there is a path to a valid subscriber.
- Each node must maintain a neighbors list containing a list of all connected neighbors in the network of brokers.
- A subscription list containing a list of all directly connected subscribers serviced by this node, and a routing table.
- Routing table maintains a list of neighbors and valid subscriptions for that pathway.
- It requires an implementation of matching on each node in the network of brokers

Filtering-based routing

<i>upon receive publish(event e) from node x</i>	<i>1</i>
<i>matchlist := match(e, subscriptions)</i>	<i>2</i>
<i>send notify(e) to matchlist;</i>	<i>3</i>
<i>fwddlist := match(e, routing);</i>	<i>4</i>
<i>send publish(e) to fwddlist - x;</i>	<i>5</i>
<i>upon receive subscribe(subscription s) from node x</i>	<i>6</i>
<i>if x is client then</i>	<i>7</i>
<i>add x to subscriptions;</i>	<i>8</i>
<i>else add(x, s) to routing;</i>	<i>9</i>
<i>send subscribe(s) to neighbours - x;</i>	<i>10</i>

Filtering-based routing

- When a **broker** receives a **publish** request from a given node,
 - It must **pass** this **notification** to **all connected nodes** where there is a corresponding matching subscription.
 - **Decide** where to **propagate** this **event** through the **network of Brokers**.
- Matching the **event** against the **subscription** list and then forwarding the **event** to all the **nodes** with matching subscriptions (the *matchlist*) (line 2,3)
- Matching the **event** against the **routing table** and forwarding **only** to the **paths** that lead to a **subscription** (the *fwdlst*). (line 4,5)
- Brokers must also deal with **incoming subscription** events. If the subscription event is from an **immediately connected subscriber**, then this **subscription** must be entered in the **subscriptions table**. (Line 7,8)
- Otherwise, the **broker** is an **intermediary node**; this node now knows that a **pathway** exists **towards** this **subscription** and hence an appropriate entry is added to the **routing table** (line 9).
- In both cases, this **subscription event** is then **passed** to all **neighbors** apart from the **originating node** (line 10).

Advertisements

- The pure filtering-based approach described above can generate a **lot** of **traffic** due to propagation of subscriptions.
- In advertisements this burden can be reduced by **propagating** the **advertisements** towards **subscribers** in a similar (actually symmetrical) way to the propagation of subscriptions.



Rendezvous-based Routing

- This view the set of all possible events as an event space and to partition responsibility for this event space between the set of brokers in the Network
- This approach defines rendezvous nodes, which are broker nodes responsible for a given subset of the event space



Rendezvous-based Routing

- First, $SN(s)$ takes a given **subscription**, s , and returns **one or more rendezvous nodes** that take responsibility for that **subscription**.
- Each such **rendezvous node** maintains a **subscription list** and **forwards** all **matching events** to the set of **subscribing nodes**



Rendezvous-based Routing

- Second, when an **event e is published**, the function **$EN(e)$** also **returns one or more rendezvous nodes**, this time responsible for **matching e against subscriptions** in the system
- Both **$SN(s)$ and $EN(e)$** return more than one node if reliability is a concern.
- Note also that this approach only works if the **intersection of $EN(e)$ and $SN(s)$ is non-empty** for a given **e that matches s**



Rendezvous-based routing

```
upon receive publish(event e) from node x at node i  
  rvlist := EN(e);  
  if i in rvlist then begin  
    matchlist := match(e, subscriptions);  
    send notify(e) to matchlist;  
  end  
  send publish(e) to rvlist - i;  
upon receive subscribe(subscription s) from node x at node i  
  rvlist := SN(s);  
  if i in rvlist then  
    add s to subscriptions;  
  else  
    send subscribe(s) to rvlist - i;
```

Gossip

- Gossip-based approaches are a popular mechanism for achieving multicast (including reliable multicast)
- They operate by nodes in the network periodically and probabilistically exchanging events (or data) with neighboring nodes.
- It is possible to take into account local information and, in particular, content to achieve what is referred to as informed gossip



Example publish-subscribe system

<i>System (and further reading)</i>	<i>Subscription model</i>	<i>Distribution model</i>	<i>Event routing</i>
CORBA Event Service (Chapter 8)	Channel-based	Centralized	-
TIB Rendezvous [Oki <i>et al.</i> 1993]	Topic-based	Distributed	Filtering
Scribe [Castro <i>et al.</i> 2002b]	Topic-based	Peer-to-peer (DHT)	Rendezvous
TERA [Baldoni <i>et al.</i> 2007]	Topic-based	Peer-to-peer	Informed gossip
Siena [Carzaniga <i>et al.</i> 2001]	Content-based	Distributed	Filtering
Gryphon [www.research.ibm.com]	Content-based	Distributed	Filtering
Hermes [Pietzuch and Bacon 2002]	Topic- and content-based	Distributed	Rendezvous and filtering
MEDYM [Cao and Singh 2005]	Content-based	Distributed	Flooding
Meghdoot [Gupta <i>et al.</i> 2004]	Content-based	Peer-to-peer	Rendezvous
Structure-less CBR [Baldoni <i>et al.</i> 2005]	Content-based	Peer-to-peer	Informed gossip

Thank You