

The background features a dark gray upper half and a white lower half, separated by a horizontal orange bar. Faint, concentric gray arcs are visible across the entire slide.

# Co-ordinator Election Algorithms

George Coulouris, Jean Dollimore and Tim Kindberg,  
“Distributed Systems Concepts and Design”, Fifth Edition,  
Pearson Education, 2012

# Election Introduction

- An algorithm for choosing a **unique process** to play a particular role is called an **Election Algorithm**.
- It is essential that **all the processes agree** on the **choice**. Afterwards, if the process that plays the role of **server** wishes to **retire** then another election is required to **choose a replacement**.
- An individual process **does not call more** than **one election** at a time, but in principle the  $N$  processes could call  $N$  concurrent elections
- A process  $p_i$  is
  - either a **participant** –it is engaged in some run of the election algorithm.
  - or a **non-participant** –it is not currently engaged in any election

# Election Introduction

- Important requirement is for the choice of **elected** process to be **unique**, even if several **processes** call **elections concurrently**.
- For instance, **two processes** could decide **independently** that a **coordinator process** has **failed**, and both **call elections**.
- We require that the elected process be chosen as the **one** with the **largest identifier**.
- Each process  $p_i$  ( $i = 1, 2, \dots, N$ ) has a variable  $\text{elected}_i = \perp$ , which will contain the
  - identifier of the elected process
- **E1: (safety)** A participant process  $p_i$  has  $\text{elected}_i = \perp$  or  $\text{elected}_i = P$ , where  $P$  is chosen as the non-crashed process at the end of the run with the largest identifier.
- **E2: (liveness)** All processes  $p_i$  participate and eventually either set  $\text{elected}_i = \perp$  - or crash.
- Performance of an election algorithm by its total **network bandwidth utilization** and by the **turnaround time** for the algorithm

# Election Algorithms


- Ring-based Election Algorithm
- Bully Algorithm

# Election Algorithms

## Ring-based Election Algorithm

# Ring-based Election Algorithm

- The **algorithm** of Chang and Roberts [1979] is suitable for a **collection** of **processes arranged** in a **logical ring**.
- Each process  $p_i$  has a **communication channel** to the **next process** in the ring,  $p_{(i+1) \bmod N}$ , and all **messages** are **sent clockwise** around the ring
- It assume that **no failures** occur, and that the **system** is **asynchronous**
- The goal of this algorithm is to **elect** a **single process** called the **coordinator**, which is the **process** with the **largest identifier**



# Ring-based Election Algorithm

- Initially, **every process** is marked as a **non-participant** in an election.
- **Any process** can **begin** an **election**. It proceeds by **marking** itself as a **participant**, **placing** its **identifier** in an election message and **sending** it to its **clockwise** neighbor.
- When a **process receives** an election message, it **compares** the **identifier** in the message with its own.
- If the **arrived identifier** is **greater**, then it **forwards** the **message** to its **neighbor**.
- If the **arrived identifier** is **smaller** and the receiver is not a participant, then it **substitutes** its **own identifier** in the message and **forwards** it;
- On forwarding an election message in any case, the process marks itself as a participant
- If, however, **the received identifier is that of the receiver itself**, then this process's **identifier must be the greatest**, and it becomes the coordinator. The coordinator **marks** itself as a non-participant once more and **sends** an **elected message** to its **neighbor**, announcing its **election** and **enclosing** its **identity**

# Ring-based Election Algorithm

- When a process  $p_i$  receives an **elected message**, it **marks** itself as a **nonparticipant**, sets its variable **elected<sub>i</sub>** to the **identifier** in the **message** and, unless it is the **new coordinator**, forwards the message to its neighbour.
- **Condition E1 is met.** All **identifiers** are **compared**, since a process must **receive** its **own identifier back** before sending an elected message.
- For any two processes, the one with the **larger identifier** will **not pass** on the other's identifier. It is therefore **impossible** that **both** **receive** their **own identifier back**.
- **Condition E2** follows **immediately** from the **guaranteed traversals** of the **ring** (there are **no failures**).



# Ring-based Election Algorithm

- The worst-performing case is when its **anti-clockwise neighbour** has the **highest identifier**.
- A total of  $N - 1$  messages are then required to reach this neighbour.
- It will not announce its election until its identifier has completed another circuit, taking a further  $N$  messages.
- The elected message is then sent  $N$  times,
- **Making  $3N - 1$  messages in all.**
- The **turnaround** time is also  $3N - 1$  , since these messages are sent sequentially.
- Drawback
  - It does **not tolerate failures** making it limited practical value.

# Election Algorithms

## Bully Algorithm


# Bully Algorithm

- The **bully algorithm** [Garcia-Molina 1982] allows **processes** to **crash** during an **election**, although it assumes that **message delivery** between **processes** is **reliable**.
- This algorithm assumes that the **system** is **synchronous**: it uses **timeouts** to **detect** a process **failure**
- **In Bully algorithm** : Each process knows which processes have higher identifiers, and that it can communicate with all such processes. 3 message types.
- **Election message** is sent to announce an election;
- **Answer message** is sent in response to an election message and
- **Coordinator message** is sent to announce the identity of the elected process

# Bully Algorithm

- A process begins an election when it notices, through timeouts, that the coordinator has failed.
- Several processes may discover this concurrently.
- Maximum message transmission delay,  $T_{trans}$ , and a maximum delay for processing a message  $T_{process}$ .
- Time  $T = 2T_{trans} + T_{process}$  that is an upper bound on the time that can elapse between sending a message to another process and receiving a response.
- If no response arrives within time  $T$ , recipient of request has failed.


# Bully Algorithm



# Bully Algorithm


- The **process** that knows it has the **highest identifier** can **elect itself** as the **coordinator** simply by **sending** a **coordinator message** to all processes with **lower identifiers**.
- On the other hand, a process with a **lower identifier** can begin an **election** by **sending** an **election message** to those processes that have a **higher identifier** and **awaiting answer** messages in response.
- If **none arrives within** time **T**, the process considers **itself** the **coordinator** and sends a **coordinator message** to all **processes** with **lower identifiers** announcing this.
- If a process  $p_i$  receives a **coordinator message**, it **sets** its variable **elected<sub>i</sub>** to the **identifier** of the **coordinator**
- If a process **receives** an **election message**, it sends back an answer message and **begins another election**

# Bully Example




2 calls election

3, 4 respond




15

## Bully Example, continued




3 calls election



16




# Bully Example, concluded



4 responds to 3

4 is the new leader



# Bully Algorithm

- Algorithm clearly **meets the liveness condition E2**, by the assumption of reliable message delivery.
- Algorithm is **not guaranteed to meet the safety condition E1** if processes that have crashed are replaced by processes with the same identifiers.
- A **process** that **replaces** a **crashed process**  $p$  may **decide** that it has the **highest identifier** just as another process (which has detected  $p$ 's crash) decides that it has the highest identifier.
- **Two processes** will therefore **announce themselves** as the **coordinator concurrently**.
- That is why this algorithm is called **Bully Algorithm**

# Bully Algorithm

- Performance of the algorithm, in the **best case** the process with the **second-highest identifier notices** the **coordinator's failure**.
- Then it can immediately elect itself and send  **$N - 2$**  coordinator messages. The turnaround time is one message.
- The bully algorithm requires  **$O(N^2)$**  messages in the worst case – that is, when the process with the **lowest identifier** first **detects** the **coordinator's failure**.
- For then  **$N - 1$**  processes altogether **begin elections**, each sending messages to processes with higher identifiers.

# Summary

## Ring-based vs Bully

	Ring Based	Bully
Asynchronous	Yes	No
Allows processes to crash	No	Yes
Satisfies Safety	Yes	Yes/No
Dynamic process identifiers	Yes	No
Dynamic configuration of processes	Maybe	Maybe
Best case performance	$2 \times N$	$N - 1$
Worst case performance	$3 \times N - 1$	$O(N^2)$

# Summary

- Electing process with highest identifier as a co-ordinator process.
- Ring-based election algorithm
- Bully Algorithm