

NAME CACHES AND LDAP

- K.SRIRAGHAV – 3rd year CSE

Need for name caches

- In UNIX systems, the system call frequency for name-mapping operations (open, stat, lstat) constitute over 50% of the file system calls.
- Researchers also observed that in a large distributed system, a substantial portion of the network traffic is naming related.
- Hence there is a need for client to cache the result of a name resolution operation for a while, rather than repeating it every time the value is needed.

What is name cache?

- Caching name servers (*DNS caches*) store DNS query results for a period of time determined in the configuration (time-to-live) of each domain-name record.
- **Advantages:**
- Name caches improve the efficiency of the DNS by **reducing DNS traffic** across the distributed system
- Reducing load on authoritative name-servers, particularly root name-servers.
- Because they can answer questions more quickly, they also increase the performance of end-user applications that use the DNS.

Characteristics of name caches

- **High degree of locality of name lookup:**
 - The property of “locality of reference” has been observed in program execution , database access and file access.
 - Due to this locality feature, a reasonable size name cache , used for caching recently used information, can provide excellent hit ratios.
- **Slow update of name information database:**
 - Naming data has a high read-to-modify ratio.
 - This behaviour implies that the cost of maintaining the consistency of cached data is significantly low.

Characteristics of name caches - [contd]

- **On-use consistency of cached information:**
 - Name cache consistency can be maintained by detecting and discarding stale cache entries on use.
 - With On-use consistency checking, there is no need to invalidate all related cache entries when a naming data update occurs .

Types of name caches

- **Directory cache:**

- Each entry consists of a directory page.
- Used in systems that use the iterative method of name resolution.
- All recently used directory pages that are brought to the client node during name resolution are cached for a while.
- Ex : LOCUS
- **Disadv:** For only one useful entry of a directory, an entire page of a directory blocks a large area of precious cache space.

Types of name caches-[contd]

- **Prefix cache:**

- Used in naming systems that use the zone-based context distribution mechanism.
- Each entry consists of a name prefix and the corresponding zone identifier.
- Ex: Sprite , V-System
- **Disadv** - Not suitable to use with structure-free context distribution approach.

Types of name caches - [contd]

- **Full-name cache:**

- Each entry consists of an object's full pathname and the identifier and location of its authoritative name server.
- So, requests for accessing an object whose name is available in local cache can be directly sent to authoritative name server
- Ex : Galaxy
- **Adv** - Used with any naming system
- **Disadv** - Larger in size than prefix name caches.

Approaches for name cache implementation

- **Cache per process:**
 - Separate name cache is maintained for each process.
 - Each cache is maintained in the corresponding process's address space.
 - Therefore, accessing of cached information is fast and no memory area of the OS is occupied by the name caches.
 - Every process must create its own name cache from scratch.
 - Hence if the processes are short-lived, caches will have short lifetimes.
 - Hit ratio will be less for process-oriented caches due to start-up misses.
 - Ex. Sprite and V-System

Approaches for name cache

implementation - s[contd]

- **A single cache for all processes of a node**
 - In this approach, a single name cache is maintained at each node for all the processes of that node.
 - When compared to the process-oriented name caches, these caches are larger in size.
 - Accessing of information is slow
 - However, cached information in a single name cache is long lived – therefore no start-up- misses.
 - So higher average hit ratio as compared to process-oriented caches.
 - Ex : Locus and Galaxy

Multicache consistency

- When a data update , related name cache entries become stale and must be invalidated or updated properly.
- Two methods for Multicache consistency are :
 - Immediate invalidate
 - On-use update
- **On-use update :**
 - Most commonly used method for maintaining name consistency.
 - No attempt is made to invalidate all related cache entries when a naming update occurs.
 - When a client uses a stale cached data, it is informed by the naming system that data being used is either incorrectly specified.
 - On receiving a negative reply, necessary steps are taken to obtain the updated data.

Immediate invalidate

First approach

- Whenever a naming update operation is done, an invalidate message identifying the data to be invalidated is broadcast to all the nodes in the system
- Each node's name cache are then examined for the presence of updated data
- If it is present, the corresponding cache entry is invalidated.

Second approach

- The storage node of naming data keeps a list of nodes on which the data is cached.
- When a storage node receives a request for a naming data update, only the nodes in the corresponding list are notified to invalidate.

LDAP

Lightweight Directory Access Protocol

What is LDAP ? What is directory?

- LDAP is an application protocol for querying and modifying directory services
 - Runs over TCP/IP
- **Directory**
 - A set of objects with similar attributes
 - Organized in a logical and hierarchical manner
 - Example:
 - Telephone directory
 - Series of names (either of persons or organizations)
 - Organized alphabetically
 - Each name has an address and phone number
 - LDAP is often used by other services for authentication
 - **LDAP directory tree is drawn.**

LDAP – Attribute Naming

Attribute	Abbr.	Value
Country	C	NL
Locality	L	Amsterdam
Organization	O	Vrije Universiteit
OrganizationalUnit	OU	Comp. Sc.
CommonName	CN	Main server
Mail_Servers	—	137.37.20.3, 130.37.24.6, 137.37.20.10
FTP_Server	—	130.37.20.20
WWW_Server	—	130.37.20.20

Figure 5-22. A simple example of an LDAP directory entry using LDAP naming conventions.

Origin and influences


- X.500 directory services
 - Traditionally accessed via the X.500 Directory Access Protocol (DAP)
 - Required the Open Systems Interconnection (OSI) protocol stack
- LDAP originally intended to be a "lightweight" alternative protocol for accessing X.500 directory services
 - Through the simpler (and now widespread) TCP/IP protocol stack

Directory structure


- Protocol accesses LDAP directories
 - Follows the 1993 edition of the [X.500](#) model:
 - directory is a tree of directory entries
 - Entry consists of a set of attributes
 - An attribute has
 - a name
 - an *attribute type* or *attribute description*
 - one or more values
 - Attributes are defined in a *schema*
 - Each entry has a unique identifier:
 - *Distinguished Name* (DN)
 - Consists of its *Relative Distinguished Name* (RDN) constructed from some attribute(s) in the entry
 - Followed by the parent entry's DN
 - Think of the DN as a full filename and the RDN as a relative filename in a folder

LDAP Data Structure

Hierarchical



Flat



dc: domain component
ou: organizational unit

Protocol overview

- Client starts an LDAP session by connecting to an LDAP server
 - Default on TCP port 389
- Client sends operation requests to the server
 - Server sends responses in turn
- With some exceptions the client need not wait for a response before sending the next request
 - Server may send the responses in any order

Protocol overview

- The client may request the following operations:
 - Start TLS
 - Optionally protect the connection with [Transport Layer Security](#) (TLS), to have a more secure connection
 - Bind - [authenticate](#) and specify LDAP protocol version
 - Search - search for and/or retrieve directory entries
 - Compare - test if a named entry contains a given attribute value
 - Add a new entry
 - Delete an entry
 - Modify an entry
 - Modify Distinguished Name (DN) - move or rename an entry
 - Abandon - abort a previous request
 - Extended Operation - generic operation used to define other operations

LDAP BACKENDS

- THE BASIC DAEMON PROCESS THAT RUNS ON THE LDAP SERVER CALLED SLAPD COMES WITH THREE DIFFERENT BACKEND DATABASES
- WE ASSUME THAT IN OUR CASE WE USE LDBM THE MOST USED ONE

LDIF

- LDIF STANDS FOR LDAP DATA INTERCHANGE FORMAT
- DIRECTORY ENTRIES IN LDAP ARE IN THE FORM OF LDIF

LDIF FORMAT

- BASIC FORM OF LDIF :
#COMMENT
DN: <DISTINGUSHED NAME>
<ATTRDESC>: <ATTRVALUE>
<ATTRDESC>: <ATTRVALUE>
.....
- EXAMPLE :
 DN: UID=ALAKESH DC=IIT
 DC=EDU

Reference

1. **Authors :**Wien-Verteilte_Systeme_VO_(Göschka)_-
Tannenbaum

Book name: Distributed systems principles and paradigms
2nd_edition

Page (237/705)