

Consensus

Three Kinds

- ▶ The problems of mutual exclusion, electing a nominee and multicast are all instances of the more general problem of consensus.
- ▶ Consensus problems more generally then are described as one of three kinds:
 1. Consensus
 2. Byzantine Generals
 3. Interactive Consensus

Global Agreement

Consensus

- ▶ A set of processes $\{p_1, p_2, \dots, p_n\}$ each begins in the *undecided* state
- ▶ Each proposes a single value v_i
- ▶ The processes then communicate, exchanging values
- ▶ To conclude, each process must set their decision variable d_i to one value and thus enter the *decided* state
- ▶ Three desired properties:
 - ▶ Termination: each process sets its *decision_i* variable
 - ▶ Agreement: If p_i and p_j are correct processes and have both entered the *decided* state, then $d_i = d_j$
 - ▶ Integrity: If the correct processes all proposed the same value v , then any correct process p_i in the *decided* state has $d_i = v$

Global Agreement

Byzantine Generals

- ▶ Imagine three or more generals are to decide whether or not to attack
- ▶ We assume that there is a commander who issues the order
- ▶ The others must decide whether or not to attack
- ▶ Either the lieutenants or the commander can be faulty and thus send incorrect values
- ▶ Three desired properties:
 - ▶ Termination: each process sets its *decision_i* variable
 - ▶ Agreement: If p_i and p_j are correct processes and have both entered the *decided* state, then $d_i = d_j$
 - ▶ Integrity: If the commander is correct then all correct processes decide on the value proposed by the commander
- ▶ When the commander is correct, *Integrity* implies *Agreement*, but the commander may not be correct

Global Agreement

Interactive Consensus

- ▶ Each process proposes its own value and the goal is for each process to agree on a vector of values
- ▶ Similar to consensus other than that each process contributes only a part of the final answer which we call the *decision vector*
- ▶ Three desired properties:
 - ▶ Termination: each process sets its *decision_i* variable
 - ▶ Agreement: The final decision vector of all processes is the same
 - ▶ Integrity: If p_i is correct and proposes v_i then all correct processes decide on v_i as the i th component of the decision vector


Global Agreement

Relating the three

- ▶ Assuming we had a solution to any of the three problems we could construct a solution to the other two
- ▶ For example, if we have a solution to Interactive Consensus, then we have a solution to Consensus, all we require is some way consistent function for choosing a single component of the decision vector
 - ▶ We might choose a majority function, maximum, minimum or some other function depending on the application
 - ▶ It only requires that the function is context independent
- ▶ If we have a solution to the Byzantine Generals then we can construct a solution to Interactive Consensus
 - ▶ To do so we simply run the Byzantine Generals solution N times, once for each process
- ▶ The point is not necessarily that this would be the way to implement such as solution (it may not be efficient)
 - ▶ However if we can determine an impossibility result for one of these problems we know that we also have the same result for the others


Global Agreement

Byzantine Generals in a Synchronous System



Global Agreement

Byzantine Generals in a Synchronous System



Global Agreement

Impossible

- ▶ Recall:
 - ▶ Agreement: If p_i and p_j are correct processes and have both entered the *decided* state, then $d_i = d_j$
 - ▶ Integrity: If the commander is correct then all correct processes decide on the value proposed by the commander
- ▶ In both scenarios, process p_2 receives different values from the commander p_1 and the other process p_3
- ▶ It can therefore know that one process is faulty but cannot know which one
- ▶ By the *Integrity* property then it is bound to choose the value given by the commander
- ▶ By symmetry the process p_3 is in the same situation when the commander is faulty.
- ▶ Hence when the commander is faulty there is no way to satisfy the *Agreement* property, so no solution exists for three processes

Global Agreement

$$N \leq 3 \times f$$

- ▶ In the above case we had three processes and at most one incorrect process, hence $N = 3$ and $f = 1$
- ▶ It has been shown, by Pease *et al* that more generally no solution can exist whenever $N \leq 3 \times f$
- ▶ However there can exist a solution whenever $N > 3 \times f$
- ▶ Such algorithms consist of *rounds* of messages
- ▶ It is known that such algorithms require at least $f + 1$ message rounds
- ▶ The complexity and cost of such algorithms suggest that they are only applicable where the threat is great
- ▶ That means either the threat of an incorrect or malicious process is great
- ▶ and/or the cost of failing due to inability to reach consensus is large

Global Agreement

Consensus in an Asynchronous System

- ▶ Fisher *et al* have shown that it is impossible to design an algorithm which is guaranteed to reach consensus in an asynchronous system, under the following condition:

Global Agreement

Consensus in an Asynchronous System

- ▶ Fisher *et al* have shown that it is impossible to design an algorithm which is guaranteed to reach consensus in an asynchronous system, under the following condition:
 - ▶ We allow a single process crash failure
- ▶ Even if we have 1000s of processes, and the failure is a crash rather than an arbitrary failure of just a single process, any consensus algorithm is not guaranteed to reach consensus
- ▶ Clearly this is a pretty benign set of circumstances
- ▶ We therefore know that there is no solution in an asynchronous system to either:
 1. Byzantine generals (and hence consensus or interactive consensus)
 2. Totally order and reliable multicast

Consensus in an Asynchronous System

So what to do?

- ▶ The important word in the previous impossibility result is: guarantee
- ▶ There is no algorithm which is guaranteed to reach consensus
- ▶ Consensus has been reached in asynchronous systems for years
- ▶ Some techniques for getting around the impossibility result:
 - ▶ Masking process failures, for example using persistent storage such that a crashed process can be replaced by one in effectively the same state
 - ▶ Thus meaning some operations appear to take a long time, but all operations do eventually complete
 - ▶ Employ failure detectors:
 - ▶ Although in an asynchronous system we cannot achieve a reliable failure detector
 - ▶ We can use one which is “perfect by design”
 - ▶ Once a process is deemed to have failed, any subsequent messages that it does send (showing that it had not failed) are ignored
 - ▶ To do this the other processes must agree that a given process has failed